ARMv8 Instruction Set Overview
Architecture Group

Document number:	**PRD03-GENC-010197 15.0**
Date of Issue:	11 November 2011

Abstract

This document provides a high-level overview of the ARMv8 instructions sets, being mainly the new A64 instruction set used in AArch64 state but also those new instructions added to the A32 and T32 instruction sets since ARMv7-A for use in AArch32 state. For A64 this document specifies the preferred architectural assembly language notation to represent the new instruction set.

Keywords

AArch64, A64, AArch32, A32, T32, ARMv8

Proprietary Notice

This specification is protected by copyright and the practice or implementation of the information herein may be protected by one or more patents or pending applications. No part of this specification may be reproduced in any form by any means without the express prior written permission of ARM. **No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this specification**.

Your access to the information in this specification is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations of the ARM architecture infringe any third party patents.

This specification is provided "as is". ARM makes no representations or warranties, either express or implied, included but not limited to, warranties of merchantability, fitness for a particular purpose, or non-infringement, that the content of this specification is suitable for any particular purpose or that any practice or implementation of the contents of the specification will not infringe any third party patents, copyrights, trade secrets, or other rights.

This specification may include technical inaccuracies or typographical errors.

To the extent not prohibited by law, in no event will ARM be liable for any damages, including without limitation any direct loss, lost revenue, lost profits or data, special, indirect, consequential, incidental or punitive damages, however caused and regardless of the theory of liability, arising out of or related to any furnishing, practicing, modifying or any use of this specification, even if ARM has been advised of the possibility of such damages.

Words and logos marked with ® or TM are registered trademarks or trademarks of ARM Limited, except as otherwise stated below in this proprietary notice. Other brands and names mentioned herein may be the trademarks of their respective owners.

Copyright © 2009-2011 ARM Limited

110 Fulbourn Road, Cambridge, England CB1 9NJ

Restricted Rights Legend: Use, duplication or disclosure by the United States Government is subject to the restrictions set forth in DFARS 252.227-7013 (c)(1)(ii) and FAR 52.227-19.

This document is Non-Confidential but any disclosure by you is subject to you providing notice to and the acceptance by the recipient of, the conditions set out above.

In this document, where the term ARM is used to refer to the company it means "ARM or any of its subsidiaries as appropriate".

Contents

1 ABOUT THIS DOCUMENT

1.1 Change control

1.1.1 Current status and anticipated changes

This document is a *beta* release specification and further changes to correct defects and improve clarity should be expected.

1.1.2 Change history

Issue	Date	By	Change
		NJS	Previous releases tracked in Domino
7.0	17th December 2010	NJS	Beta0 release
8.0	25th February 2011	NJS	Beta0 update 1
9.0	20th April 2011	NJS	Beta1 release
10.0	14th July 2011	NJS	Beta2 release
11.0	9th September 2011	NJS	Beta2 update 1
12.0	28th September 2011	NJS	Beta3 release
13.0	28th October 2011	NJS	Beta3 update 1
14.0	28th October 2011	NJS	Restructured and incorporated new AArch32 instructions.
15.0	11th November 2011	NJS	First non-confidential release. Describe partial deprecation of the `IT` instruction. Rename `DRET` to `DRPS` and clarify its behavior.

1.2 References

This document refers to the following documents.

Reference	Author	Document number	Title
[v7A]	ARM	ARM DDI 0406	ARM® Architecture Reference Manual, ARMv7-A and ARMv7-R edition
[AES]	NIST	FIPS 197	Announcing the Advanced Encryption Standard (AES)
[SHA]	NIST	FIPS 180-2	Announcing the Secure Hash Standard (SHA)
[GCM]	McGrew and Viega	n/a	The Galois/Counter Mode of Operation (GCM)

1.3 Terms and abbreviations

This document uses the following terms and abbreviations.

Term	Meaning
AArch64	The 64-bit general purpose register width state of the ARMv8 architecture.
AArch32	The 32-bit general purpose register width state of the ARMv8 architecture, broadly compatible with the ARMv7-A architecture.
	Note: The register width state can change only upon a change of exception level.
A64	The new instruction set available when in AArch64 state, and described in this document.
A32	The instruction set named ARM in the ARMv7 architecture, which uses 32-bit instructions. The new A32 instructions added by ARMv8 are described in §6.
T32	The instruction set named Thumb in the ARMv7 architecture, which uses 16-bit and 32-bit instructions. The new T32 instructions added by ARMv8 are described in §6.
UNALLOCATED	Describes an opcode or combination of opcode fields which do not select a valid instruction at the current privilege level. Executing an UNALLOCATED encoding will usually result in taking an Undefined Instruction exception.
RESERVED	Describes an instruction field value within an otherwise allocated instruction which should not be used within this specific instruction context, for example a value which selects an unsupported data type or addressing mode. An instruction encoding which contains a RESERVED field value is an UNALLOCATED encoding.

2 INTRODUCTION

This document provides an overview of the ARMv8 instruction sets. Most of the document forms a description of the new A64 instruction set used when the processor is operating in AArch64 register width state, and defines its preferred architectural assembly language.

Section 6 below lists the extensions introduced by ARMv8 to the A32 and T32 instruction sets – known in ARMv7 as the ARM and Thumb instruction sets respectively – which are available when the processor is operating in AArch32 register width state. The A32 and T32 assembly language syntax is unchanged from ARMv7.

In the syntax descriptions below the following conventions are used:

UPPER   UPPER-CASE text is fixed, while lower-case text is variable. So register name Xn indicates that the `X' is required, followed by a variable register number, e.g. X29.

< >   Any item bracketed by < and > is a short description of a type of value to be supplied by the user in that position. A longer description of the item is normally supplied by subsequent text.

{ }   Any item bracketed by curly braces { and } is optional. A description of the item and of how its presence or absence affects the instruction is normally supplied by subsequent text. In some cases curly braces are actual symbols in the syntax, for example surrounding a register list, and such cases will be called out in the surrounding text.

[ ]   A list of alternative characters may be bracketed by [and]. A single one of the characters can be used in that position and the the subsequent text will describe the meaning of the alternatives. In some cases the symbols [and] are part of the syntax itself, such as addressing modes and vector elements, and such cases will be called out in the surrounding text.

a | b   Alternative words are separated by a vertical bar | and may be surrounded by parentheses to delimit them, e.g. U(ADD|SUB)W represents UADDW or USUBW.

+/-   This indicates an optional + or - sign. If neither is coded, then + is assumed.

3 A64 OVERVIEW

The A64 instruction set provides similar functionality to the A32 and T32 instruction sets in AArch32 or ARMv7. However just as the addition of 32-bit instructions to the T32 instruction set rationalized some of the ARM ISA behaviors, the A64 instruction set includes further rationalizations. The highlights of the new instruction set are as follows:

- A clean, fixed length instruction set – instructions are 32 bits wide, register fields are contiguous bit fields at fixed positions, immediate values mostly occupy contiguous bit fields.

- Access to a larger general-purpose register file with 31 unbanked registers (0-30), with each register extended to 64 bits. General registers are encoded as 5-bit fields with register number 31 (0b11111) being a special case representing:

  - *Zero Register*: in most cases register number 31 *reads as zero* when used as a source register, and *discards the result* when used as a destination register.

  - *Stack Pointer*: when used as a load/store base register, and in a small selection of arithmetic instructions, register number 31 provides access to the *current stack pointer*.

- The PC is never accessible as a named register. Its use is implicit in certain instructions such as PC-relative load and address generation. The only instructions which cause a non-sequential change to the PC are the designated Control Flow instructions (see §5.1) and exceptions. The PC cannot be specified as the destination of a data processing instruction or load instruction.

- The procedure call link register (LR) is unbanked, general-purpose register 30; exceptions save the restart PC to the target exception level's ELR system register.

- Scalar load/store addressing modes are uniform across all sizes and signedness of scalar integer, floating point and vector registers.

- A load/store immediate offset may be scaled by the access size, increasing its effective offset range.

- A load/store index register may contain a 64-bit or 32-bit signed/unsigned value, optionally scaled by the access size.

- Arithmetic instructions for address generation which mirror the load/store addressing modes, see §3.3.

- PC-relative load/store and address generation with a range of ±4GiB is possible using just two instructions without the need to load an offset from a literal pool.

- PC-relative offsets for literal pool access and most conditional branches are extended to ±1MiB, and for unconditional branches and calls to ±128MiB.

- There are no multiple register LDM, STM, PUSH and POP instructions, but load-store of a non-contiguous pair of registers is available.

- Unaligned addresses are permitted for most loads and stores, including paired register accesses, floating point and SIMD registers, with the exception of exclusive and ordered accesses (see §3.5.2).

- Reduced conditionality. Fewer instructions can set the condition flags. Only conditional branches, and a handful of data processing instructions read the condition flags. Conditional or predicated execution is not provided, and there is no equivalent of T32's IT instruction (see §3.2).

- A shift option for the final register operand of data processing instructions is available:
    - Immediate shifts only (as in T32).
    - No RRX shift, and no ROR shift for ADD/SUB.
    - The ADD/SUB/CMP instructions can first sign or zero-extend a byte, halfword or word in the final register operand, followed by an optional left shift of 1 to 4 bits.

- Immediate generation replaces A32's rotated 8-bit immediate with operation-specific encodings:
    - Arithmetic instructions have a simple 12-bit immediate, with an optional left shift by 12.
    - Logical instructions provide sophisticated replicating bit mask generation.
    - Other immediates may be constructed inline in 16-bit "chunks", extending upon the MOVW and MOVT instructions of AArch32.

- Floating point support is similar to AArch32 VFP but with some extensions, as described in §3.6.

- Floating point and Advanced SIMD processing share a register file, in a similar manner to AArch32, but extended to thirty-two 128-bit registers. Smaller registers are no longer packed into larger registers, but are mapped one-to-one to the low-order bits of the 128-bit register, as described in §4.4.2.

- There are no SIMD or saturating arithmetic instructions which operate on the general purpose registers, such operations being available only as part of the Advanced SIMD processing, described in §5.7.

- There is no access to CPSR as a single register, but new system instructions provide the ability to atomically modify individual processor state fields, see §5.8.2.

- The concept of a "coprocessor" is removed from the architecture. A set of system instructions described in §5.8 provides:
    - System register access
    - Cache/TLB management
    - VA⇔PA translation
    - Barriers and CLREX
    - Architectural hints (WFI, etc)
    - Debug

3.1 Distinguishing 32-bit and 64-bit Instructions

Most integer instructions in the A64 instruction set have two forms, which operate on either 32-bit or 64-bit values within the 64-bit general-purpose register file. Where a 32-bit instruction form is selected, the following holds true:

- The upper 32 bits of the source registers are ignored;
- The upper 32 bits of the destination register are set to ZERO;
- Right shifts/rotates inject at bit 31, instead of bit 63;
- The condition flags, where set by the instruction, are computed from the lower 32 bits.

This distinction applies even when the result(s) of a 32-bit instruction form would be indistinguishable from the lower 32 bits computed by the equivalent 64-bit instruction form. For example a 32-bit bitwise `ORR` could be performed using a 64-bit `ORR`, and simply ignoring the top 32 bits of the result. But the A64 instruction set includes separate 32 and 64-bit forms of the `ORR` instruction.

Rationale: The C/C++ LP64 and LLP64 data models – expected to be the most commonly used on AArch64 – both define the frequently used `int`, `short` and `char` types to be 32 bits or less. By maintaining this semantic information in the instruction set, implementations can exploit this information to avoid expending energy or cycles to compute, forward and store the unused upper 32 bits of such data types. Implementations are free to exploit this freedom in whatever way they choose to save energy.

As well as distinct sign/zero-extend instructions, the A64 instruction set also provides the ability to extend and shift the final source register of an `ADD`, `SUB` or `CMP` instruction and the index register of a load/store instruction. This allows for an efficient implementation of array index calculations involving a 64-bit array pointer and 32-bit array index.

The assembly language notation is designed to allow the identification of registers holding 32-bit values as distinct from those holding 64-bit values. As well as aiding readability, tools may be able to use this to perform limited type checking, to identify programming errors resulting from the change in register size.

3.2 Conditional Instructions

The A64 instruction set does not include the concept of predicated or conditional execution. Benchmarking shows that modern branch predictors work well enough that predicated execution of instructions does not offer sufficient benefit to justify its significant use of opcode space, and its implementation cost in advanced implementations.

A very small set of "conditional data processing" instructions are provided. These instructions are unconditionally executed but use the condition flags as an extra input to the instruction. This set has been shown to be beneficial in situations where conditional branches predict poorly, or are otherwise inefficient.

The conditional instruction types are:

- Conditional branch: the traditional ARM conditional branch, together with compare and branch if register zero/non-zero, and test single bit in register and branch if zero/non-zero – all with increased displacement.
- Add/subtract with carry: the traditional ARM instructions, for multi-precision arithmetic, checksums, etc.
- Conditional select with increment, negate or invert: conditionally select between one source register and a second incremented/negated/inverted/unmodified source register. Benchmarking reveals these to be the highest frequency uses of single conditional instructions, e.g. for counting, absolute value, etc. These instructions also implement:
    - Conditional select (move): sets the destination to one of two source registers, selected by the condition flags. Short conditional sequences can be replaced by unconditional instructions followed by a conditional select.
    - Conditional set: conditionally select between 0 and 1 or -1, for example to materialize the condition flags as a Boolean value or mask in a general register.
- Conditional compare: sets the condition flags to the result of a comparison if the original condition was true, else to an immediate value. Permits the flattening of nested conditional expressions without using conditional branches or performing Boolean arithmetic within general registers.

3.3 Addressing Features

The prime motivation for a 64-bit architecture is access to a larger virtual address space. The AArch64 memory translation system supports a 49-bit virtual address (48 bits per translation table). Virtual addresses are sign-extended from 49 bits, and stored within a 64-bit pointer. Optionally, under control of a system register, the most significant 8 bits of a 64-bit pointer may hold a "tag" which will be ignored when used as a load/store address or the target of an indirect branch.

3.3.1 Register Indexed Addressing

The A64 instruction set extends on 32-bit T32 addressing modes, allowing a 64-bit index register to be added to the 64-bit base register, with optional scaling of the index by the access size. Additionally it provides for sign or zero-extension of a 32-bit value within an index register, again with optional scaling.

These register index addressing modes provide a useful performance gain if they can be performed within a single cycle, and it is believed that at least some implementations will be able to do this. However, based on implementation experience with AArch32, it is expected that other implementations will need an additional cycle to execute such addressing modes.

Rationale: The architects intend that implementations should be free to fine-tune the performance trade-offs within each implementation, and note that providing an instruction which in some implementations takes two cycles, is preferable to requiring the dynamic grouping of two independent instructions in an implementation that can perform this address arithmetic in a single cycle.

3.3.2 PC-relative Addressing

There is improved support for position-independent code and data addressing:

- PC-relative literal loads have an offset range of ±1MiB. This permits fewer literal pools, and more sharing of literal data between functions – reducing I-cache and TLB pollution.

- Most conditional branches have a range of ±1MiB, expected to be sufficient for the majority of conditional branches which take place within a single function.

- Unconditional branches, including branch and link, have a range of ±128MiB. Expected to be sufficient to span the static code segment of most executable load modules and shared objects, without needing linker-inserted trampolines or "veneers".

- PC-relative load/store and address generation with a range of ±4GiB may be performed inline using only two instructions, i.e. without the need to load an offset from a literal pool.

3.4 The Program Counter (PC)

The current Program Counter (PC) cannot be referred to by number as if part of the general register file and therefore cannot be used as the source or destination of arithmetic instructions, or as the base, index or transfer register of load/store instructions. The only instructions which read the PC are those whose function is to compute a PC-relative address (ADR, ADRP, literal load, and direct branches), and the branch-and-link instructions which store it in the link register (BL and BLR). The only way to modify the Program Counter is using explicit control flow instructions: conditional branch, unconditional branch, exception generation and exception return instructions.

Where the PC is read by an instruction to compute a PC-relative address, then its value is the address of the instruction, i.e. unlike A32 and T32 there is no implied offset of 4 or 8 bytes.

3.5 Memory Load-Store

3.5.1 Bulk Transfers

The LDM, STM, PUSH and POP instructions do not exist in A64, however bulk transfers can be constructed using the LDP and STP instructions which load and store a pair of independent registers from consecutive memory locations, and which support unaligned addresses when accessing normal memory. The LDNP and

STNP instructions additionally provide a "streaming" or "non-temporal" hint that the data does not need to be retained in caches. The PRFM (prefetch memory) instructions also include hints for "streaming" or "non-temporal" accesses, and allow targeting of a prefetch to a specific cache level.

3.5.2 Exclusive Accesses

Exclusive load-store of a byte, halfword, word and doubleword. Exclusive access to a pair of doublewords permit atomic updates of a pair of pointers, for example circular list inserts. All exclusive accesses must be naturally aligned, and exclusive pair access must be aligned to twice the data size (i.e. 16 bytes for a 64-bit pair).

3.5.3 Load-Acquire, Store-Release

Explicitly synchronising load and store instructions implement the release-consistency (RCsc) memory model, reducing the need for explicit memory barriers, and providing a good match to emerging language standards for shared memory. The instructions exist in both exclusive and non-exclusive forms, and require natural address alignment. See §5.2.8 for more details.

3.6 Integer Multiply/Divide

Including 32 and 64-bit multiply, with accumulation:

- $32 \pm (32 \times 32) \rightarrow 32$
- $64 \pm (64 \times 64) \rightarrow 64$
- $\pm (32 \times 32) \rightarrow 32$
- $\pm (64 \times 64) \rightarrow 64$

Widening multiply (signed and unsigned), with accumulation:

- $64 \pm (32 \times 32) \rightarrow 64$
- $\pm (32 \times 32) \rightarrow 64$
- $(64 \times 64) \rightarrow \text{hi64} <127:64>$

Multiply instructions write a single register. A 64×64 to 128-bit multiply requires a sequence of two instructions to generate a pair of 64-bit result registers:

- $+ (64 \times 64) \rightarrow <63:0>$
- $(64 \times 64) \rightarrow <127:64>$

Signed and unsigned 32- and 64-bit divide are also provided. A remainder instruction is not provided, but a remainder may be computed easily from the dividend, divisor and quotient using an MSUB instruction. There is no hardware check for "divide by zero", but this check can be performed in the shadow of a long latency division. A divide by zero writes zero to the destination register.

3.7 Floating Point

AArch64 mandates hardware floating point wherever floating point arithmetic is required – there is no "soft-float" variant of the AArch64 Procedure Calling Standard (PCS).

Floating point functionality is similar to AArch32 VFP, with the following changes:

- The deprecated "small vector" feature of VFP is removed.
- There are 32 S registers and 32 D registers. The S registers are not packed into D registers, but occupy the low 32 bits of the corresponding D register. For example S31=D31<31:0>, not D15<63:32>.
- Load/store addressing modes identical to integer load/stores.
- Load/store of a pair of floating point registers.
- Floating point FCSEL and FCCMP equivalent to the integer CSEL and CCMP.

- Floating point `FCMP` and `FCCMP` instructions set the integer condition flags directly, and do not modify the condition flags in the FPSR.

- All floating-point multiply-add and multiply-subtract instructions are "fused".

- Convert between 64-bit integer and floating point.

- Convert FP to integer with explicit rounding direction (towards zero, towards +Inf, towards -Inf, to nearest with ties to even, and to nearest with ties away from zero).

- Round FP to nearest integral FP with explicit rounding direction (as above).

- Direct conversion between half-precision and double-precision.

- `FMINNM` & `FMAXNM` implementing the IEEE754-2008 `minNum()` and `maxNum()` operations, returning the numerical value if one of the operands is a quiet NaN.

3.8 Advanced SIMD

See §5.7 below for a detailed description.

4 A64 ASSEMBLY LANGUAGE

4.1 Basic Structure

The letter `W` is shorthand for a 32-bit *word*, and `X` for a 64-bit *extended word*. The letter `X` (*extended*) is used rather than `D` (*double*), since `D` conflicts with its use for floating point and SIMD "double-precision" registers and the T32 load/store "double-register" instructions (e.g. `LDRD`).

An A64 assembler will recognise both upper and lower-case variants of instruction mnemonics and register names, but not mixed case. An A64 disassembler may output either upper or lower-case mnemonics and register names. The case of program and data labels **is** significant.

The fundamental statement format and operand order follows that used by AArch32 UAL assemblers and disassemblers, i.e. a single statement per source line, consisting of one or more optional program labels, followed by an instruction mnemonic, then a *destination* register and one or more *source* operands separated by commas.

```
{label:*} {opcode {dest{, source1{, source2{, source3}}}}}}
```

This dest/source ordering is reversed for store instructions, in common with AArch32 UAL.

The A64 assembly language does not require the '#' symbol to introduce immediate values, though an assembler must allow it. An A64 disassembler shall always output a '#' before an immediate value for readability.

Where a user-defined symbol or label is identical to a pre-defined register name (e.g. "X0") then if it is used in a context where its interpretation is ambiguous – for example in an operand position that would accept either a register name or an immediate expression – then an assembler must interpret it as the register name. A symbol may be disambiguated by using it within an expression context, i.e. by placing it within parentheses and/or prefixing it with an explicit '#' symbol.

In the examples below the sequence "//" is used as a comment leader, though A64 assemblers are also expected to to support their legacy ARM comment syntax.

4.2 Instruction Mnemonics

An A64 instruction form can be identified by the following combination of attributes:

- The operation *name* (e.g. `ADD`) which indicates the instruction semantics.
- The operand *container*, usually the register type. An instruction writes to the whole container, but if it is not the largest in its class, then the remainder of the largest container in the class is set to ZERO.
- The operand data *subtype*, where some operand(s) are a different size from the primary *container*.
- The final source operand type, which may be a register or an immediate value.

The *container* is one of:

Integer Class	
W	32-bit integer
X	64-bit integer
SIMD Scalar & Floating Point Class	
B	8-bit scalar
H	16-bit scalar & half-precision float
S	32-bit scalar & single-precision float
D	64-bit scalar & double-precision float
Q	128-bit scalar

The *subtype* is one of:

Load-Store / Sign-Zero Extend	
B	byte
SB	signed byte
H	halfword
SH	signed halfword
W	word
SW	signed word
Register Width Changes	
H	High (dst gets top half)
N	Narrow (dst < src)
L	Long (dst > src)
W	Wide (dst == src1, src1 > src2)
etc	

These attributes are combined in the assembly language notation to identify the specific instruction form. In order to retain a close look and feel to the existing ARM assembly language, the following format has been adopted:

```
<name>{<subtype>}        <container>
```

In other words the operation *name* and *subtype* are described by the instruction mnemonic, and the *container* size by the operand name(s). Where *subtype* is omitted, it is inherited from *container*.

In this way an assembler programmer can write an instruction without having to remember a multitude of new mnemonics; and the reader of a disassembly listing can straightforwardly read an instruction and see at a glance the type and size of each operand.

The implication of this is that the A64 assembly language *overloads* instruction mnemonics, and distinguishes between the different forms of an instruction based on the operand register names. For example the ADD instructions below all have different opcodes, but the programmer only has to remember one mnemonic and the assembler automatically chooses the correct opcode based on the operands − with the disassembler doing the reverse.

```
ADD    W0, W1, W2              // add 32-bit register
ADD    X0, X1, X2              // add 64-bit register
ADD    X0, X1, W2, SXTW        // add 64-bit extending register
ADD    X0, X1, #42             // add 64-bit immediate
```

4.3 Condition Codes

In AArch32 assembly language conditionally executed instructions are represented by directly appending the condition to the mnemonic, without a delimiter. This leads to some ambiguity which can make assembler code difficult to parse: for example ADCS, BICS, LSLS and TEQ look at first glance like conditional instructions.

The A64 ISA has far fewer instructions which set or test condition codes. Those that do will be identified as follows:

1. Instructions which set the condition flags are notionally different instructions, and will continue to be identified by appending an 'S' to the base mnemonic, e.g. ADDS.

2. Instructions which are truly conditionally executed (i.e. when the condition is false they have no effect on the architectural state, aside from advancing the program counter) have the condition appended to the instruction with a '.' delimiter. For example `B.EQ`.
3. If there is more than one instruction extension, then the conditional extension is always last.
4. Where a conditional instruction has qualifiers, the qualifiers follow the condition.
5. Instructions which are unconditionally executed, but use the condition flags as a source operand, will specify the condition to test in their final operand position, e.g. `CSEL Wd,Wm,Wn,NE`

To aid portability an A64 assembler may also provide the old UAL conditional mnemonics, so long as they have direct equivalents in the A64 ISA. However, the UAL mnemonics will not be generated by an A64 disassembler – their use is deprecated in 64-bit assembler code, and may cause a warning or error if backward compatibility is not explicitly requested by the programmer.

The full list of condition codes is as follows:

Encoding	Name (& alias)	Meaning (integer)	Meaning (floating point)	Flags
0000	EQ	Equal	Equal	Z==1
0001	NE	Not equal	Not equal, or unordered	Z==0
0010	HS (CS)	Unsigned higher or same (Carry set)	Greater than, equal, or unordered	C==1
0011	LO (CC)	Unsigned lower (Carry clear)	Less than	C==0
0100	MI	Minus (negative)	Less than	N==1
0101	PL	Plus (positive or zero)	Greater than, equal, or unordered	N==0
0110	VS	Overflow set	Unordered	V==1
0111	VC	Overflow clear	Ordered	V==0
1000	HI	Unsigned higher	Greater than, or unordered	C==1 && Z==0
1001	LS	Unsigned lower or same	Less than or equal	!(C==1 && Z==0)
1010	GE	Signed greater than or equal	Greater than or equal	N==V
1011	LT	Signed less than	Less than or unordered	N!=V
1100	GT	Signed greater than	Greater than	Z==0 && N==V
1101	LE	Signed less than or equal	Less than, equal, or unordered	!(Z==0 && N==V)
1110	AL	Always	Always	Any
1111	NV[†]	Always	Always	Any

†The condition code `NV` exists only to provide a valid disassembly of the '1111b' encoding, and otherwise behaves identically to `AL`.

4.4 Register Names

4.4.1 General purpose (integer) registers

The thirty one general purpose registers in the main integer register bank are named `R0` to `R30`, with special register number 31 having different names, depending on the context in which it is used. However, when the registers are used in a specific instruction form, they must be further qualified to indicate the operand data size (32 or 64 bits) – and hence the instruction's data size.

The qualified names for the general purpose registers are as follows, where 'n' is the register number 0 to 30:

Size (bits)	32b	64b
Name	Wn	Xn

Where register number 31 represents *read zero* or *discard result* (aka the "*zero register")*:

Size (bits)	32b	64b
Name	WZR	XZR

Where register number 31 represents the *stack pointer*:

Size (bits)	32b	64b
Name	WSP	SP

In more detail:

- The names `Xn` and `Wn` refer to the same architectural register.
- There is no register named `W31` or `X31`.
- For instruction operands where register 31 in interpreted as the 64-bit *stack pointer*, it is represented by the name `SP`. For operands which do not interpret register 31 as the 64-bit stack pointer this name shall cause an assembler error.
- The name `WSP` represents register 31 as the *stack pointer* in a 32-bit context. It is provided only to allow a valid disassembly, and should not be seen in correctly behaving 64-bit code.
- For instruction operands which interpret register 31 as the *zero register,* it is represented by the name `XZR` in 64-bit contexts, and `WZR` in 32-bit contexts. In operand positions which do not interpret register 31 as the zero register these names shall cause an assembler error.
- Where a mnemonic is overloaded (i.e. can generate different instruction encodings depending on the data size), then an assembler shall determine the precise form of the instruction from the size of the *first* register operand. Usually the other operand registers should match the size of the first operand, but in some cases a register may have a different size (e.g. an address base register is always 64 bits), and a source register may be smaller than the destination if it contains a word, halfword or byte that is being widened by the instruction to 64 bits.
- The architecture does not define a special name for register 30 that reflects its special role as the link register on procedure calls. Such software names may be defined as part of the Procedure Calling Standard.

4.4.2 FP/SIMD registers

The thirty two registers in the FP/SIMD register bank named V0 to V31 are used to hold floating point operands for the scalar floating point instructions, and both scalar and vector operands for the Advanced SIMD instructions. As with the general purpose integer registers, when they are used in a specific instruction form the names must be further qualified to indicate the data *shape* (i.e. the data element size and number of elements or lanes) held within them.

Note however that the data *type*, i.e. the interpretation of the bits within each register or vector element – integer (signed, unsigned or irrelevant), floating point, polynomial or cryptographic hash – is not described by the register name, but by the instruction mnemonics which operate on them. For more details see the Advanced SIMD description in §5.7.

4.4.2.1 SIMD scalar register

In Advanced SIMD and floating point instructions which operate on scalar data the FP/SIMD registers behave similarly to the main general-purpose integer registers, i.e. only the lower bits are accessed, with the unused high bits ignored on a read and set to zero on a write. The qualified names for scalar FP/SIMD names indicate the number of significant bits as follows, where 'n' is a register number 0 to 31:

Size (bits)	8b	16b	32b	64b	128b
Name	Bn	Hn	Sn	Dn	Qn

4.4.2.2 SIMD vector register

When a register holds multiple data elements on which arithmetic will be performed in a parallel, SIMD fashion, then a qualifier describes the vector shape: i.e. the element size, and the number of elements or "lanes". Where "bits×lanes" does not equal 128, the upper 64 bits of the register are ignored when read and set to zero on a write.

The fully qualified SIMD vector register names are as follows, where 'n' is the register number 0 to 31:

Shape (bits×lanes)	8b×8	8b×16	16b×4	16b×8	32b×2	32b×4	64b×1	64b×2
Name	Vn.8B	Vn.16B	Vn.4H	Vn.8H	Vn.2S	Vn.4S	Vn.1D	Vn.2D

4.4.2.3 SIMD vector element

Where a single element from a SIMD vector register is used as a scalar operand, this is indicated by appending a constant, zero-based "element index" to the vector register name, inside square brackets. The number of lanes is not represented, since it is not encoded, and may only be inferred from the index value.

Size (bits)	8b	16b	32b	64b
Name	Vn.B[i]	Vn.H[i]	Vn.S[i]	Vn.D[i]

However an assembler shall accept a fully qualified SIMD vector register name as in §4.4.2.2, so long as the number of lanes is greater than the index value. For example the following forms will both be accepted by an assembler as the name for the 32-bit element in bits <63:32> of SIMD register 9:

V9.S[1]	*standard disassembly*
V9.2S[1]	*optional number of lanes*
V9.4S[1]	*optional number of lanes*

Note that the vector register element name `Vn.S[0]` is <u>not</u> equivalent to the scalar register name `Sn`. Although they represent the same bits in the register, they select different instruction encoding forms, i.e. vector element vs scalar form.

4.4.2.4 SIMD vector register list

Where an instruction operates on a "list" of vector registers – for example vector load-store and table lookup – the registers are specified as a list within curly braces. This list consists of either a sequence of registers separated by commas, or a register range separated by a hyphen. The registers must be numbered in increasing order (modulo 32), in increments of one or two. The hyphenated form is preferred for disassembly if there are more than two registers in the list, and the register numbers are monotonically increasing in increments of one. The following are equivalent representations of a set of four registers `V4` to `V7`, each holding four lanes of 32-bit elements:

```
{V4.4S – V7.4S}                    standard disassembly
{V4.4S, V5.4S, V6.4S, V7.4S}       alternative representation
```

4.4.2.5 SIMD vector element list

It is also possible for registers in a list to have a vector element form, for example `LD4` loading one element into each of four registers, in which case the index is appended to the list, as follows:

```
{V4.S - V7.S}[3]                   standard disassembly
{V4.4S, V5.4S, V6.4S, V7.4S}[3]    alternative with optional number of lanes
```

4.5 Load/Store Addressing Modes

Load/store addressing modes in the A64 instruction set broadly follows T32 consisting of a 64-bit `base` register (`Xn` or `SP`) plus an immediate or register offset.

Type	Immediate Offset	Register Offset	Extended Register Offset
Simple register (exclusive)	`[base{,#0}]`	n/a	n/a
Offset	`[base{,#imm}]`	`[base,Xm{,LSL #imm}]`	`[base,Wm,(S|U)XTW {#imm}]`
Pre-indexed	`[base,#imm]!`	n/a	n/a
Post-indexed	`[base],#imm`	n/a	n/a
PC-relative (literal) load	`label`	n/a	n/a

- An immediate offset is encoded in various ways, depending on the type of load/store instruction:

Bits	Sign	Scaling	Write-back?	Load/Store Type
0	-	-	-	exclusive / acquire / release
7	signed	scaled	option	register pair
9	signed	unscaled	option	single register
12	unsigned	scaled	no	single register

- Where an immediate offset is scaled, it is encoded as a multiple of the data access size (except PC-relative loads, where it is always a word multiple). The assembler always accepts a **byte** offset, which is converted to the scaled offset for encoding, and a disassembler decodes the scaled offset encoding and displays it as a byte offset. The range of byte offsets supported therefore varies according to the type of load/store instruction and the data access size.
- The "post-indexed" forms mean that the memory address is the `base` register value, then `base` plus offset is written back to the `base` register.
- The "pre-indexed" forms mean that the memory address is the `base` register value plus offset, then the computed address is written back to the `base` register.
- A "register offset" means that the memory address is the `base` register value plus the value of 64-bit index register `Xm` optionally scaled by the access size (in bytes), i.e. shifted left by `log2(size)`.
- An "extended register offset" means that the memory address is the `base` register value plus the value of 32-bit index register `Wm`, sign or zero extended to 64 bits, then optionally scaled by the access size.
- An assembler should accept `Xm` as an extended index register, though `Wm` is preferred.
- The pre/post-indexed forms are not available with a register offset.
- There is no "down" option, so subtraction from the base register requires a negative signed immediate offset (two's complement) or a negative value in the index register.
- When the base register is `SP` the stack pointer is required to be quadword (16 byte, 128 bit) aligned prior to the address calculation and write-backs – misalignment will cause a stack alignment fault. The stack pointer may not be used as an index register.
- Use of the program counter (`PC`) as a base register is implicit in literal load instructions and not permitted in other load or store instructions. Literal loads do not include byte and halfword forms. See section 5 below for the definition of `label`.

5 A64 INSTRUCTION SET

.The following syntax terms are used frequently throughout the A64 instruction set description. See also the syntax notation described in section 2 above.

Xn   Unless otherwise indicated a general register operand Xn or Wn interprets register 31 as the *zero register,* represented by the names XZR or WZR respectively.

Xn|SP   A general register operand of the form Xn|SP or Wn|WSP interprets register 31 as the *stack pointer,* represented by the names SP or WSP respectively.

cond   A standard ARM condition EQ, NE, CS|HS, CC|LO, MI, PL, VS, VC, HI, LS, GE, LT, GT, LE, AL or NV with the same meanings as in AArch32. Note that although AL and NV represent different encodings, as in AArch32 they are both interpreted as the "always true" condition. Unless stated AArch64 instructions do not set or use the condition flags, but those that do set all of the condition flags. If used in a pseudo-code expression this symbol represents a Boolean whose value is the truth of the specified condition test.

invert(cond)

            The inverse of cond, for example the inverse of GT is LE.

uimm*n*   An *n*-bit unsigned (positive) immediate value.

simm*n*   An *n*-bit two's complement signed immediate value (where *n* includes the sign bit).

label   Represents a pc-relative reference from an instruction to a target code or data location. The precise syntax is likely to be specific to individual toolchains, but the preferred form is "*pcsym*" or "*pcsym±offs*", where *pcsym* is:

            a. The preferred architectural notation which is (at the choice of the disassembler) the character '.' or string "{pc}" representing the referencing instruction's address or offset.

            b. For a programmers' view where the instruction's address in memory or offset within a relocatable image is known and a list of symbols is available, then the symbol name whose value is nearest to, and preferably less than or equal to the target location's address or offset.

            c. For a programmers' view where the instruction's address or offset is known but a list of symbols is not available, then the target address or offset as a hexadecimal constant.

            And where in all cases "*±offs*" gives the byte offset from *pcsym* to the target location's address or offset, which may be omitted if the offset is zero.

addr   Represents an addressing mode that is some subset (documented for each class of instruction) of the addressing modes in section 4.5 above.

lshift   Represents an optional shift operator performed on the final source operand of a logical instruction, taking chosen from LSL, LSR, ASR, or ROR, followed by a constant shift amount #imm in the range 0 to *regwidth*-1. If omitted the default is "LSL #0".

ashift   Represents an optional shift operator to be performed on the final source operand of an arithmetic instruction chosen from LSL, LSR, or ASR, followed by a constant shift amount #imm in the range 0 to *regwidth*-1. If omitted the default is "LSL #0".

5.1 Control Flow

5.1.1 Conditional Branch

Unless stated, conditional branches have a branch offset range of ±1MiB from the program counter.

`B.cond label`

    Branch: conditionally jumps to program-relative label if cond is true.

`CBNZ Wn, label`

    Compare and Branch Not Zero: conditionally jumps to program-relative `label` if `Wn` is not equal to zero.

`CBNZ Xn, label`

    Compare and Branch Not Zero (extended): conditionally jumps to `label` if `Xn` is not equal to zero.

`CBZ Wn, label`

    Compare and Branch Zero: conditionally jumps to `label` if `Wn` is equal to zero.

`CBZ Xn, label`

    Compare and Branch Zero (extended): conditionally jumps to `label` if `Xn` is equal to zero.

`TBNZ Xn|Wn, #uimm6, label`

    Test and Branch Not Zero: conditionally jumps to `label` if bit number `uimm6` in register `Xn` is not zero. The bit number implies the width of the register, which may be written and should be disassembled as `Wn` if `uimm` is less than 32. Limited to a branch offset range of ±32KiB.

`TBZ Xn|Wn, #uimm6, label`

    Test and Branch Zero: conditionally jumps to `label` if bit number `uimm6` in register `Xn` is zero. The bit number implies the width of the register, which may be written and should be disassembled as `Wn` if `uimm6` is less than 32. Limited to a branch offset range of ±32KiB.

5.1.2 Unconditional Branch (immediate)

Unconditional branches support an immediate branch offset range of ±128MiB.

`B label`

    Branch: unconditionally jumps to pc-relative `label`.

`BL label`

    Branch and Link: unconditionally jumps to pc-relative `label`, writing the address of the next sequential instruction to register `X30`.

5.1.3 Unconditional Branch (register)

`BLR Xm`

    Branch and Link Register: unconditionally jumps to address in `Xm`, writing the address of the next sequential instruction to register `X30`.

`BR Xm`

    Branch Register: jumps to address in `Xm`, with a hint to the CPU that this <u>is not</u> a subroutine return.

`RET {Xm}`

    Return: jumps to register `Xm`, with a hint to the CPU that this <u>is</u> a subroutine return. An assembler shall default to register `X30` if `Xm` is omitted.

5.2 Memory Access

Aside from exclusive and explicitly ordered loads and stores, addresses may have arbitrary alignment unless strict alignment checking is enabled (SCTLR.A==1). However if SP is used as the base register then the value of the stack pointer prior to adding any offset must be quadword (16 byte) aligned, or else a stack alignment exception will be generated.

A memory read or write generated by the load or store of a single general-purpose register aligned to the size of the transfer is atomic. Memory reads or writes generated by the non-exclusive load or store of a pair of general-purpose registers aligned to the size of the register are treated as two atomic accesses, one for each register. In all other cases, unless otherwise stated, there are no atomicity guarantees.

5.2.1 Load-Store Single Register

The most general forms of load-store support a variety of addressing modes, consisting of base register Xn or SP, plus one of:

- Scaled, 12-bit, unsigned immediate offset, <u>without</u> pre- and post-index options.
- Unscaled, 9-bit, signed immediate offset with pre- or post-index writeback.
- Scaled or unscaled 64-bit register offset.
- Scaled or unscaled 32-bit extended register offset.

If a Load instruction specifies writeback and the register being loaded is also the base register, then one of the following behaviours can occur:

- The instruction is UNALLOCATED

- The instruction is treated as a NOP

- The instruction performs the load using the specified addressing mode and the base register becomes UNKNOWN. In addition, if an exception occurs during such an instruction, the base address might be corrupted such that the instruction cannot be repeated.

If a Store instruction performs a writeback and the register being stored is also the base register, then one of the following behaviours can occur:

- The instruction is UNALLOCATED

- The instruction is treated as a NOP

- The instruction performs the stores of the register specified using the specified addressing mode but the value stored is UNKNOWN

LDR Wt, addr

 Load Register: loads a word from memory addressed by addr to Wt.

LDR Xt, addr

 Load Register (extended): loads a doubleword from memory addressed by addr to Xt.

LDRB Wt, addr

 Load Byte: loads a byte from memory addressed by addr, then zero-extends it to Wt.

LDRSB Wt, addr

 Load Signed Byte: loads a byte from memory addressed by addr, then sign-extends it into Wt.

LDRSB Xt, addr

 Load Signed Byte (extended): loads a byte from memory addressed by addr, then sign-extends it into Xt.

LDRH Wt, addr

 Load Halfword: loads a halfword from memory addressed by addr, then zero-extends it into Wt.

LDRSH Wt, addr

Load Signed Halfword: loads a halfword from memory addressed by addr, then sign-extends it into Wt.

LDRSH Xt, addr

Load Signed Halfword (extended): loads a halfword from memory addressed by addr, then sign-extends it into Xt.

LDRSW Xt, addr

Load Signed Word (extended): loads a word from memory addressed by addr, then sign-extends it into Xt.

STR Wt, addr

Store Register: stores word from Wt to memory addressed by addr.

STR Xt, addr

Store Register (extended): stores doubleword from Xt to memory addressed by addr.

STRB Wt, addr

Store Byte: stores byte from Wt to memory addressed by addr.

STRH Wt, addr

Store Halfword: stores halfword from Wt to memory addressed by addr.

5.2.2 Load-Store Single Register (unscaled offset)

The load-store single register (unscaled offset) instructions support an addressing mode of base register Xn or SP, plus:

- Unscaled, 9-bit, signed immediate offset, <u>without</u> pre- and post-index options

These instructions use unique mnemonics to distinguish them from normal load-store instructions due to the overlap of functionality with the scaled 12-bit unsigned immediate offset addressing mode when the offset is positive and naturally aligned.

A programmer-friendly assembler could generate these instructions in response to the standard LDR/STR mnemonics when the immediate offset is unambiguous, i.e. when it is negative or unaligned. Similarly a disassembler could display these instructions using the standard LDR/STR mnemonics when the encoded immediate is negative or unaligned. However this behaviour is not required by the architectural assembly language.

LDUR Wt, [base,#simm9]

Load (Unscaled) Register: loads a word from memory addressed by base+simm9 to Wt.

LDUR Xt, [base,#simm9]

Load (Unscaled) Register (extended): loads a doubleword from memory addressed by base+simm9 to Xt.

LDURB Wt, [base,#simm9]

Load (Unscaled) Byte: loads a byte from memory addressed by base+simm9, then zero-extends it into Wt.

LDURSB Wt, [base,#simm9]

Load (Unscaled) Signed Byte: loads a byte from memory addressed by base+simm9, then sign-extends it into Wt.

LDURSB Xt, [base,#simm9]

Load (Unscaled) Signed Byte (extended): loads a byte from memory addressed by base+simm9, then sign-extends it into Xt.

```
LDURH Wt, [base,#simm9]
```

Load (Unscaled) Halfword: loads a halfword from memory addressed by base+simm9, then zero-extends it into Wt.

```
LDURSH Wt, [base,#simm9]
```

Load (Unscaled) Signed Halfword: loads a halfword from memory addressed by base+simm9, then sign-extends it into Wt.

```
LDURSH Xt, [base,#simm9]
```

Load (Unscaled) Signed Halfword (extended): loads a halfword from memory addressed by base+simm9, then sign-extends it into Xt.

```
LDURSW Xt, [base,#simm9]
```

Load (Unscaled) Signed Word (extended): loads a word from memory addressed by base+simm9, then sign-extends it into Xt.

```
STUR Wt, [base,#simm9]
```

Store (Unscaled) Register: stores word from Wt to memory addressed by base+simm9.

```
STUR Xt, [base,#simm9]
```

Store (Unscaled) Register (extended): stores doubleword from Xt to memory addressed by base+simm9.

```
STURB Wt, [base,#simm9]
```

Store (Unscaled) Byte: stores byte from Wt to memory addressed by base+simm9.

```
STURH Wt, [base,#simm9]
```

Store (Unscaled) Halfword: stores halfword from Wt to memory addressed by base+simm9.

5.2.3 Load Single Register (pc-relative, literal load)

The pc-relative address from which to load is encoded as a 19-bit signed **word** offset which is shifted left by 2 and added to the program counter, giving access to any word-aligned location within ±1MiB of the PC.

As a convenience assemblers will typically permit the notation "=value" in conjunction with the pc-relative literal load instructions to automatically place an immediate value or symbolic address in a nearby literal pool and generate a hidden label which references it. But that syntax is not architectural and will never appear in a disassembly. A64 has other instructions to construct immediate values (section 5.3.3) and addresses (section 5.3.4) in a register which may be preferable to loading them from a literal pool.

```
LDR Wt, label | =value
```

Load Literal Register (32-bit): loads a word from memory addressed by label to Wt.

```
LDR Xt, label | =value
```

Load Literal Register (64-bit): loads a doubleword from memory addressed by label to Xt.

```
LDRSW Xt, label | =value
```

Load Literal Signed Word (extended): loads a word from memory addressed by label, then sign-extends it into Xt.

5.2.4 Load-Store Pair

The load-store pair instructions support an addressing mode consisting of base register Xn or SP, plus:

- Scaled 7-bit signed immediate offset, with pre- and post-index writeback options

If a Load Pair instruction specifies the same register for the two registers that are being loaded, then one of the following behaviours can occur:

- The instruction is UNALLOCATED

- The instruction is treated as a NOP

- The instruction performs all of the loads using the specified addressing mode and the register being loaded takes an UNKNOWN value

If a Load Pair instruction specifies writeback and one of the registers being loaded is also the base register, then one of the following behaviours can occur:

- The instruction is UNALLOCATED

- The instruction is treated as a NOP

- The instruction performs all of the loads using the specified addressing mode and the base register becomes UNKNOWN. In addition, if an exception occurs during such an instruction, the base address might be corrupted such that the instruction cannot be repeated.

If a Store Pair instruction performs a writeback and one of the registers being stored is also the base register, then one of the following behaviours can occur:

- The instruction is UNALLOCATED

- The instruction is treated as a NOP

- The instruction performs all of the stores of the registers specified using the specified addressing mode but the value stored for the base register is UNKNOWN

```
LDP Wt1, Wt2, addr
```
> Load Pair Registers: loads two words from memory addressed by addr to Wt1 and Wt2.

```
LDP Xt1, Xt2, addr
```
> Load Pair Registers (extended): loads two doublewords from memory addressed by addr to Xt1 and Xt2.

```
LDPSW Xt1, Xt2, addr
```
> Load Pair Signed Words (extended) loads two words from memory addressed by addr, then sign-extends them into Xt1 and Xt2.

```
STP Wt1, Wt2, addr
```
> Store Pair Registers: stores two words from Wt1 and Wt2 to memory addressed by addr.

```
STP Xt1, Xt2, addr
```
> Store Pair Registers (extended): stores two doublewords from Xt1 and Xt2 to memory addressed by addr.

5.2.5 Load-Store Non-temporal Pair

The LDNP and STNP non-temporal pair instructions provide a hint to the memory system that an access is "non-temporal" or "streaming" and unlikely to be accessed again in the near future so need not be retained in data caches. However depending on the memory type they may permit memory reads to be preloaded and memory writes to be gathered, in order to accelerate bulk memory transfers.

Furthermore, as a special exception to the normal memory ordering rules, where an address dependency exists between two memory reads and the second read was generated by a Load Non-temporal Pair instruction then, in the absence of any other barrier mechanism to achieve order, those memory accesses can be observed in any order by other observers within the shareability domain of the memory addresses being accessed.

The LDNP and STNP instructions support an addressing mode of base register `Xn` or `SP`, plus:

- Scaled 7-bit signed immediate offset, <u>without</u> pre- and post-index options

If a Load Non-temporal Pair instruction specifies the same register for the two registers that are being loaded, then one of the following behaviours can occur:

- The instruction is UNALLOCATED
- The instruction is treated as a NOP
- The instruction performs all of the loads using the specified addressing mode and the register being loaded takes an UNKNOWN value

`LDNP Wt1, Wt2, [base,#imm]`

Load Non-temporal Pair: loads two words from memory addressed by `base+imm` to `Wt1` and `Wt2`, with a non-temporal hint.

`LDNP Xt1, Xt2, [base,#imm]`

Load Non-temporal Pair (extended): loads two doublewords from memory addressed by `base+imm` to `Xt1` and `Xt2`, with a non-temporal hint.

`STNP Wt1, Wt2, [base,#imm]`

Store Non-temporal Pair: stores two words from `Wt1` and `Wt2` to memory addressed by `base+imm`, with a non-temporal hint.

`STNP Xt1, Xt2, [base,#imm]`

Store Non-temporal Pair (extended): stores two doublewords from `Xt1` and `Xt2` to memory addressed by `base+imm`, with a non-temporal hint.

5.2.6 Load-Store Unprivileged

The load-store unprivileged instructions may be used when the processor is at the EL1 exception level to perform a memory access as if it were at the EL0 (unprivileged) exception level. If the processor is at any other exception level, then a normal memory access for that level is performed. (The letter 'T' in these mnemonics is based on an historical ARM convention which described an access to an unprivileged virtual address as being "translated").

The load-store unprivileged instructions support an addressing mode of base register `Xn` or `SP`, plus:

- Unscaled, 9-bit, signed immediate offset, <u>without</u> pre- and post-index options

`LDTR Wt, [base,#simm9]`

Load Unprivileged Register: loads word from memory addressed by `base+simm9` to `Wt`, using EL0 privileges when at EL1.

`LDTR Xt, [base,#simm9]`

Load Unprivileged Register (extended): loads doubleword from memory addressed by `base+simm9` to `Xt`, using EL0 privileges when at EL1.

`LDTRB Wt, [base,#simm9]`

Load Unprivileged Byte: loads a byte from memory addressed by `base+simm9`, then zero-extends it into `Wt`, using EL0 privileges when at EL1.

`LDTRSB Wt, [base,#simm9]`

Load Unprivileged Signed Byte: loads a byte from memory addressed by `base+simm9`, then sign-extends it into `Wt`, using EL0 privileges when at EL1.

```
LDTRSB  Xt, [base,#simm9]
```

> Load Unprivileged Signed Byte (extended): loads a byte from memory addressed by `base+simm9`, then sign-extends it into `Xt`, using EL0 privileges when at EL1.

```
LDTRH   Wt, [base,#simm9]
```

> Load Unprivileged Halfword: loads a halfword from memory addressed by `base+simm9`, then zero-extends it into `Wt`, using EL0 privileges when at EL1.

```
LDTRSH  Wt, [base,#simm9]
```

> Load Unprivileged Signed Halfword: loads a halfword from memory addressed by `base+simm9`, then sign-extends it into `Wt`, using EL0 privileges when at EL1.

```
LDTRSH  Xt, [base,#simm9]
```

> Load Unprivileged Signed Halfword (extended): loads a halfword from memory addressed by `base+simm9`, then sign-extends it into `Xt`, using EL0 privileges when at EL1.

```
LDTRSW  Xt, [base,#simm9]
```

> Load Unprivileged Signed Word (extended): loads a word from memory addressed by `base+simm9`, then sign-extends it into `Xt`, using EL0 privileges when at EL1.

```
STTR    Wt, [base,#simm9]
```

> Store Unprivileged Register: stores a word from `Wt` to memory addressed by `base+simm9`, using EL0 privileges when at EL1.

```
STTR    Xt, [base,#simm9]
```

> Store Unprivileged Register (extended): stores a doubleword from `Xt` to memory addressed by `base+simm9`, using EL0 privileges when at EL1.

```
STTRB   Wt, [base,#simm9]
```

> Store Unprivileged Byte: stores a byte from `Wt` to memory addressed by `base+simm9`, using EL0 privileges when at EL1.

```
STTRH   Wt, [base,#simm9]
```

> Store Unprivileged Halfword: stores a halfword from `Wt` to memory addressed by `base+simm9`, using EL0 privileges when at EL1.

5.2.7 Load-Store Exclusive

The load exclusive instructions mark the accessed physical address being accessed as an exclusive access, which is checked by the store exclusive, permitting the construction of "atomic" read-modify-write operations on shared memory variables, semaphores, mutexes, spinlocks, etc.

The load-store exclusive instructions support a simple addressing mode of base register `Xn` or `SP` only. An optional offset of `#0` must be accepted by the assembler, but may be omitted on disassembly.

Natural alignment is required: an unaligned address will cause an alignment fault. A memory access generated by a load exclusive pair or store exclusive pair must be aligned to the size of the pair, and when a store exclusive pair succeeds it will cause a single-copy atomic update of the entire memory location.

```
LDXR    Wt, [base{,#0}]
```

> Load Exclusive Register: loads a word from memory addressed by `base` to `Wt`. Records the physical address as an exclusive access.

```
LDXR    Xt, [base{,#0}]
```

> Load Exclusive Register (extended): loads a doubleword from memory addressed by `base` to `Xt`. Records the physical address as an exclusive access.

```
LDXRB Wt, [base{,#0}]
```

Load Exclusive Byte: loads a byte from memory addressed by `base`, then zero-extends it into `Wt`. Records the physical address as an exclusive access.

```
LDXRH Wt, [base{,#0}]
```

Load Exclusive Halfword: loads a halfword from memory addressed by `base`, then zero-extends it into `Wt`. Records the physical address as an exclusive access.

```
LDXP Wt, Wt2, [base{,#0}]
```

Load Exclusive Pair Registers: loads two words from memory addressed by `base`, and to `Wt` and `Wt2`. Records the physical address as an exclusive access.

```
LDXP Xt, Xt2, [base{,#0}]
```

Load Exclusive Pair Registers (extended): loads two doublewords from memory addressed by `base` to `Xt` and `Xt2`. Records the physical address as an exclusive access.

```
STXR Ws, Wt, [base{,#0}]
```

Store Exclusive Register: stores word from `Wt` to memory addressed by `base`, and sets `Ws` to the returned exclusive access status.

```
STXR Ws, Xt, [base{,#0}]
```

Store Exclusive Register (extended): stores doubleword from `Xt` to memory addressed by `base`, and sets `Ws` to the returned exclusive access status.

```
STXRB Ws, Wt, [base{,#0}]
```

Store Exclusive Byte: stores byte from `Wt` to memory addressed by `base`, and sets `Ws` to the returned exclusive access status.

```
STXRH Ws, Wt, [base{,#0}]
```

Store Exclusive Halfword: stores halfword from `Xt` to memory addressed by `base`, and sets `Ws` to the returned exclusive access status.

```
STXP Ws, Wt, Wt2, [base{,#0}]
```

Store Exclusive Pair: stores two words from `Wt` and `Wt2` to memory addressed by `base`, and sets `Ws` to the returned exclusive access status.

```
STXP Ws, Xt, Xt2, [base{,#0}]
```

Store Exclusive Pair (extended): stores two doublewords from `Xt` and `Xt2` to memory addressed by `base`, and sets `Ws` to the returned exclusive access status.

5.2.8 Load-Acquire / Store-Release

A load-acquire is a load where it is guaranteed that all loads and stores appearing in program order after the load-acquire will be observed by each observer after that observer observes the load-acquire, but says nothing about loads and stores appearing before the load-acquire.

A store-release will be observed by each observer after that observer observes any loads or stores that appear in program order before the store-release, but says nothing about loads and stores appearing after the store-release.

In addition, a store-release followed by a load-acquire will be observed by each observer in program order.

A further consideration is that all store-release operations must be multi-copy atomic: that is, if one agent has seen a store-release, then all agents have seen the store-release. There are no requirements for ordinary stores to be multi-copy atomic.

The load-acquire and store-release instructions support the simple addressing mode of base register `Xn` or `SP` only. An optional offset of `#0` must be accepted by the assembler, but may be omitted on disassembly.

Natural alignment is required: an unaligned address will cause an alignment fault.

5.2.8.1 Non-exclusive

LDAR Wt, [base{,#0}]

> Load-Acquire Register: loads a word from memory addressed by base to Wt.

LDAR Xt, [base{,#0}]

> Load-Acquire Register (extended): loads a doubleword from memory addressed by base to Xt.

LDARB Wt, [base{,#0}]

> Load-Acquire Byte: loads a byte from memory addressed by base, then zero-extends it into Wt.

LDARH Wt, [base{,#0}]

> Load-Acquire Halfword: loads a halfword from memory addressed by base, then zero-extends it into Wt.

STLR Wt, [base{,#0}]

> Store-Release Register: stores a word from Wt to memory addressed by base.

STLR Xt, [base{,#0}]

> Store-Release Register (extended): stores a doubleword from Xt to memory addressed by base.

STLRB Wt, [base{,#0}]

> Store-Release Byte: stores a byte from Wt to memory addressed by base.

STLRH Wt, [base{,#0}]

> Store-Release Halfword: stores a halfword from Wt to memory addressed by base.

5.2.8.2 Exclusive

LDAXR Wt, [base{,#0}]

> Load-Acquire Exclusive Register: loads word from memory addressed by base to Wt. Records the physical address as an exclusive access.

LDAXR Xt, [base{,#0}]

> Load-Acquire Exclusive Register (extended): loads doubleword from memory addressed by base to Xt. Records the physical address as an exclusive access.

LDAXRB Wt, [base{,#0}]

> Load-Acquire Exclusive Byte: loads byte from memory addressed by base, then zero-extends it into Wt. Records the physical address as an exclusive access.

LDAXRH Wt, [base{,#0}]

> Load-Acquire Exclusive Halfword: loads halfword from memory addressed by base, then zero-extends it into Wt. Records the physical address as an exclusive access.

LDAXP Wt, Wt2, [base{,#0}]

> Load-Acquire Exclusive Pair Registers: loads two words from memory addressed by base to Wt and Wt2. Records the physical address as an exclusive access.

LDAXP Xt, Xt2, [base{,#0}]

> Load-Acquire Exclusive Pair Registers (extended): loads two doublewords from memory addressed by base to Xt and Xt2. Records the physical address as an exclusive access.

STLXR Ws, Wt, [base{,#0}]

> Store-Release Exclusive Register: stores word from Wt to memory addressed by base, and sets Ws to the returned exclusive access status.

```
STLXR Ws, Xt, [base{,#0}]
```

>   Store-Release Exclusive Register (extended): stores doubleword from Xt to memory addressed by base, and sets Ws to the returned exclusive access status.

```
STLXRB Ws, Wt, [base{,#0}]
```

>   Store-Release Exclusive Byte: stores byte from Wt to memory addressed by base, and sets Ws to the returned exclusive access status.

```
STLXRH Ws, Xt|Wt, [base{,#0}]
```

>   Store-Release Exclusive Halfword: stores the halfword from Wt to memory addressed by base, and sets Ws to the returned exclusive access status.

```
STLXP Ws, Wt, Wt2, [base{,#0}]
```

>   Store-Release Exclusive Pair: stores two words from Wt and Wt2 to memory addressed by base, and sets Ws to the returned exclusive access status.

```
STLXP Ws, Xt, Xt2, [base{,#0}]
```

>   Store-Release Exclusive Pair (extended): stores two doublewords from Xt and Xt2 to memory addressed by base, and sets Ws to the returned exclusive access status.

5.2.9 Prefetch Memory

The prefetch memory instructions signal the memory system that memory accesses from a specified address are likely in the near future. The memory system can respond by taking actions that are expected to speed up the memory accesses when they do occur, such as pre-loading the specified address into one or more caches. Since these are only hints, it is valid for the CPU to treat any or all prefetch instructions as a no-op.

The prefetch instructions support a wide range of addressing modes, consisting of a base register Xn or SP, plus one of:

- Scaled, 12-bit, unsigned immediate offset, <u>without</u> pre- and post-index options.
- Unscaled, 9-bit, signed immediate offset, <u>without</u> pre- and post-index options.
- Scaled or unscaled 64-bit register offset.
- Scaled or unscaled 32-bit extended register offset.

Additionally:

- A PC-relative address or label, within ±1MB of the current PC.
- Where an offset is scaled it is as if for an access size of 8 bytes.

```
PRFM <prfop>, addr|label
```

>   Prefetch Memory, using the <prfop> hint, where <prfop> is one of:
>   ```
>       PLDL1KEEP, PLDL1STRM, PLDL2KEEP, PLDL2STRM, PLDL3KEEP, PLDL3STRM
>       PSTL1KEEP, PSTL1STRM, PSTL2KEEP, PSTL2STRM, PSTL3KEEP, PSTL3STRM
>   ```
>
>   ```
>   <prfop>     ::= <type><target><policy> | #uimm5
>   <type>      ::= "PLD" (prefetch for load) | "PST" (prefetch for store)
>   <target>    ::= "L1" (L1 cache) | "L2" (L2 cache) | "L3" (L3 cache)
>   <policy>    ::= "KEEP" (retained or temporal prefetch, i.e. allocate in cache normally)
>                 | "STRM" (streaming or non-temporal prefetch, i.e. memory used only once)
>   #uimm5      ::= represents the unallocated hint encodings as a 5-bit immediate
>   ```

```
PRFUM <prfop>, addr
```
> Prefetch Memory (unscaled offset), explicitly uses the 9-bit, signed, unscaled immediate offset addressing mode, as described in section 5.2.2

5.3 Data Processing (immediate)

The following instruction groups are supported:

- Arithmetic (immediate)
- Logical (immediate)
- Move (immediate)
- Bitfield (operations)
- Shift (immediate)
- Sign/zero extend

5.3.1 Arithmetic (immediate)

These instructions accept an *arithmetic immediate* shown as `aimm`, which is encoded as a 12-bit unsigned immediate shifted left by 0 or 12 bits. In the assembly language this may be written as:

```
#uimm12, LSL #sh
```
> A 12-bit unsigned immediate, explicitly shifted left by 0 or 12.

```
#uimm24
```
> A 24-bit unsigned immediate. An assembler shall determine the appropriate value of `uimm12` with lowest possible shift of 0 or 12 which generates the requested value; if the value contains non-zero bits in bits<23:12> and in bits<11:0> then an error shall result.

```
#nimm25
```
> A "programmer-friendly" assembler may accept a negative immediate between $-(2^{24}-1)$ and -1 inclusive, causing it to convert a requested `ADD` operation to a `SUB`, or *vice versa*, and then encode the absolute value of the immediate as for `uimm24`. However this behaviour is not required by the architectural assembly language.

A disassembler should normally output the arithmetic immediate using the `uimm24` form, unless the encoded shift amount is not the lowest possible shift that could have been used (for example `#0,LSL #12` could not be output using the `uimm24` form).

The arithmetic instructions which do not set condition flags may read and/or write the current stack pointer, for example to adjust the stack pointer in a function prologue or epilogue; the flag setting instructions can read the stack pointer, but not write it.

```
ADD Wd|WSP, Wn|WSP, #aimm
```
> Add (immediate): `Wd|WSP = Wn|WSP + aimm`.

```
ADD Xd|SP, Xn|SP, #aimm
```
> Add (extended immediate): `Xd|SP = Xn|SP + aimm`.

```
ADDS Wd, Wn|WSP, #aimm
```
> Add and set flags (immediate): `Wd = Wn|WSP + aimm`, setting the condition flags.

```
ADDS Xd, Xn|SP, #aimm
```

Add and set flags (extended immediate): `Xd = Xn|SP + aimm`, setting the condition flags.

```
SUB Wd|WSP, Wn|WSP, #aimm
```

Subtract (immediate): `Wd|WSP = Wn|WSP - aimm`.

```
SUB Xd|SP, Xn|SP, #aimm
```

Subtract (extended immediate): `Xd|SP = Xn|SP - aimm`.

```
SUBS Wd, Wn|WSP, #aimm
```

Subtract and set flags (immediate): `Wd = Wn|WSP - aimm`, setting the condition flags.

```
SUBS Xd, Xn|SP, #aimm
```

Subtract and set flags (extended immediate): `Xd = Xn|SP - aimm`, setting the condition flags.

```
CMP Wn|WSP, #aimm
```

Compare (immediate): alias for `SUBS WZR,Wn|WSP,#aimm`.

```
CMP Xn|SP, #aimm
```

Compare (extended immediate): alias for `SUBS XZR,Xn|SP,#aimm`.

```
CMN Wn|WSP, #aimm
```

Compare negative (immediate): alias for `ADDS WZR,Wn|WSP,#aimm`.

```
CMN Xn|SP, #aimm
```

Compare negative (extended immediate): alias for `ADDS XZR,Xn|SP,#aimm`.

```
MOV Wd|WSP, Wn|WSP
```

Move (register): alias for `ADD Wd|WSP,Wn|WSP,#0`, but only when one or other of the registers is `WSP`. In other cases the `ORR Wd,WZR,Wn` instruction is used.

```
MOV Xd|SP, Xn|SP
```

Move (extended register): alias for `ADD Xd|SP,Xn|SP,#0`, but only when one or other of the registers is `SP`. In other cases the `ORR Xd,XZR,Xn` instruction is used.

5.3.2 Logical (immediate)

The logical immediate instructions accept a *bitmask immediate* `bimm32` or `bimm64`. Such an immediate consists EITHER of a single consecutive sequence with at least one non-zero bit, and at least one zero bit, within an element of 2, 4, 8, 16, 32 or 64 bits; the element then being replicated across the register width, or the bitwise inverse of such a value. The immediate values of all-zero and all-ones may not be encoded as a bitmask immediate, so an assembler must either generate an error for a logical instruction with such an immediate, or a programmer-friendly assembler may transform it into some other instruction which achieves the intended result.

The logical (immediate) instructions may write to the current stack pointer, for example to align the stack pointer in a function prologue.

Note: Apart from `ANDS`, logical immediate instructions do not set the condition flags, but "interesting" results can usually directly control a `CBZ`, `CBNZ`, `TBZ` or `TBNZ` conditional branch.

```
AND Wd|WSP, Wn, #bimm32
```

Bitwise AND (immediate): `Wd|WSP = Wn AND bimm32`.

```
AND Xd|SP, Xn, #bimm64
```

Bitwise AND (extended immediate): `Xd|SP = Xn AND bimm64`.

PRD03-GENC-010197   Page 33 of 112

`ANDS Wd, Wn, #bimm32`

> Bitwise AND and Set Flags (immediate): `Wd = Wn AND bimm32`, setting N & Z condition flags based on the result and clearing the C & V flags.

`ANDS Xd, Xn, #bimm64`

> Bitwise AND and Set Flags (extended immediate): `Xd = Xn AND bimm64`, setting N & Z condition flags based on the result and clearing the C & V flags.

`EOR Wd|WSP, Wn, #bimm32`

> Bitwise exclusive OR (immediate): `Wd|WSP = Wn EOR bimm32`.

`EOR Xd|SP, Xn, #bimm64`

> Bitwise exclusive OR (extended immediate): `Xd|SP = Xn EOR bimm64`.

`ORR Wd|WSP, Wn, #bimm32`

> Bitwise inclusive OR (immediate): `Wd|WSP = Wn OR bimm32`.

`ORR Xd|SP, Xn, #bimm64`

> Bitwise inclusive OR (extended immediate): `Xd|SP = Xn OR bimm64`.

`MOVI Wd, #bimm32`

> Move bitmask (immediate): alias for `ORR Wd,WZR,#bimm32`, but may disassemble as `MOV`, see below.

`MOVI Xd, #bimm64`

> Move bitmask (extended immediate): alias for `ORR Xd,XZR,#bimm64`, but may disassemble as `MOV`, see below.

`TST Wn, #bimm32`

> Bitwise test (immediate): alias for `ANDS WZR,Wn,#bimm32`.

`TST Xn, #bimm64`

> Bitwise test (extended immediate): alias for `ANDS XZR,Xn,#bimm64`

5.3.3 Move (wide immediate)

These instructions insert a 16-bit immediate (or inverted immediate) into a 16-bit aligned position in the destination register, with the value of the other destination register bits depending on the variant used. The shift amount `pos` may be any multiple of 16 less than the register size. Omitting "`LSL #pos`" implies a shift of 0.

`MOVZ Wt, #uimm16{, LSL #pos}`

> Move with Zero (immediate): `Wt = LSL(uimm16, pos)`.
> Usually disassembled as `MOV`, see below.

`MOVZ Xt, #uimm16{, LSL #pos}`

> Move with Zero (extended immediate): `Xt = LSL(uimm16, pos)`.
> Usually disassembled as `MOV`, see below.

`MOVN Wt, #uimm16{, LSL #pos}`

> Move with NOT (immediate): `Wt = NOT(LSL(uimm16, pos))`.
> Usually disassembled as `MOV`, see below.

`MOVN Xt, #uimm16{, LSL #pos}`

> Move with NOT (extended immediate): `Xt = NOT(LSL(uimm16, pos))`.
> Usually disassembled as `MOV`, see below.

`MOVK Wt, #uimm16{, LSL #pos}`

> Move with Keep (immediate): `Wt<pos+15:pos> = uimm16`.

```
MOVK Xt, #uimm16{, LSL #pos}
```

> Move with Keep (extended immediate): `Xt<pos+15:pos> = uimm16`.

5.3.3.1 Move (immediate)

```
MOV Wd, #simm32
```

> A synthetic assembler instruction which generates a single `MOVZ`, `MOVN` or `MOVI` instruction that loads a 32-bit immediate value into register `Wd`. An assembler error shall result if the immediate cannot be created by a single one of these instructions. If there is a choice, then to ensure reversability an assembler must prefer a `MOVZ` to `MOVN`, and `MOVZ` or `MOVN` to `MOVI`. A disassembler may output `MOVI`, `MOVZ` and `MOVN` as a `MOV` mnemonic, except when `MOVI` has an immediate that could be generated by a `MOVZ` or `MOVN` instruction, or where a `MOVN` has an immediate that could be encoded by `MOVZ`, or where `MOVZ/MOVN #0` have a shift amount other than `LSL #0`, in which case the machine-instruction mnemonic must be used.

```
MOV Xd, #simm64
```

> As `MOV` but for loading a 64-bit immediate into register `Xd`.

5.3.4 Address Generation

```
ADRP Xd, label
```

> Address of Page: sign extends a 21-bit offset, shifts it left by 12 and adds it to the value of the PC with its bottom 12 bits cleared, writing the result to register `Xd`. This computes the base address of the 4KiB aligned memory region containing label, and is designed to be used in conjunction with a load, store or `ADD` instruction which supplies the bottom 12 bits of the label's address. This permits position-independent addressing of any location within ±4GiB of the PC using two instructions, providing that dynamic relocation is done with a minimum granularity of 4KiB (i.e. the bottom 12 bits of the label's address are unaffected by the relocation). The term "page" is short-hand for the 4KiB relocation granule, and is not necessarily related to the virtual memory page size.

```
ADR Xd, label
```

> Address: adds a 21-bit signed byte offset to the program counter, writing the result to register `Xd`. Used to compute the effective address of any location within ±1MiB of the PC.

5.3.5 Bitfield Operations

```
BFM Wd, Wn, #r, #s
```

> Bitfield Move: if `s>=r` then `Wd<s-r:0> = Wn<s:r>`, else `Wd<32+s-r,32-r> = Wn<s:0>`. Leaves other bits in `Wd` unchanged.

```
BFM Xd, Xn, #r, #s
```

> Bitfield Move: if `s>=r` then `Xd<s-r:0> = Xn<s:r>`, else `Xd<64+s-r,64-r> = Xn<s:0>`. Leaves other bits in `Xd` unchanged.

```
SBFM Wd, Wn, #r, #s
```

> Signed Bitfield Move: if `s>=r` then `Wd<s-r:0> = Wn<s:r>`, else `Wd<32+s-r,32-r> = Wn<s:0>`. Sets bits to the left of the destination bitfield to copies of its leftmost bit, and bits to the right to zero.

```
SBFM Xd, Xn, #r, #s
```

> Signed Bitfield Move: if `s>=r` then `Xd<s-r:0> = Xn<s:r>`, else `Xd<64+s-r,64-r> = Xn<s:0>`. Sets bits to the left of the destination bitfield to copies of its leftmost bit, and bits to the right to zero.

```
UBFM Wd, Wn, #r, #s
```

> Unsigned Bitfield Move: if `s>=r` then `Wd<s-r:0> = Wn<s:r>`, else `Wd<32+s-r,32-r> = Wn<s:0>`. Sets bits to the left and right of the destination bitfield to zero.

```
UBFM Xd, Xn, #r, #s
```

> Unsigned Bitfield Move: if s>=r then Xd<s-r:0> = Xn<s:r>, else Xd<32+s-r,32-r> = Xn<s:0>.
> Sets bits to the left and right of the destination bitfield to zero.

The following aliases provide more familiar bitfield insert and extract mnemonics, with conventional bitfield `lsb` and `width` operands, which must satisfy the constraints `lsb >= 0 && width >= 1 && lsb+width <= reg.size`

```
BFI Wd, Wn, #lsb, #width
```

> Bitfield Insert: alias for `BFM Wd,Wn,#((32-lsb)&31),#(width-1)`.
> Preferred for disassembly when s < r.

```
BFI Xd, Xn, #lsb, #width
```

> Bitfield Insert (extended): alias for `BFM Xd,Xn,#((64-lsb)&63),#(width-1)`.
> Preferred for disassembly when s < r.

```
BFXIL Wd, Wn, #lsb, #width
```

> Bitfield Extract and Insert Low: alias for `BFM Wd,Wn,#lsb,#(lsb+width-1)`.
> Preferred for disassembly when s >= r.

```
BFXIL Xd, Xn, #lsb, #width
```

> Bitfield Extract and Insert Low (extended): alias for `BFM Xd,Xn,#lsb,#(lsb+width-1)`.
> Preferred for disassembly when s >= r.

```
SBFIZ Wd, Wn, #lsb, #width
```

> Signed Bitfield Insert in Zero: alias for) `SBFM Wd,Wn,#((32-lsb)&31),#(width-1)`.
> Preferred for disassembly when s < r.

```
SBFIZ Xd, Xn, #lsb, #width
```

> Signed Bitfield Insert in Zero (extended): alias for `SBFM Xd,Xn,#((64-lsb)&63),#(width-1)`.
> Preferred for disassembly when s < r.

```
SBFX Wd, Wn, #lsb, #width
```

> Signed Bitfield Extract: alias for `SBFM Wd,Wn,#lsb,#(lsb+width-1)`.
> Preferred for disassembly when s >= r.

```
SBFX Xd, Xn, #lsb, #width
```

> Signed Bitfield Extract (extended): alias for `SBFM Xd,Xn,#lsb,#(lsb+width-1)`.
> Preferred for disassembly when s >= r.

```
UBFIZ Wd, Wn, #lsb, #width
```

> Unsigned Bitfield Insert in Zero: alias for `UBFM Wd,Wn,#((32-lsb)&31),#(width-1)`.
> Preferred for disassembly when s < r.

```
UBFIZ Xd, Xn, #lsb, #width
```

> Unsigned Bitfield Insert in Zero (extended): alias for `UBFM Xd,Xn,#((64-lsb)&63),#(width-1)`.
> Preferred for disassembly when s < r.

```
UBFX Wd, Wn, #lsb, #width
```

> Unsigned Bitfield Extract: alias for `UBFM Wd,Wn,#lsb,#(lsb+width-1)`.
> Preferred for disassembly when s >= r.

```
UBFX Xd, Xn, #lsb, #width
```

> Unsigned Bitfield Extract (extended): alias for `UBFM Xd,Xn,#lsb,#(lsb+width-1)`.
> Preferred for disassembly when s >= r.

5.3.6 Extract (immediate)

`EXTR Wd, Wn, Wm, #lsb`

> Extract: `Wd = Wn:Wm<lsb+31,lsb>`. The bit position `lsb` must be in the range 0 to `31`.

`EXTR Xd, Xn, Xm, #lsb`

> Extract (extended): `Xd = Xn:Xm<lsb+63,lsb>`. The bit position `lsb` must be in the range 0 to `63`.

5.3.7 Shift (immediate)

All immediate shifts and rotates are aliases, implemented using the Bitfield or Extract instructions. In all cases the immediate shift amount `uimm` must be in the range 0 to (reg.size - 1).

`ASR Wd, Wn, #uimm`

> Arithmetic Shift Right (immediate): alias for `SBFM Wd,Wn,#uimm,#31`.

`ASR Xd, Xn, #uimm`

> Arithmetic Shift Right (extended immediate): alias for `SBFM Xd,Xn,#uimm,#63`.

`LSL Wd, Wn, #uimm`

> Logical Shift Left (immediate): alias for `UBFM Wd,Wn,#((32-uimm)&31),#(31-uimm)`.

`LSL Xd, Xn, #uimm`

> Logical Shift Left (extended immediate): alias for `UBFM Xd,Xn,#((64-uimm)&63),#(63-uimm)`

`LSR Wd, Wn, #uimm`

> Logical Shift Left (immediate): alias for `UBFM Wd,Wn,#uimm,#31`.

`LSR Xd, Xn, #uimm`

> Logical Shift Left (extended immediate): alias for `UBFM Xd,Xn,#uimm,#31`.

`ROR Wd, Wm, #uimm`

> Rotate Right (immediate): alias for `EXTR Wd,Wm,Wm,#uimm`.

`ROR Xd, Xm, #uimm`

> Rotate Right (extended immediate): alias for `EXTR Xd,Xm,Xm,#uimm`.

5.3.8 Sign/Zero Extend

`SXT[BH] Wd, Wn`

> Signed Extend Byte|Halfword: alias for `SBFM Wd,Wn,#0,#7|15`.

`SXT[BHW] Xd, Wn`

> Signed Extend Byte|Halfword|Word (extended): alias for `SBFM Xd,Xn,#0,#7|15|31`.

`UXT[BH] Wd, Wn`

> Unsigned Extend Byte|Halfword: alias for `UBFM Wd,Wn,#0,#7|15`.

`UXT[BHW] Xd, Wn`

> Unsigned Extend Byte|Halfword|Word (extended): alias for `UBFM Xd,Xn,#0,#7|15|31`.

5.4 Data Processing (register)

The following instruction groups are supported:

- Arithmetic (shifted register)
- Arithmetic (extending register)

- Logical (shifted register)
- Arithmetic (unshifted register)
- Shift (register)
- Bitwise operations

5.4.1 Arithmetic (shifted register)

The shifted register instructions apply an optional shift to the final source operand value before performing the arithmetic operation. The register size of the instruction controls where new bits are fed in to the intermediate result on a right shift or rotate (i.e. bit 63 or 31).

The shift operators LSL, ASR and LSR accept an immediate shift amount in the range 0 to reg.size - 1.

Omitting the shift operator implies "LSL #0" (i.e. no shift), and "LSL #0" should not be output by a disassembler; all other shifts by zero must be output.

The register names SP and WSP may not be used with this class of instructions, instead see section 5.4.2.

ADD Wd, Wn, Wm{, ashift #imm}

      Add (register): Wd = Wn + ashift(Wm, imm).

ADD Xd, Xn, Xm{, ashift #imm}

      Add (extended register): Xd = Xn + ashift(Xm, imm).

ADDS Wd, Wn, Wm{, ashift #imm}

      Add and Set Flags (register): Wd = Wn + ashift(Wm, imm), setting condition flags.

ADDS Xd, Xn, Xm{, ashift #imm}

      Add and Set Flags (extended register): Xd = Xn + ashift(Xm, imm), setting condition flags.

SUB Wd, Wn, Wm{, ashift #imm}

      Subtract (register): Wd = Wn - ashift(Wm, imm).

SUB Xd, Xn, Xm{, ashift #imm}

      Subtract (extended register): Xd = Xn - ashift(Xm, imm).

SUBS Wd, Wn, Wm{, ashift #imm}

      Subtract and Set Flags (register): Wd = Wn - ashift(Wm, imm), setting condition flags.

SUBS Xd, Xn, Xm{, ashift #imm}

      Subtract and Set Flags (extended register): Xd = Xn - ashift(Xm, imm), setting condition flags.

CMN Wn, Wm{, ashift #imm}

      Compare Negative (register): alias for ADDS WZR, Wn, Wm{, ashift #imm}.

CMN Xn, Xm{, ashift #imm}

      Compare Negative (extended register): alias for ADDS XZR, Xn, Xm{, ashift #imm}.

CMP Wn, Wm{, ashift #imm}

      Compare (register): alias for SUBS WZR, Wn, Wm{,ashift #imm}.

CMP Xn, Xm{, ashift #imm}

      Compare (extended register): alias for SUBS XZR, Xn, Xm{, ashift #imm}.

NEG Wd, Wm{, ashift #imm}

      Negate: alias for SUB Wd, WZR, Wm{, ashift #imm}.

```
NEG Xd, Xm{, ashift #imm}
```

> Negate (extended): alias for `SUB Xd, XZR, Xm{, ashift #imm}`.

```
NEGS Wd, Wm{, ashift #imm}
```

> Negate and Set Flags: alias for `SUBS Wd, WZR, Wm{, ashift #imm}`.

```
NEGS Xd, Xm{, ashift #imm}
```

> Negate and Set Flags (extended): alias for `SUBS Xd, XZR, Xm{, ashift #imm}`.

5.4.2 Arithmetic (extending register)

The extending register instructions differ from the shifted register forms in that:

1. Non-flag setting variants permit use of the stack pointer as either or both of the destination and first source register. The flag setting variants only permit the stack pointer as the first source register.
2. They provide an optional sign or zero-extension of a portion of the second source register value, followed by an optional immediate left shift between 1 and 4 inclusive.

The "extending shift" is described by the mandatory `extend` operator SXTB, SXTH, SXTW, SXTX, UXTB, UXTH, UXTW or UXTX, which is followed by an optional left shift amount. If the shift amount is omitted then it defaults to zero, and a zero shift amount should not be output by a disassembler.

For 64-bit instruction forms the operators UXTX and SXTX (UXTX preferred) both perform a "no-op" extension of the second source register, followed by optional shift. If and only if UXTX used in combination with the register name SP in at least one operand, then the alias LSL is preferred, and in this case both the operator and shift amount may be omitted, implying "LSL #0".

Similarly for 32-bit instruction forms the operators UXTW and SXTW (UXTW preferred) both perform a "no-op" extension of the second source register, followed by optional shift. If and only if UXTW is used in combination with the register name WSP in at least one operand, then the alias LSL is preferred. In the 64-bit form of these instructions the final register operand is written as Wm for all but the (possibly omitted) UXTX/LSL and SXTX operators. For example:

```
CMP   X4, W5, SXTW
ADD   X1, X2, W3, UXTB #2
SUB   SP, SP, X1                 // SUB SP, SP, X1, UXTX #0
```

```
ADD Wd|WSP, Wn|WSP, Wm, extend {#imm}
```

> Add (register, extending): `Wd|WSP = Wn|WSP + LSL(extend(Wm),imm)`.

```
ADD Xd|SP, Xn|SP, Wm, extend {#imm}
```

> Add (extended register, extending): `Xd|SP = Xn|SP + LSL(extend(Wm),imm)`.

```
ADD Xd|SP, Xn|SP, Xm{, UXTX|LSL #imm}
```

> Add (extended register, extending): `Xd|SP = Xn|SP + LSL(Xm,imm)`.

```
ADDS Wd, Wn|WSP, Wm, extend {#imm}
```

> Add and Set Flags (register, extending): `Wd = Wn|WSP + LSL(extend(Wm),imm)`, setting the condition flags.

```
ADDS Xd, Xn|SP, Wm, extend {#imm}
```

> Add and Set Flags (extended register, extending): `Xd = Xn|SP + LSL(extend(Wm),imm)`, setting the condition flags.

```
ADDS Xd, Xn|SP, Xm{, UXTX|LSL #imm}
```
> Add and Set Flags (extended register, extending): `Xd = Xn|SP + LSL(Xm,imm)`, setting the condition flags.

```
SUB Wd|WSP, Wn|WSP, Wm, extend {#imm}
```
> Subtract (register, extending): `Wd|WSP = Wn|WSP - LSL(extend(Wm),imm)`.

```
SUB Xd|SP, Xn|SP, Wm, extend {#imm}
```
> Subtract (extended register, extending): `Xd|SP = Xn|SP - LSL(extend(Wm),imm)`.

```
SUB Xd|SP, Xn|SP, Xm{, UXTX|LSL #imm}
```
> Subtract (extended register, extending): `Xd|SP = Xn|SP - LSL(Xm,imm)`.

```
SUBS Wd, Wn|WSP, Wm, extend {#imm}
```
> Subtract and Set Flags (register, extending): `Wd = Wn|WSP - LSL(extend(Wm),imm)`, setting the condition flags.

```
SUBS Xd, Xn|SP, Wm, extend {#imm}
```
> Subtract and Set Flags (extended register, extending): `Xd = Xn|SP - LSL(extend(Wm),imm)`, setting the condition flags.

```
SUBS Xd, Xn|SP, Xm{, UXTX|LSL #imm}
```
> Subtract and Set Flags (extended register, extending): `Xd = Xn|SP - LSL(Xm,imm)`, setting the condition flags.

```
CMN Wn|WSP, Wm, extend {#imm}
```
> Compare Negative (register, extending): alias for `ADDS WZR,Wn,Wm,extend {#imm}`.

```
CMN Xn|SP, Wm, extend {#imm}
```
> Compare Negative (extended register, extending): alias for `ADDS XZR,Xn,Wm,extend {#imm}`.

```
CMN Xn|SP, Xm{, UXTX|LSL #imm}
```
> Compare Negative (extended register, extending): alias for `ADDS XZR,Xn,Xm{,UXTX|LSL #imm}`.

```
CMP Wn|WSP, Wm, extend {#imm}
```
> Compare (register, extending): alias for `SUBS WZR,Wn,Wm,extend {#imm}`.

```
CMP Xn|SP, Wm, extend {#imm}
```
> Compare (extended register, extending): alias for `SUBS XZR,Xn,Wm,extend {#imm}`.

```
CMP Xn|SP, Xm{, UXTX|LSL #imm}
```
> Compare (extended register, extending): alias for `SUBS XZR,Xn,Xm{,UXTX|LSL #imm}`.

5.4.3 Logical (shifted register)

The logical (shifted register) instructions apply an optional shift operator to their final source operand before performing the main operation. The register size of the instruction controls where new bits are fed in to the intermediate result on a right shift or rotate (i.e. bit 63 or 31).

The shift operators LSL, ASR, LSR and ROR accept an immediate shift amount in the range 0 to reg.size - 1.

Omitting the shift operator implies "LSL #0" (i.e. no shift), and an "LSL #0" should not be output by a disassembler – however all other shifts by zero must be output.

Note: Apart from ANDS and BICS the logical instructions do not set the condition flags, but "interesting" results can usually directly control a CBZ, CBNZ, TBZ or TBNZ conditional branch.

```
AND Wd, Wn, Wm{, lshift #imm}
```
  Bitwise AND (register): `Wd = Wn AND lshift(Wm, imm)`.

```
AND Xd, Xn, Xm{, lshift #imm}
```
  Bitwise AND (extended register): `Xd = Xn AND lshift(Xm, imm)`.

```
ANDS Wd, Wn, Wm{, lshift #imm}
```
  Bitwise AND and Set Flags (register): `Wd = Wn AND lshift(Wm, imm)`, setting N & Z condition flags based on the result and clearing the C & V flags.

```
ANDS Xd, Xn, Xm{, lshift #imm}
```
  Bitwise AND and Set Flags (extended register): `Xd = Xn AND lshift(Xm, imm)`, setting N & Z condition flags based on the result and clearing the C & V flags.

```
BIC Wd, Wn, Wm{, lshift #imm}
```
  Bit Clear (register): `Wd = Wn AND NOT(lshift(Wm, imm))`.

```
BIC Xd, Xn, Xm{, lshift #imm}
```
  Bit Clear (extended register): `Xd = Xn AND NOT(lshift(Xm, imm))`.

```
BICS Wd, Wn, Wm{, lshift #imm}
```
  Bit Clear and Set Flags (register): `Wd = Wn AND NOT(lshift(Wm, imm))`, setting N & Z condition flags based on the result and clearing the C & V flags.

```
BICS Xd, Xn, Xm{, lshift #imm}
```
  Bit Clear and Set Flags (extended register): `Xd = Xn AND NOT(lshift(Xm, imm))`, setting N & Z condition flags based on the result and clearing the C & V flags.

```
EON Wd, Wn, Wm{, lshift #imm}
```
  Bitwise exclusive OR NOT (register): `Wd = Wn EOR NOT(lshift(Wm, imm))`.

```
EON Xd, Xn, Xm{, lshift #imm}
```
  Bitwise exclusive OR NOT (extended register): `Xd = Xn EOR NOT(lshift(Xm, imm))`.

```
EOR Wd, Wn, Wm{, lshift #imm}
```
  Bitwise exclusive OR (register): `Wd = Wn EOR lshift(Wm, imm)`.

```
EOR Xd, Xn, Xm{, lshift #imm}
```
  Bitwise exclusive OR (extended register): `Xd = Xn EOR lshift(Xm, imm)`.

```
ORR Wd, Wn, Wm{, lshift #imm}
```
  Bitwise inclusive OR (register): `Wd = Wn OR lshift(Wm, imm)`.

```
ORR Xd, Xn, Xm{, lshift #imm}
```
  Bitwise inclusive OR (extended register): `Xd = Xn OR lshift(Xm, imm)`.

```
ORN Wd, Wn, Wm{, lshift #imm}
```
  Bitwise inclusive OR NOT (register): `Wd = Wn OR NOT(lshift(Wm, imm))`.

```
ORN Xd, Xn, Xm{, lshift #imm}
```
  Bitwise inclusive OR NOT (extended register): `Xd = Xn OR NOT(lshift(Xm, imm))`.

```
MOV Wd, Wm
```
  Move (register): alias for `ORR Wd,WZR,Wm`.

```
MOV Xd, Xm
```
  Move (extended register): alias for `ORR Xd,XZR,Xm`.

```
MVN Wd, Wm{, lshift #imm}
```
  Move NOT (register): alias for `ORN Wd,WZR,Wm{,lshift #imm}`.

```
MVN Xd, Xm{, lshift #imm}
```

> Move NOT (extended register): alias for `ORN Xd,XZR,Xm{,lshift #imm}`.

```
TST Wn, Wm{, lshift #imm}
```

> Bitwise Test (register): alias for `ANDS WZR,Wn,Wm{,lshift #imm}`.

```
TST Xn, Xm{, lshift #imm}
```

> Bitwise Test (extended register): alias for `ANDS XZR,Xn,Xm{,lshift #imm}`.

5.4.4 Variable Shift

The variable shift amount in `Wm` or `Xm` is positive, and modulo the register size. For example an extended 64-bit shift with `Xm` containing the value 65 will result in a shift by (65 MOD 64) = 1 bit. The machine instructions are as follows:

```
ASRV Wd, Wn, Wm
```

> Arithmetic Shift Right Variable: `Wd = ASR(Wn, Wm & 0x1f)`.

```
ASRV Xd, Xn, Xm
```

> Arithmetic Shift Right Variable (extended): `Xd = ASR(Xn, Xm & 0x3f)`.

```
LSLV Wd, Wn, Wm
```

> Logical Shift Left Variable: `Wd = LSL(Wn, Wm & 0x1f)`.

```
LSLV Xd, Xn, Xm
```

> Logical Shift Left Variable (extended register): `Xd = LSL(Xn, Xm & 0x3f)`.

```
LSRV Wd, Wn, Wm
```

> Logical Shift Right Variable: `Wd = LSR(Wn, Wm & 0x1f)`.

```
LSRV Xd, Xn, Xm
```

> Logical Shift Right Variable (extended): `Xd = LSR(Xn, Xm & 0x3f)`.

```
RORV Wd, Wn, Wm
```

> Rotate Right Variable: `Wd = ROR(Wn, Wm & 0x1f)`.

```
RORV Xd, Xn, Xm
```

> Rotate Right Variable (extended): `Xd = ROR(Xn, Xm & 0x3f)`.

However the "Variable Shift" machine instructions have a preferred set of "Shift (register)" aliases which match the Shift (immediate) aliases described elsewhere:

```
ASR Wd, Wn, Wm
```

> Arithmetic Shift Right (register): preferred alias for `ASRV Wd, Wn, Wm`.

```
ASR Xd, Xn, Xm
```

> Arithmetic Shift Right (extended register): preferred alias for `ASRV Xd, Xn, Xm`.

```
LSL Wd, Wn, Wm
```

> Logical Shift Left (register): preferred alias for `LSLV Wd, Wn, Wm`.

```
LSL Xd, Xn, Xm
```

> Logical Shift Left (extended register): preferred alias for `LSLV Xd, Xn, Xm`.

```
LSR Wd, Wn, Wm
```

> Logical Shift Right (register): preferred alias for `LSRV Wd, Wn, Wm`.

```
LSR Xd, Xn, Xm
```

> Logical Shift Right (extended register): preferred alias for `LSRV Xd, Xn, Xm`.

```
ROR Wd, Wn, Wm
```

>   Rotate Right (register): preferred alias for `RORV Wd, Wn, Wm`.

```
ROR Xd, Xn, Xm
```

>   Rotate Right (extended register): preferred alias for `RORV Xd, Xn, Xm`.

5.4.5 Bit Operations

```
CLS Wd, Wm
```

>   Count Leading Sign Bits: sets `Wd` to the number of consecutive bits following the topmost bit in `Wm`, that are the same as the topmost bit. The count does not include the topmost bit itself, so the result will be in the range 0 to 31 inclusive.

```
CLS Xd, Xm
```

>   Count Leading Sign Bits (extended): sets `Xd` to the number of consecutive bits following the topmost bit in `Xm`, that are the same as the topmost bit. The count does not include the topmost bit itself, so the result will be in the range 0 to 63 inclusive.

```
CLZ Wd, Wm
```

>   Count Leading Zeros: sets `Wd` to the number of binary zeros at the most significant end of `Wm`. The result will be in the range 0 to 32 inclusive.

```
CLZ Xd, Xm
```

>   Count Leading Zeros: (extended) sets `Xd` to the number of binary zeros at the most significant end of `Xm`. The result will be in the range 0 to 64 inclusive.

```
RBIT Wd, Wm
```

>   Reverse Bits: reverses the 32 bits from `Wm`, writing to `Wd`.

```
RBIT Xd, Xm
```

>   Reverse Bits (extended): reverses the 64 bits from `Xm`, writing to `Xd`.

```
REV Wd, Wm
```

>   Reverse Bytes: reverses the 4 bytes in `Wm`, writing to `Wd`.

```
REV Xd, Xm
```

>   Reverse Bytes (extended): reverses 8 bytes in `Xm`, writing to `Xd`.

```
REV16 Wd, Wm
```

>   Reverse Bytes in Halfwords: reverses the 2 bytes in each 16-bit element of `Wm`, writing to `Wd`.

```
REV16 Xd, Xm
```

>   Reverse Bytes in Halfwords (extended): reverses the 2 bytes in each 16-bit element of `Xm`, writing to `Xd`.

```
REV32 Xd, Xm
```

>   Reverse Bytes in Words (extended): reverses the 4 bytes in each 32-bit element of `Xm`, writing to `Xd`.

5.4.6 Conditional Data Processing

These instructions support two unshifted source registers, with the condition flags as a third source. Note that the instructions are not conditionally executed: the destination register is always written.

```
ADC Wd, Wn, Wm
```

>   Add with Carry: `Wd = Wn + Wm + C`.

```
ADC Xd, Xn, Xm
```

>   Add with Carry (extended): `Xd = Xn + Xm + C`.

```
ADCS Wd, Wn, Wm
```

Add with Carry and Set Flags: `Wd = Wn + Wm + C`, setting the condition flags.

```
ADCS Xd, Xn, Xm
```

Add with Carry and Set Flags (extended): `Xd = Xn + Xm + C`, setting the condition flags.

```
CSEL Wd, Wn, Wm, cond
```

Conditional Select: `Wd = if cond then Wn else Wm`.

```
CSEL Xd, Xn, Xm, cond
```

Conditional Select (extended): `Xd = if cond then Xn else Xm`.

```
CSINC Wd, Wn, Wm, cond
```

Conditional Select Increment: `Wd = if cond then Wn else Wm+1`.

```
CSINC Xd, Xn, Xm, cond
```

Conditional Select Increment (extended): `Xd = if cond then Xn else Xm+1`.

```
CSINV Wd, Wn, Wm, cond
```

Conditional Select Invert: `Wd = if cond then Wn else NOT(Wm)`.

```
CSINV Xd, Xn, Xm, cond
```

Conditional Select Invert (extended): `Xd = if cond then Xn else NOT(Xm)`.

```
CSNEG Wd, Wn, Wm, cond
```

Conditional Select Negate: `Wd = if cond then Wn else -Wm`.

```
CSNEG Xd, Xn, Xm, cond
```

Conditional Select Negate (extended): `Xd = if cond then Xn else -Xm`.

```
CSET Wd, cond
```

Conditional Set: `Wd = if cond then 1 else 0`.
Alias for `CSINC Wd,WZR,WZR,invert(cond)`.

```
CSET Xd, cond
```

Conditional Set (extended): `Xd = if cond then 1 else 0`.
Alias for `CSINC Xd,XZR,XZR,invert(cond)`

```
CSETM Wd, cond
```

Conditional Set Mask: `Wd = if cond then -1 else 0`.
Alias for `CSINV Wd,WZR,WZR,invert(cond)`.

```
CSETM Xd, cond
```

Conditional Set Mask (extended): `Xd = if cond then -1 else 0`.
Alias for `CSINV Xd,WZR,WZR,invert(cond)`.

```
CINC Wd, Wn, cond
```

Conditional Increment: `Wd = if cond then Wn+1 else Wn`.
Alias for `CSINC Wd,Wn,Wn,invert(cond)`.

```
CINC Xd, Xn, cond
```

Conditional Increment (extended): `Xd = if cond then Xn+1 else Xn`.
Alias for `CSINC Xd,Xn,Xn,invert(cond)`.

```
CINV Wd, Wn, cond
```

Conditional Invert: `Wd = if cond then NOT(Wn) else Wn`.
Alias for `CSINV Wd,Wn,Wn,invert(cond)`.

```
CINV Xd, Xn, cond
```
> Conditional Invert (extended): `Xd = if cond then NOT(Xn) else Xn`.
> Alias for `CSINV Xd,Xn,Xn,invert(cond)`.

```
CNEG Wd, Wn, cond
```
> Conditional Negate: `Wd = if cond then -Wn else Wn`.
> Alias for `CSNEG Wd,Wn,Wn,invert(cond)`.

```
CNEG Xd, Xn, cond
```
> Conditional Negate (extended): `Xd = if cond then -Xn else Xn`.
> Alias for `CSNEG Xd,Xn,Xn,invert(cond)`.

```
SBC Wd, Wn, Wm
```
> Subtract with Carry: `Wd = Wn - Wm - 1 + C`.

```
SBC Xd, Xn, Xm
```
> Subtract with Carry (extended): `Xd = Xn - Xm - 1 + C`.

```
SBCS Wd, Wn, Wm
```
> Subtract with Carry and Set Flags: `Wd = Wn - Wm - 1 + C`, setting the condition flags.

```
SBCS Xd, Xn, Xm
```
> Subtract with Carry and Set Flags (extended): `Xd = Xn - Xm - 1 + C`, setting the condition flags.

```
NGC Wd, Wm
```
> Negate with Carry: `Wd = -Wm - 1 + C`.
> Alias for `SBC Wd,WZR,Wm`.

```
NGC Xd, Xm
```
> Negate with Carry (extended): `Xd = -Xm - 1 + C`.
> Alias for `SBC Xd,XZR,Xm`.

```
NGCS Wd, Wm
```
> Negate with Carry and Set Flags: `Wd = -Wm - 1 + C`, setting the condition flags.
> Alias for `SBCS Wd,WZR,Wm`.

```
NGCS Xd, Xm
```
> Negate with Carry and Set Flags (extended): `Xd = -Xm - 1 + C`, setting the condition flags.
> Alias for `SBCS Xd,XZR,Xm`.

5.4.7 Conditional Comparison

Conditional comparison provides a "conditional select" for the `NZCV` condition flags, setting the flags to the result of a comparison if the input condition is true, or to an immediate value if the input condition is false. There are register and immediate forms, with the immediate form accepting a small 5-bit unsigned value.

The `#uimm4` operand is the bitmask used to set the `NZCV` flags when the input condition is false, with bit 3 the new value of the `N` flag, bit 2 the `Z` flag, bit 1 the `C` flag, and bit 0 the `V` flag.

```
CCMN Wn, Wm, #uimm4, cond
```
> Conditional Compare Negative (register):
> `NZCV = if cond then CMP(Wn,-Wm) else uimm4`.

```
CCMN Xn, Xm, #uimm4, cond
```
> Conditional Compare Negative (extended register):
> `NZCV = if cond then CMP(Xn,-Xm) else uimm4`.

```
CCMN Wn, #uimm5, #uimm4, cond
```

Conditional Compare Negative (immediate):
```
NZCV = if cond then CMP(Wn,-uimm5) else uimm4.
```

```
CCMN Xn, #uimm5, #uimm4, cond
```

Conditional Compare Negative (extended immediate):
```
NZCV = if cond then CMP(Xn,-uimm5) else uimm4.
```

```
CCMP Wn, Wm, #uimm4, cond
```

Conditional Compare (register):
```
NZCV = if cond then CMP(Wn,Wm) else uimm4.
```

```
CCMP Xn, Xm, #uimm4, cond
```

Conditional Compare (extended register):
```
NZCV = if cond then CMP(Xn,Xm) else uimm4.
```

```
CCMP Wn, #uimm5, #uimm4, cond
```

Conditional Compare (immediate):
```
NZCV = if cond then CMP(Wn,uimm5) else uimm4.
```

```
CCMP Xn, #uimm5, #uimm4, cond
```

Conditional Compare (extended immediate):
```
NZCV = if cond then CMP(Xn,uimm5) else uimm4.
```

5.5 Integer Multiply / Divide

5.5.1 Multiply

```
MADD Wd, Wn, Wm, Wa
```

Multiply-Add: `Wd = Wa + (Wn × Wm)`.

```
MADD Xd, Xn, Xm, Xa
```

Multiply-Add (extended): `Xd = Xa + (Xn × Xm.)`

```
MSUB Wd, Wn, Wm, Wa
```

Multiply-Subtract: `Wd = Wa − (Wn × Wm)`.

```
MSUB Xd, Xn, Xm, Xa
```

Multiply-Subtract (extended): `Xd = Xa − (Xn × Xm)`.

```
MNEG Wd, Wn, Wm
```

Multiply-Negate: `Wd = −(Wn × Wm)`.
Alias for `MSUB Wd, Wn, Wm, WZR`.

```
MNEG Xd, Xn, Xm
```

Multiply-Negate (extended): `Xd = −(Xn × Xm)`.
Alias for `MSUB Xd, Xn, Xm, XZR`.

```
MUL Wd, Wn, Wm
```

Multiply: `Wd = Wn × Wm`.
Alias for `MADD Wd, Wn, Wm, WZR`.

```
MUL Xd, Xn, Xm
```

Multiply (extended): `Xd = Xn × Xm`.
Alias for `MADD Xd, Xn, Xm, XZR`.

```
SMADDL Xd, Wn, Wm, Xa
```

Signed Multiply-Add Long: `Xd = Xa + (Wn × Wm)`, treating source operands as signed.

```
SMSUBL Xd, Wn, Wm, Xa
```

Signed Multiply-Subtract Long: `Xd = Xa – (Wn × Wm)`, treating source operands as signed.

```
SMNEGL Xd, Wn, Wm
```

Signed Multiply-Negate Long: `Xd = -(Wn × Wm)`, treating source operands as signed.
Alias for `SMSUBL Xd, Wn, Wm, XZR`.

```
SMULL Xd, Wn, Wm
```

Signed Multiply Long: `Xd = Wn × Wm`, treating source operands as signed.
Alias for `SMADDL Xd, Wn, Wm, XZR`.

```
SMULH Xd, Xn, Xm
```

Signed Multiply High: `Xd = (Xn × Xm)<127:64>`, treating source operands as signed.

```
UMADDL Xd, Wn, Wm, Xa
```

Unsigned Multiply-Add Long: `Xd = Xa + (Wn × Wm)`, treating source operands as unsigned.

```
UMSUBL Xd, Wn, Wm, Xa
```

Unsigned Multiply-Subtract Long: `Xd = Xa – (Wn × Wm)`, treating source operands as unsigned.

```
UMNEGL Xd, Wn, Wm
```

Unsigned Multiply-Negate Long: `Xd = -(Wn × Wm)`, treating source operands as unsigned.
Alias for `UMSUBL Xd, Wn, Wm, XZR`.

```
UMULL Xd, Wn, Wm
```

Unsigned Multiply Long: `Xd = Wn × Wm`, treating source operands as unsigned.
Alias for `UMADDL Xd, Wn, Wm, XZR`.

```
UMULH Xd, Xn, Xm
```

Unsigned Multiply High: `Xd = (Xn × Xm)<127:64>`, treating source operands as unsigned.

5.5.2 Divide

The integer divide instructions compute (numerator÷denominator) and deliver the quotient, which is rounded towards zero. The remainder may then be computed as numerator–(quotient×denominator) using the `MSUB` instruction.

If a signed integer division (INT_MIN ÷ -1) is performed, where INT_MIN is the most negative integer value representable in the selected register size, then the result will overflow the signed integer range. No indication of this overflow is produced and the result written to the destination register will be INT_MIN.

NOTE: The divide instructions do not generate a trap upon division by zero, but write zero to the destination register.

```
SDIV Wd, Wn, Wm
```

Signed Divide: `Wd = Wn ÷ Wm`, treating source operands as signed.

```
SDIV Xd, Xn, Xm
```

Signed Divide (extended): `Xd = Xn ÷ Xm`, treating source operands as signed.

```
UDIV Wd, Wn, Wm
```

Unsigned Divide: `Wd = Wn ÷ Wm`, treating source operands as unsigned.

```
UDIV Xd, Xn, Xm
```

Unsigned Divide (extended): `Xd = Xn ÷ Xm`, treating source operands as unsigned.

5.6 Scalar Floating-point

The A64 scalar floating point instruction set is based closely on ARM VFPv4, and unless explicitly mentioned in individual instruction descriptions the handling and generation of denormals, infinities, non-numerics, and floating point exceptions, replicates the behaviour of the equivalent VFPv4 instructions. Full details may be found in the floating point pseudocode.

5.6.1 Floating-point/SIMD Scalar Memory Access

The FP/SIMD scalar load-store instructions operate on the scalar form of the FP/SIMD registers as described in §4.4.2.1. The available memory addressing modes (see §4.5) are identical to the general-purpose register load-store instructions, and like those instructions permit arbitrary address alignment unless strict alignment checking is enabled. However, unlike the general-purpose load-store instructions, the FP/SIMD load-store instructions make no guarantee of atomicity, even when the address is naturally aligned to the size of data.

5.6.1.1 Load-Store Single FP/SIMD Register

The most general forms of load-store support a range of addressing modes, consisting of base register `Xn` or `SP`, plus one of:

- Scaled, 12-bit, unsigned immediate offset, <u>without</u> pre- and post-index options.
- Unscaled, 9-bit, signed immediate offset, with pre- and post-index options.
- Scaled or unscaled 64-bit register offset.
- Scaled or unscaled 32-bit extended register offset.

Additionally:

- For loads of 32 bits or larger only, a PC-relative address within ±1MiB of the program counter.

`LDR Bt, addr`

> Load Register (byte): load a byte from memory addressed by `addr` to 8-bit `Bt`.

`LDR Ht, addr`

> Load Register (half): load a halfword from memory addressed by `addr` to 16-bit `Ht`.

`LDR St, addr`

> Load Register (single): load a word from memory addressed by `addr` to 32-bit `St`.

`LDR Dt, addr`

> Load Register (double): load a doubleword from memory addressed by `addr` to 64-bit `Dt`.

`LDR Qt, addr`

> Load Register (quad): load a quadword from memory addressed by `addr` and pack into 128-bit `Qt`.

`STR Bt, addr`

> Store Register (byte): store byte from 8-bit `Bt` to memory addressed by `addr`.

`STR Ht, addr`

> Store Register (half): store halfword from 16-bit `Ht` to memory addressed by `addr`.

`STR St, addr`

> Store Register (single): store word from 32-bit `St` to memory addressed by `addr`.

`STR Dt, addr`

> Store Register (double): store doubleword from 64-bit `Dt` to memory addressed by `addr`.

```
STR Qt, addr
```
>   Store Register (quad): store quadword from 128-bit `Qt` to memory addressed by `addr`.

5.6.1.2 Load-Store Single FP/SIMD Register (unscaled offset)

Provides explicit access to the unscaled, 9-bit, signed offset form of load/store instruction, see §5.2.2 for more information about this mnemonic.

```
LDUR Bt, [base,#simm9]
```
>   Load (Unscaled) Register (byte): load a byte from memory addressed by `base+simm9` to 8-bit `Bt`.

```
LDUR Ht, [base,#simm9]
```
>   Load (Unscaled) Register (half): load a halfword from memory addressed by `base+simm9` to 16-bit `Ht`.

```
LDUR St, [base,#simm9]
```
>   Load (Unscaled) Register (single): load a word from memory addressed by `base+simm9` to 32-bit `St`.

```
LDUR Dt, [base,#simm9]
```
>   Load (Unscaled) Register (double): load a doubleword from memory addressed by `base+simm9` to 64-bit `Dt`.

```
LDUR Qt, [base,#simm9]
```
>   Load (Unscaled) Register (quad): load a quadword from memory addressed by `base+simm9` and pack into 128-bit `Qt`.

```
STUR Bt, [base,#simm9]
```
>   Store (Unscaled) Register (byte): store byte from 8-bit `Bt` to memory addressed by `base+simm9`.

```
STUR Ht, [base,#simm9]
```
>   Store (Unscaled) Register (half): store halfword from 16-bit `Ht` to memory addressed by `base+simm9`.

```
STUR St, [base,#simm9]
```
>   Store (Unscaled) Register (single): store word from 32-bit `St` to memory addressed by `base+simm9`.

```
STUR Dt, [base,#simm9]
```
>   Store (Unscaled) Register (double): store doubleword from 64-bit `Dt` to memory addressed by `base+simm9`.

```
STUR Qt, [base,#simm9]
```
>   Store (Unscaled) Register (quad): store quadword from 128-bit `Qt` to memory addressed by `base+simm9`.

5.6.1.3 Load-Store FP/SIMD Pair

The load-store pair instructions support an addressing mode consisting of base register `Xn` or `SP`, plus:

- Scaled, 7-bit, signed immediate offset, with pre- and post-index options

If a Load Pair instruction specifies the same register for the two registers that are being loaded, then one of the following behaviours can occur:

- The instruction is UNALLOCATED
- The instruction is treated as a NOP
- The instruction performs all of the loads using the specified addressing mode and the register being loaded takes an UNKNOWN value

---

LDP St1, St2, addr

> Load Pair (single): load two consecutive words from memory addressed by `addr` to 32-bit `St1` and `St2`.

LDP Dt1, Dt2, addr

> Load Pair (double): load two consecutive doublewords from memory addressed by `addr` to 64-bit `Dt1` and `Dt2`.

LDP Qt1, Qt2, addr

> Load Pair (quad): load two consecutive quadwords from memory addressed by `addr` and to 128-bit `Qt1` and `Qt2`.

STP St1, St2, addr

> Store Pair (single): store two consecutive words from 32-bit `St1` and `St2` to memory addressed by `addr`.

STP Dt1, Dt2, addr

> Store Pair (double): store two consecutive doublewords from 64-bit `Dt1` and `Dt2` to memory addressed by `addr`.

STP Qt1, Qt2, addr

> Store Pair (quad): store two consecutive quadwords from 128-bit `Qt1` and `Qt2` to memory addressed by `addr`.

5.6.1.4 Load-Store FP/SIMD Non-Temporal Pair

The load-store non-temporal pair instructions provide a hint to the memory system that the data being accessed is "non-temporal", i.e. it is a "streaming" access to memory which is unlikely to be referenced again in the near future, and need not be retained in data caches.

As a special exception to the normal memory ordering rules, where an address dependency exists between two memory reads and the second read was generated by a Load Non-temporal Pair instruction then, in the absence of any other barrier mechanism to achieve order, those memory accesses can be observed in any order by other observers within the shareability domain of the memory addresses being accessed.

The load-store non-temporal pair instructions support an addressing mode of base register `Xn` or `SP`, plus:

- Scaled, 7-bit, signed immediate offset, <u>without</u> pre- and post-index options

If a Load Non-temporal Pair instruction specifies the same register for the two registers that are being loaded, then one of the following behaviours can occur:

- The instruction is UNALLOCATED
- The instruction is treated as a NOP
- The instruction performs all of the loads using the specified addressing mode and the register being loaded takes an UNKNOWN value

LDNP St1, St2, [base,#imm]

> Load Non-temporal Pair (single): load two consecutive words from memory addressed by `base+imm` to 32-bit `St1` and `St2`, with a non-temporal hint.

LDNP Dt1, Dt2, [base,#imm]

> Load Non-temporal Pair (double): load two consecutive doublewords from memory addressed by `base+imm` to 64-bit `Dt1` and `Dt2`, with a non-temporal hint.

LDNP Qt1, Qt2, [base,#imm]

> Load Non-temporal Pair (quad): load two consecutive quadwords from memory addressed by `base+imm` to 128-bit `Qt1` and `Qt2`, with a non-temporal hint.

```
STNP St1, St2, [base,#imm]
```

> Store Non-temporal Pair (single): store two consecutive words from 32-bit `St1` and `St2` to memory addressed by `base+imm`, with a non-temporal hint.

```
STNP Dt1, Dt2, [base,#imm]
```

> Store Non-temporal Pair (double): store two consecutive doublewords from 64-bit `Dt1` and `Dt2` to memory addressed by `base+imm`, with a non-temporal hint.

```
STNP Qt1, Qt2, [base,#imm]
```

> Store Non-temporal Pair (quad): store two consecutive quadwords from 128-bit `Qt1` and `Qt2` to memory addressed by `base+imm`, with a non-temporal hint.

5.6.2 Floating-point Move (register)

```
FMOV Sd, Sn
```

> Move 32 bits unchanged from `Sn` to `Sd`.

```
FMOV Dd, Dn
```

> Move 64 bits unchanged from `Dn` to `Dd`.

```
FMOV Wd, Sn
```

> Move 32 bits unchanged from `Sn` to `Wd`.

```
FMOV Sd, Wn
```

> Move 32 bits unchanged from `Wn` to `Sd`.

```
FMOV Xd, Dn
```

> Move 64 bits unchanged from `Dn` to `Xd`.

```
FMOV Dd, Xn
```

> Move 64 bits unchanged from `Xn` to `Dd`.

```
FMOV Xd, Vn.D[1]
```

> Move 64 bits unchanged from `Vn<127:64>` to `Xd`.

```
FMOV Vd.D[1], Xn
```

> Move 64 bits unchanged from `Xn` to `Vd<127:64>`, leaving the other bits in `Vd` unchanged.

5.6.3 Floating-point Move (immediate)

The floating point constant `fpimm` may be specified either in decimal notation (e.g. "12.0" or "-1.2e1"), or as a string beginning "0x" followed by the hexadecimal representation of its IEEE754 encoding. A disassembler should prefer the decimal notation, so long as the value can be displayed precisely.

The floating point value must be expressable as $\pm n \div 16 \times 2^r$, where n and r are integers such that $16 \leq n \leq 31$ and $-3 \leq r \leq 4$, i.e. a normalized binary floating point encoding with 1 sign bit, 4 bits of fraction and a 3-bit exponent.

Note that this encoding does not include the value 0.0, however this value may be loaded using a floating-point move (register) instruction of the form `FMOV Sd,WZR`.

```
FMOV Sd, #fpimm
```

> Single-precision floating-point move immediate `Sd = fpimm`.

```
FMOV Dd, #fpimm
```

> Double-precision floating-point move immediate `Dd = fpimm`.

5.6.4 Floating-point Convert

5.6.4.1 Convert to/from Floating-point

```
FCVT Sd, Hn
```

Convert from half-precision scalar in `Hn` to single-precision in `Sd`.

```
FCVT Hd, Sn
```

Convert from single-precision scalar in `Sn` to half-precision in `Hd`.

```
FCVT Dd, Hn
```

Convert from half-precision scalar in `Hn` to double-precision in `Dd`.

```
FCVT Hd, Dn
```

Convert from double-precision scalar in `Dn` to half-precision in `Hd`.

```
FCVT Dd, Sn
```

Convert from single-precision scalar in `Sn` to double-precision in `Dd`.

```
FCVT Sd, Dn
```

Convert from double-precision scalar in `Dn` to single-precision in `Sd`.

5.6.4.2 Convert to/from Integer

These instructions raise the Invalid Operation exception (`FPSR.IOC`) in response to a floating point input of NaN, Infinity, or a numerical value that cannot be represented within the destination register. An out of range integer result will also be saturated to the destination size. A numeric result which differs from the input will raise the Inexact exception (`FPSR.IXC`). When flush-to-zero mode is enabled a denormal input will be replaced by a zero and will raise the Input Denormal exception (`FPSR.IDC`).

```
FCVTAS Wd, Sn
```

Convert single-precision scalar in `Sn` to nearest signed 32-bit integer in `Wd`, with halfway cases rounding away from zero.

```
FCVTAS Xd, Sn
```

Convert single-precision scalar in `Sn` to nearest signed 64-bit integer in `Xd`, with halfway cases rounding away from zero.

```
FCVTAS Wd, Dn
```

Convert double-precision scalar in `Dn` to nearest signed 32-bit integer in `Wd`, with halfway cases rounding away from zero.

```
FCVTAS Xd, Dn
```

Convert double-precision scalar in `Dn` to nearest signed 64-bit integer in `Xd`, with halfway cases rounding away from zero.

```
FCVTAU Wd, Sn
```

Convert single-precision scalar in `Sn` to nearest unsigned 32-bit integer in `Wd`, with halfway cases rounding away from zero.

```
FCVTAU Xd, Sn
```

Convert single-precision scalar in `Sn` to nearest unsigned 64-bit integer in `Xd`, with halfway cases rounding away from zero.

```
FCVTAU Wd, Dn
```

Convert double-precision scalar in `Dn` to nearest unsigned 32-bit integer in `Wd`, with halfway cases rounding away from zero.

```
FCVTAU Xd, Dn
```

Convert double-precision scalar in `Dn` to nearest unsigned 64-bit integer in `Xd`, with halfway cases rounding away from zero.

```
FCVTMS Wd, Sn
```

Convert single-precision scalar in `Sn` to signed 32-bit integer in `Wd`, rounding towards -∞ (RM).

```
FCVTMS Xd, Sn
```

Convert single-precision scalar in `Sn` to signed 64-bit integer in `Xd`, rounding towards -∞ (RM).

```
FCVTMS Wd, Dn
```

Convert double-precision scalar in `Dn` to signed 32-bit integer in `Wd`, rounding towards -∞ (RM).

```
FCVTMS Xd, Dn
```

Convert double-precision scalar in `Dn` to signed 64-bit integer in `Xd`, rounding towards -∞ (RM).

```
FCVTMU Wd, Sn
```

Convert single-precision scalar in `Sn` to unsigned 32-bit integer in `Wd`, rounding towards -∞ (RM).

```
FCVTMU Xd, Sn
```

Convert single-precision scalar in `Sn` to unsigned 64-bit integer in `Xd`, rounding towards -∞ (RM).

```
FCVTMU Wd, Dn
```

Convert double-precision scalar in `Dn` to unsigned 32-bit integer in `Wd`, rounding towards -∞ (RM).

```
FCVTMU Xd, Dn
```

Convert double-precision scalar in `Dn` to unsigned 64-bit integer in `Xd`, rounding towards -∞ (RM).

```
FCVTNS Wd, Sn
```

Convert single-precision scalar in `Sn` to signed 32-bit integer in `Wd`, with halfway cases rounding to even (RN).

```
FCVTNS Xd, Sn
```

Convert single-precision scalar in `Sn` to signed 64-bit integer in `Xd`, with halfway cases rounding to even (RN).

```
FCVTNS Wd, Dn
```

Convert double-precision scalar in `Dn` to nearest signed 32-bit integer in `Wd`, with halfway cases rounding to even (RN).

```
FCVTNS Xd, Dn
```

Convert double-precision scalar in `Dn` to nearest signed 64-bit integer in `Xd`, with halfway cases rounding to even (RN).

```
FCVTNU Wd, Sn
```

Convert single-precision scalar in `Sn` to nearest unsigned 32-bit integer in `Wd`, with halfway cases rounding to even (RN).

```
FCVTNU Xd, Sn
```

Convert single-precision scalar in `Sn` to nearest unsigned 64-bit integer in `Xd`, with halfway cases rounding to even (RN).

```
FCVTNU Wd, Dn
```

Convert double-precision scalar in `Dn` to nearest unsigned 32-bit integer in `Wd`, with halfway cases rounding to even (RN).

```
FCVTNU Xd, Dn
```

Convert double-precision scalar in `Dn` to nearest unsigned 64-bit integer in `Xd`, with halfway cases rounding to even (RN).

```
FCVTPS Wd, Sn
```

Convert single-precision scalar in `Sn` to signed 32-bit integer in `Wd`, rounding towards +∞ (RP).

FCVTPS Xd, Sn

Convert single-precision scalar in Sn to signed 64-bit integer in Xd, rounding towards +∞ (RP).

FCVTPS Wd, Dn

Convert double-precision scalar in Dn to signed 32-bit integer in Wd, rounding towards +∞ (RP).

FCVTPS Xd, Dn

Convert double-precision scalar in Dn to signed 64-bit integer in Xd, rounding towards +∞ (RP).

FCVTPU Wd, Sn

Convert single-precision scalar in Sn to unsigned 32-bit integer in Wd, rounding towards +∞ (RP).

FCVTPU Xd, Sn

Convert single-precision scalar in Sn to unsigned 64-bit integer in Xd, rounding towards +∞ (RP).

FCVTPU Wd, Dn

Convert double-precision scalar in Dn to unsigned 32-bit integer in Wd, rounding towards +∞ (RP).

FCVTPU Xd, Dn

Convert double-precision scalar in Dn to unsigned 64-bit integer in Xd, rounding towards +∞ (RP).

FCVTZS Wd, Sn

Convert single-precision scalar in Sn to signed 32-bit integer in Wd, rounding towards zero (RZ).

FCVTZS Xd, Sn

Convert single-precision scalar in Sn to signed 64-bit integer in Xd, rounding towards zero (RZ).

FCVTZS Wd, Dn

Convert double-precision scalar in Dn to signed 32-bit integer in Wd, rounding towards zero (RZ).

FCVTZS Xd, Dn

Convert double-precision scalar in Dn to signed 64-bit integer in Xd, rounding towards zero (RZ).

FCVTZU Wd, Sn

Convert single-precision scalar in Sn to unsigned 32-bit integer in Wd, rounding towards zero (RZ).

FCVTZU Xd, Sn

Convert single-precision scalar in Sn to unsigned 64-bit integer in Xd, rounding towards zero (RZ).

FCVTZU Wd, Dn

Convert double-precision scalar in Dn to unsigned 32-bit integer in Wd, rounding towards zero (RZ).

FCVTZU Xd, Dn

Convert double-precision scalar in Dn to unsigned 64-bit integer in Xd, rounding towards zero (RZ).

SCVTF Sd, Wn

Convert signed 32-bit integer in Wn to single-precision scalar in Sd, using FPCR rounding mode.

SCVTF Sd, Xn

Convert signed 64-bit integer in Xn to single-precision scalar in Sd, using FPCR rounding mode.

SCVTF Dd, Wn

Convert signed 32-bit integer in Wn to double-precision scalar in Dd, using FPCR rounding mode.

SCVTF Dd, Xn

Convert signed 64-bit integer in Xn to double-precision scalar in Dd, using FPCR rounding mode.

UCVTF Sd, Wn

Convert unsigned 32-bit integer in Wn to single-precision scalar in Sd, using FPCR rounding mode.

UCVTF Sd, Xn

> Convert unsigned 64-bit integer in Xn to single-precision scalar in Sd, using FPCR rounding mode.

UCVTF Dd, Wn

> Convert unsigned 32-bit integer in Wn to double-precision scalar in Dd, using FPCR rounding mode.

UCVTF Dd, Xn

> Convert unsigned 64-bit integer in Xn to double-precision scalar in Dd, using FPCR rounding mode.

5.6.4.3 Convert to/from Fixed-point

The #fbits operand indicates that the general register holds a fixed-point number with fbits bits after the binary point, where fbits is in the range 1 to 32 for a 32-bit general register, or 1 to 64 for a 64-bit general register.

These instructions raise the Invalid Operation exception (FPSR.IOC) in response to a floating point input of NaN, Infinity, or a numerical value that cannot be represented within the destination register. An out of range fixed-point result will also be saturated to the destination size. A numeric result which differs from the input will raise the Inexact exception (FPSR.IXC). When flush-to-zero mode is enabled a denormal input will be replaced by a zero and will raise the Input Denormal exception (FPSR.IDC).

FCVTZS Wd, Sn, #fbits

> Convert single-precision scalar in Sn to signed 32-bit fixed-point in Wd, rounding towards zero.

FCVTZS Xd, Sn, #fbits

> Convert single-precision scalar in Sn to signed 64-bit fixed-point in Xd, rounding towards zero.

FCVTZS Wd, Dn, #fbits

> Convert double-precision scalar in Dn to signed 32-bit fixed-point in Wd, rounding towards zero.

FCVTZS Xd, Dn, #fbits

> Convert double-precision scalar in Dn to signed 64-bit fixed-point in Xd, rounding towards zero.

FCVTZU Wd, Sn, #fbits

> Convert single-precision scalar in Sn to unsigned 32-bit fixed-point in Wd, rounding towards zero.

FCVTZU Xd, Sn, #fbits

> Convert single-precision scalar in Sn to unsigned 64-bit fixed-point in Xd, rounding towards zero.

FCVTZU Wd, Dn, #fbits

> Convert double-precision scalar in Dn to unsigned 32-bit fixed-point in Wd, rounding towards zero.

FCVTZU Xd, Dn, #fbits

> Convert double-precision scalar in Dn to unsigned 64-bit fixed-point in Xd, rounding towards zero.

SCVTF Sd, Wn, #fbits

> Convert signed 32-bit fixed-point in Wn to single-precision scalar in Sd, using FPCR rounding mode.

SCVTF Sd, Xn, #fbits

> Convert signed 64-bit fixed-point in Xn to single-precision scalar in Sd, using FPCR rounding mode.

SCVTF Dd, Wn, #fbits

> Convert signed 32-bit fixed-point in Wn to double-precision scalar in Dd, using FPCR rounding mode.

SCVTF Dd, Xn, #fbits

> Convert signed 64-bit fixed-point in Xn to double-precision scalar in Dd, using FPCR rounding mode.

```
UCVTF   Sd, Wn, #fbits
```

Convert unsigned 32-bit fixed-point in `Wn` to single-precision scalar in `Sd`, using FPCR rounding mode.

```
UCVTF   Sd, Xn, #fbits
```

Convert unsigned 64-bit fixed-point in `Xn` to single-precision scalar in `Sd`, using FPCR rounding mode.

```
UCVTF   Dd, Wn, #fbits
```

Convert unsigned 32-bit fixed-point in `Wn` to double-precision scalar in `Dd`, using FPCR rounding mode.

```
UCVTF   Dd, Xn, #fbits
```

Convert unsigned 64-bit fixed-point in `Xn` to double-precision scalar in `Dd`, using FPCR rounding mode.

5.6.5 Floating-point Round to Integral

The round to integral instructions round a floating-point value to an integral floating-point value of the same size. The only FPSR exception flags that can be raised by these instructions are: FPSR.IOC (Invalid Operation) for a Signaling NaN input; FPSR.IDC (Input Denormal) for a denormal input when flush-to-zero mode is enabled; for FRINTX only the FPSR.IXC (Inexact) exception if the result is numeric and does not have the same numerical value as the source. A zero input gives a zero result with the same sign, an infinite input gives an infinite result with the same sign, and a NaN is propagated as in normal arithmetic.

```
FRINTA  Sd, Sn
```

Round to nearest integral with halfway cases rounding away from zero, single-precision, from `Sn` to `Sd`.

```
FRINTA  Dd, Dn
```

Round to nearest integral with halfway cases rounding away from zero, double-precision, from `Dn` to `Dd`.

```
FRINTI  Sd, Sn
```

Round to integral using FPCR rounding mode, single-precision, from `Sn` to `Sd`.

```
FRINTI  Dd, Dn
```

Round to integral using FPCR rounding mode, double-precision, from `Dn` to `Dd`,.

```
FRINTM  Sd, Sn
```

Round to integral towards -∞, single-precision, from `Sn` to `Sd`.

```
FRINTM  Dd, Dn
```

Round to integral towards -∞, double-precision, from `Dn` to `Dd`.

```
FRINTN  Sd, Sn
```

Round to nearest integral with halfway cases rounding to even, single-precision, from `Sn` to `Sd`,

```
FRINTN  Dd, Dn
```

Round to nearest integral with halfway cases rounding to even, double-precision from `Dn` to `Dd`.

```
FRINTP  Sd, Sn
```

Round to integral towards +∞, single-precision, from `Sn` to `Sd`.

```
FRINTP  Dd, Dn
```

Round to integral towards +∞, double-precision, from `Dn` to `Dd`.

```
FRINTX  Sd, Sn
```

Round to integral exact using FPCR rounding mode, single-precision, from `Sn` to `Sd`.
For a numerical input sets the Inexact flag if result does not have the same value as the input.

```
FRINTX  Dd, Dn
```

Round to integral exact using FPCR rounding mode, double-precision, from `Dn` to `Dd`.
For a numerical input sets the Inexact flag if result does not have the same value as the input.

`FRINTZ Sd, Sn`

> Round to integral towards zero, single-precision, from `Sn` to `Sd`.

`FRINTZ Dd, Dn`

> Round to integral towards zero, double-precision, from `Dn` to `Dd`.

5.6.6 Floating-point Arithmetic (1 source)

`FABS Sd, Sn`

> Single-precision floating-point scalar absolute value: `Sd = abs(Sn)`.

`FABS Dd, Dn`

> Double-precision floating-point scalar absolute value: `Dd = abs(Dn)`.

`FNEG Sd, Sn`

> Single-precision floating-point scalar negation: `Sd = -Sn`.

`FNEG Dd, Dn`

> Double-precision floating-point scalar negation: `Dd = -Dn`.

`FSQRT Sd, Sn`

> Single-precision floating-point scalar square root: `Sd = sqrt(Sn)`.

`FSQRT Dd, Dn`

> Double-precision floating-point scalar square root: `Dd = sqrt(Dn)`.

5.6.7 Floating-point Arithmetic (2 source)

`FADD Sd, Sn, Sm`

> Single-precision floating-point scalar addition: `Sd = Sn + Sm`.

`FADD Dd, Dn, Dm`

> Double-precision floating-point scalar addition: `Dd = Dn + Dm`.

`FDIV Sd, Sn, Sm`

> Single-precision floating-point scalar division: `Sd = Sn / Sm`.

`FDIV Dd, Dn, Dm`

> Double-precision floating-point scalar division: `Dd = Dn / Dm`.

`FMUL Sd, Sn, Sm`

> Single-precision floating-point scalar multiply: `Sd = Sn * Sm`.

`FMUL Dd, Dn, Dm`

> Double-precision floating-point scalar multipy: `Dd = Dn * Dm`.

`FNMUL Sd, Sn, Sm`

> Single-precision floating-point scalar multiply-negate: `Sd = -(Sn * Sm)`.

`FNMUL Dd, Dn, Dm`

> Double-precision floating-point scalar multiply-negate: `Dd = -(Dn * Dm)`.

`FSUB Sd, Sn, Sm`

> Single-precision floating-point scalar subtraction: `Sd = Sn - Sm`.

`FSUB Dd, Dn, Dm`

> Double-precision floating-point scalar subtraction: `Dd = Dn - Dm`.

5.6.8 Floating-point Min/Max

The `min(x,y)` and `max(x,y)` operations behave similarly to the ARM v7 VMIN.F and VMAX.F instructions and return a quiet NaN when either `x` or `y` is a NaN. In flush-to-zero mode subnormal operands are flushed to zero before comparison, and if a flushed value is then the appropriate result the zero value is returned. Where both `x` and `y` are zero (or subnormal values flushed to zero) with differing sign, then +0.0 is returned by `max()` and -0.0 by `min()`.

The `minNum(x,y)` and `maxNum(x,y)` operations follow the IEEE 754-2008 standard and return the numerical operand when one operand is numerical and the other a quiet NaN. Apart from this additional handling of a single quiet NaN the result is then identical to `min(x,y)` and `max(x,y)`.

FMAX Sd, Sn, Sm

> Single-precision floating-point scalar maximum: `Sd = max(Sn,Sm)`.

FMAX Dd, Dn, Dm

> Double-precision floating-point scalar maximum: `Dd = max(Dn,Dm)`.

FMAXNM Sd, Sn, Sm

> Single-precision floating-point scalar max number: `Sd = maxNum(Sn,Sm)`.

FMAXNM Dd, Dn, Dm

> Double-precision floating-point scalar max number: `Dd = maxNum(Dn,Dm)`.

FMIN Sd, Sn, Sm

> Single-precision floating-point scalar minimum: `Sd = min(Sn,Sm)`.

FMIN Dd, Dn, Dm
> Double-precision floating-point scalar minimum: `Dd = min(Dn,Dm)`.

FMINNM Sd, Sn, Sm

> Single-precision floating-point scalar min number: `Sd = minNum(Sn,Sm)`.

FMINNM Dd, Dn, Dm

> Double-precision floating-point scalar min number: `Dd = minNum(Dn,Dm)`.

5.6.9 Floating-point Multiply-Add

FMADD Sd, Sn, Sm, Sa

> Single-precision floating-point scalar fused multiply-add: `Sd = Sa + Sn*Sm`.

FMADD Dd, Dn, Dm, Da

> Double-precision floating-point scalar fused multiply-add: `Dd = Da + Dn*Dm`.

FMSUB Sd, Sn, Sm, Sa

> Single-precision floating-point scalar fused multiply-subtract: `Sd = Sa + (-Sn)*Sm`.

FMSUB Dd, Dn, Dm, Da

> Double-precision floating-point scalar fused multiply-subtract: `Dd = Da + (-Dn)*Dm`.

FNMADD Sd, Sn, Sm, Sa

> Single-precision floating-point scalar negated fused multiply-add: `Sd = (-Sa) + (-Sn)*Sm`.

FNMADD Dd, Dn, Dm, Da

> Double-precision floating-point scalar negated fused multiply-add: `Dd = (-Da) + (-Dn)*Dm`.

FNMSUB Sd, Sn, Sm, Sa

> Single-precision floating-point scalar negated fused multiply-subtract: `Sd = (-Sa) + Sn*Sm`.

FNMSUB Dd, Dn, Dm, Da

> Double-precision floating-point scalar negated fused multiply-subtract: `Dd = (-Da) + Dn*Dm`.

5.6.10 Floating-point Comparison

These instructions set the integer NZCV condition flags directly, and do not alter the condition flags in the FPSR. In the conditional compare instructions, the `#uimm4` operand is a bitmask used to set the NZCV flags when the input condition is false, with bit 3 setting the N flag, bit 2 the Z flag, bit 1 the C flag, and bit 0 the V flag. If floating-point comparisons are *unordered* the C and V flag bits are set and the N and Z bits cleared.

FCMP   Sn, Sm|#0.0

> Single-precision compare: set condition flags from floating point comparison of Sn with Sm or 0.0.
> Invalid Operation exception only on signaling NaNs.

FCMP   Dn, Dm|#0.0

> Double-precision compare: set condition flags from floating point comparison of Dn with Dm or 0.0.
> Invalid Operation exception only on signaling NaNs.

FCMPE   Sn, Sm|#0.0

> Single-precision compare, exceptional: set flags from floating point comparison of Sn with Sm or 0.0.
> Invalid Operation exception on all NaNs.

FCMPE   Dn, Dm|#0.0

> Double-precision compare, exceptional: set flags from floating point comparison of Dn with Dm or 0.0.
> Invalid Operation exception on all NaNs.

FCCMP   Sn, Sm, #uimm4, cond

> Single-precision conditional compare: `NZCV = if cond then FPCompare(Sn, Sm) else uimm4.`
> Invalid Operation exception only on signaling NaNs when cond holds true.

FCCMP   Dn, Dm, #uimm4, cond

> Double-precision conditional compare: `NZCV = if cond then FPcompare(Dn, Dm) else uimm4.`
> Invalid Operation exception only on signaling NaNs when cond holds true.

FCCMPE Sn, Sm, #uimm4, cond

> Single-precision conditional compare, exceptional:
> `NZCV = if cond then FPCompare(Sn, Sm) else uimm4.`
> Invalid Operation exception on all NaNs when cond holds true.

FCCMPE Dn, Dm, #uimm4, cond

> Double-precision conditional compare, exceptional:
> `NZCV = if cond then FPCompare(Dn, Dm) else uimm4.`
> Invalid Operation exception on all NaNs when cond holds true.

5.6.11 Floating-point Conditional Select

FCSEL   Sd, Sn, Sm, cond

> Single-precision conditional select: `Sd = if cond then Sn else Sm.`

FCSEL   Dd, Dn, Dm, cond

> Double-precision conditional select: `Dd = if cond then Dn else Dm.`

5.7 Advanced SIMD

5.7.1 Overview

AArch64 Advanced SIMD is based upon the existing AArch32 Advanced SIMD extension, with the following changes:

- In AArch64 Advanced SIMD, there are thirty two 128-bit wide vector registers, whereas AArch32 Advanced SIMD had sixteen 128-bit wide registers.

- There are thirty two 64-bit vectors and these are held in the lower 64 bits of each 128-bit register.

- Writes of 64 bits or less to a vector register result in the higher bits being zeroed (except for lane inserts).

- New lane insert and extract instructions have been added to support the new register packing scheme.

- Additional widening instructions are provided for generating the top 64 bits of a 128-bit vector register.

- Data-processing instructions which would generate more than one result register (e.g. widening a 128-bit vector), or consume more than three sources (e.g. narrowing a 128-bit vector), have been split into separate instructions.

- A set of scalar instructions have been added to implement loop heads and tails, but only where the instruction does not already exist in the main scalar floating-point instruction set, and only when "over-computing" using a vector form might have the side effect of setting the saturation or floating point exception flags if there was "garbage" in unused higher lanes. Scalar operations on 64-bit integers are also provided in this section, to avoid the cost of over-computing using a 128-bit vector.

- A new set of vector "reduction" operations provide across-lane sum, minimum and maximum.

- Some existing instructions have been extended to support 64-bit integer values: e.g. comparison, addition, absolute value and negate, including saturating versions.

- Advanced SIMD now supports both single-precision (32-bit) and double-precision (64-bit) floating-point vector data types and arithmetic as defined by the IEEE 754 floating-point standard, honoring the FPCR Rounding Mode field, the Default NaN control, the Flush-to-Zero control, and (where supported by the implementation) the Exception trap enable bits.

- The ARMv7 SIMD "chained" floating-point multiply-accumulate instructions have been replaced with IEEE754 "fused" multiply-add. This includes the reciprocal step and reciprocal square root step instructions.

- Convert float to integer (FCVTxU, FCVTxS) encode a directed rounding mode: towards zero, towards +Inf, towards –Inf, to nearest with ties to even, and to nearest with ties away from zero.

- Round float to nearest integer in floating-point format (FRINTx) has been added, with the same directed rounding modes, as well as rounding according to the ambient rounding mode.

- A new double to single precision down-convert instruction with "exact" rounding, suitable for ongoing single to half-precision down-conversion with correct double to half rounding (FCVTXN).

- IEEE 754-2008 minNum() and maxNum() instructions have been added (FMINNM, FMAXNM).

- Instructions to accelerate floating point vector normalisation have been added (FRECPX, FMULX).

- Saturating instructions have been extended to include unsigned accumulate into signed, and vice-versa.

5.7.2 Advanced SIMD Mnemonics

Although derived from the AArch32 Advanced SIMD syntax, a number of changes have been made to harmonise with the AArch64 core integer and floating point instruction set syntax, and to unify AArch32's divergent "architectural" and "programmers'" notations:

- The 'V' mnemonic prefix has been removed, and S/U/F/P added to indicate signed/unsigned/floating-point/polynomial data type. The mnemonic always indicates the data type(s) of the operation.

- The vector organisation (element size and number of lanes) is described by the register qualifiers and never by a mnemonic qualifier. See the description of the vector register syntax in §4.4.2 above.

- The 'P' prefix for "pairwise" operations becomes a suffix.

- A 'V' suffix has been added for the new reduction (across-all-lanes) operations

- A '2' suffix has been added for the new widening/narrowing "second part" instructions, described below.

- Vector compares now use the integer condition code names to indicate whether an integer comparison is signed or unsigned (e.g. CMLT, CMLO, CMGE, CMHI, etc)

- Some mnemonics have been renamed where the removal of the V prefix caused clash with the core instruction set mnemonics.

With the exception of the above changes, the mnemonics are based closely on AArch32 Advanced SIMD. As such, the learning curve for existing Advanced SIMD programmers is reduced. A full list of the equivalent AArch32 mnemonics can be found in §5.7.23 below.

Widening instructions with a '2' suffix implement the "second" or "top" part of a widening operation that would otherwise need to write two 128-bit vectors: they get their input data from the high numbered lanes of the 128-bit source vectors, and write the expanded results to the 128-bit destination.

Narrowing instructions with a '2' suffix implement the "second" or "top" part of a narrowing operation that would otherwise need to read two 128-bit vectors for each source operand: they get their input data from the 128-bit source operands and insert their narrowed results into the high numbered lanes of the 128-bit destination, leaving the lower lanes unchanged.

5.7.3 Data Movement

DUP Vd.<Td>, Vn.<Ts>[index]

> Duplicate element (vector). Replicate single vector element from Vn to all elements of Vd. Where <Td>/<Ts> may be 8B/B, 16B/B, 4H/H, 8H/H, 2S/S, 4S/S or 2D/D. The immediate index is a value in the range 0 to nelem(<Ts>)-1.

DUP Vd.<T>, Wn

> Duplicate 32-bit general register (vector). Replicate low order bits from 32-bit general register Wn to all elements of vector Vd. Where <T> may be 8B, 16B, 4H, 8H, 2S or 4S.

DUP Vd.2D, Xn

> Duplicate 64-bit general register (vector). Replicate 64-bit general register Xn to both elements of vector Vd.

DUP <V>d, Vn.<T>[index]

> Duplicate element (scalar). Copy single vector element from Vn to scalar register <V>d. Where <V>/<T> may be B/B, H/H, S/S or D/D. The immediate index is a value in the range 0 to nelem(<T>)-1. Normally disassembled as MOV.

INS Vd.<T>[index], Vn.<T>[index2]

> Insert element (vector). Inserts a single vector element from Vn into a single element of Vd. Where <T> may be B, H, S or D. Both immediates index and index2 are values in the range 0 to nelem(<T>)-1. Normally disassembled as MOV.

`INS Vd.<T>[index], Wn`

> Insert 32-bit general register (vector). Inserts low order bits from 32-bit general register Wn into a single vector element of Vd. Where <T> may be 8B, 16B, 4H, 8H, 2S or 4S. The immediate index is a value in the range 0 to nelem(<T>)-1. Normally disassembled as MOV.

`INS Vd.D[index], Xn`

> Insert 64-bit general register (vector). Inserts 64-bit general register Xn into a single vector element of Vd. The immediate index is a value in the range 0 to 1. Normally disassembled as MOV.

`MOV Vd.<T>[index], Vn.<T>[index2]`

> Move element. Moves a vector element from Vn to a vector element in Vd: alias for INS Vd.<T>[index],Vn.<T>[index2].

`MOV Vd.<T>[index], Wn`

> Move 32-bit general register to element. Moves a 32-bit general register Wn to vector element in Vd: alias for INS Vd.<T>[index],Wn.

`MOV Vd.2D[index], Xn`

> Move 64-bit general register to element. Moves a 64-bit general register Xn to a vector element in Vd: alias for INS Vd.D[index],Xn.

`MOV <V>d, Vn.<T>[index]`

> Move element (scalar). Moves a vector element from Vn to scalar register <V>d: alias for DUP <V>d,Vn.<T>[index].

`MOV <V>d, <V>n`

> Move (scalar). Moves a scalar register <V>n to scalar register <V>d: alias for DUP <V>d,Vn.<V>[0].

`UMOV Wd, Vn.<Ts>[index]`

> Unsigned integer move element to 32-bit general register. Zero-extends an integer vector element from Vn into 32-bit general register Wd. Where <Ts> may be 8B, 16B, 4H, 8H, 2S or 4S. The index is in the range 0 to nelem(<Ts>)-1.

`UMOV Xd, Vn.D[index]`

> Unsigned integer move element to 64-bit general register. Moves an unsigned 64-bit integer vector element from Vn into 64-bit general register Wd. The immediate index is in the range 0 to 1.

`SMOV Wd, Vn.<T>[index]`

> Signed integer move element to 32-bit general register. Sign-extends an integer vector element from Vn into 32-bit general register Wd. Where <T> may be B or H. The index is a value is in the range 0 to nelem(<T>)-1.

`SMOV Xd, Vn.<T>[index]`

> Signed integer move element to 64-bit general register. Sign-extends an integer vector element from Vn into 64-bit general register Xd. Where <T> may be B, H or S. The index is in the range 0 to nelem(<T>)-1.

5.7.4 Vector Arithmetic

`UABA Vd.<T>, Vn.<T>, Vm.<T>`

> Unsigned integer absolute difference and accumulate (vector). Subtracts the elements of Vm from the corresponding elements of Vn, and accumulates the absolute values of the results into the elements of Vd. Operand and result elements are all unsigned integers of the same length: <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`SABA Vd.<T>, Vn.<T>, Vm.<T>`

> Signed integer absolute difference and accumulate (vector). Subtracts the elements of Vm from the corresponding elements of Vn, and accumulates the absolute values of the results into the elements of Vd. Operand and result elements are all signed integers of the same length: <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`UABD Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer absolute difference (vector). Subtracts the elements of Vm from the corresponding elements of Vn, and places the absolute values of the results in the elements of Vd. Operand and result elements are all integers of the same length: <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`SABD Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer absolute difference (vector). Subtracts the elements of Vm from the corresponding elements of Vn, and places the absolute values of the results in the elements of Vd. Operand and result elements are all integers of the same length: <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`FABD Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point absolute difference (vector). Subtracts the elements of Vm from the corresponding elements of Vn, and places the absolute values of the results in the elements of Vd. Operand and result elements are all of the same length: <T> is 2S, 4S or 2D.

`ADD Vd.<T>, Vn.<T>, Vm.<T>`

Integer add (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`FADD Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point add (vector). Where <T> is 2S, 4S or 2D.

`AND Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise AND (vector). Where <T> is 8B or 16B (though an assembler should accept any valid format).

`BIC Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise bit clear (vector). Where <T> is 8B or 16B (though an assembler should accept any valid format).

`BIF Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise insert if false (vector). Where <T> is 8B or 16B (though an assembler should accept any valid format).

`BIT Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise insert if true (vector). Where <T> is 8B or 16B (though an assembler should accept any valid format).

`BSL Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise select (vector). Where <T> is 8B or 16B (though an assembler should accept any valid format).

`CMEQ Vd.<T>, Vn.<T>, Vm.<T>`

Integer compare mask equal (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

`CMEQ Vd.<T>, Vn.<T>, #0`

Integer compare mask equal to zero (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

`CMHS Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer compare mask higher or same (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

`CMGE Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer compare mask greater than or equal (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

`CMGE Vd.<T>, Vn.<T>, #0`

Signed integer compare mask greater than or equal to zero (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

`CMHI Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer compare mask higher (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

`CMGT Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer compare mask greater than (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
CMGT  Vd.<T>, Vn.<T>, #0
```

Signed integer compare mask greater than zero (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
CMLS  Vd.<T>, Vn.<T>, Vm.<T>
```

Unsigned integer compare mask lower or same (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.
Alias for CMHS with operands reversed.

```
CMLE  Vd.<T>, Vn.<T>, Vm.<T>
```

Signed integer compare mask less than or equal (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.
Alias for CMGE with operands reversed.

```
CMLE  Vd.<T>, Vn.<T>, #0
```

Signed integer compare mask less than or equal to zero (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
CMLO  Vd.<T>, Vn.<T>, Vm.<T>
```

Unsigned integer compare mask lower (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.
Alias for CMHI with operands reversed.

```
CMLT  Vd.<T>, Vn.<T>, Vm.<T>
```

Signed integer compare mask less than (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.
Alias for CMGT with operands reversed.

```
CMLT  Vd.<T>, Vn.<T>, #0
```

Signed integer compare mask less than zero (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
CMTST  Vd.<T>, Vn.<T>, Vm.<T>
```

Integer compare mask bitwise test (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
FCMEQ  Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point compare mask equal (vector). Where <T> is 2S, 4S or 2D.

```
FCMEQ  Vd.<T>, Vn.<T>, #0
```

Floating-point compare mask equal to zero (vector). Where <T> is 2S, 4S or 2D.

```
FCMGE  Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point compare mask greater than or equal (vector). Where <T> is 2S, 4S or 2D.

```
FCMGE  Vd.<T>, Vn.<T>, #0
```

Floating-point compare mask greater than or equal to zero (vector). Where <T> is 2S, 4S or 2D.

```
FCMGT  Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point compare mask greater than (vector). Where <T> is 2S, 4S or 2D.

```
FCMGT  Vd.<T>, Vn.<T>, #0
```

Floating-point compare mask greater than zero (vector). Where <T> is 2S, 4S or 2D.

```
FCMLE  Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point compare mask less than or equal (vector). Where <T> is 2S, 4S or 2D.
Alias for FCMGE with operands reversed.

```
FCMLE  Vd.<T>, Vn.<T>, #0
```

Floating-point compare mask less than or equal to zero (vector). Where <T> is 2S, 4S or 2D.

```
FCMLT  Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point compare mask less than (vector). Where <T> is 2S, 4S or 2D.
Alias for FCMGT with operands reversed.

```
FCMLT  Vd.<T>, Vn.<T>, #0
```

Floating-point compare mask less than zero (vector). Where <T> is 2S, 4S or 2D.

`FACGE Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point absolute compare mask greater than or equal (vector). Where <T> is 2S, 4S or 2D.

`FACGT Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point absolute compare mask greater than (vector). Where <T> is 2S, 4S or 2D.

`FACLE Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point absolute compare mask less than or equal (vector). Where <T> is 2S, 4S or 2D.
Alias for FACGE with operands reversed.

`FACLT Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point absolute compare mask less than (vector). Where <T> is 2S, 4S or 2D.
Alias for FACGT with operands reversed.

`FDIV Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point divide (vector). Where <T> is 2S, 4S or 2D.

`EOR Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise exclusive OR (vector). Where <T> is 8B or 16B (an assembler should accept any valid arrangement).

`UHADD Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer halving add (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`SHADD Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer halving add (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`UHSUB Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer halving subtract (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`SHSUB Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer halving subtract (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`UMAX Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer maximum (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`SMAX Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer maximum (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`FMAX Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point maximum (vector). Where <T> is 2S, 4S or 2D.

`FMAXNM Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point maxNum (vector). Where <T> is 2S, 4S or 2D.

`UMIN Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer minimum (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`SMIN Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer minimum (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`FMIN Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point minimum (vector). Where <T> is 2S, 4S or 2D.

`FMINNM Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point minNum (vector). Where <T> is 2S, 4S or 2D.

`MLA Vd.<T>, Vn.<T>, Vm.<T>`

Integer multiply-accumulate (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`FMLA Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point fused multiply-accumulate (vector). Where <T> is 2S, 4S or 2D.

`MLS Vd.<T>, Vn.<T>, Vm.<T>`

Integer multiply-subtract from accumulator (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`FMLS Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point fused multiply-subtract from accumulator (vector). Where <T> is 2S, 4S or 2D.

`MUL Vd.<T>, Vn.<T>, Vm.<T>`

Integer multiply (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`FMUL Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point multiply (vector). Where <T> is 2S, 4S or 2D.

`FMULX Vd.<T>, Vn.<T>, Vm.<T>`

Floating-point multiply extended, like FMUL but $0 \times \pm\infty \to \pm 2$ (vector). Where <T> is 2S, 4S or 2D.

`PMUL Vd.<T>, Vn.<T>, Vm.<T>`

Polynomial multiply (vector). Where <T> is 8B or 16B.

`ORN Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise OR NOT (vector). Where <T> is 8B or 16B (an assembler should accept any valid arrangement).

`ORR Vd.<T>, Vn.<T>, Vm.<T>`

Bitwise OR (vector). Where <T> is 8B or 16B (an assembler should accept any valid arrangement).

`SQADD Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer saturating add (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

`UQADD Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer saturating add (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`SQDMULH Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer saturating doubling multiply high half (vector). Where <T> is 4H, 8H, 2S or 4S.

`SQRDMULH Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer saturating rounding doubling multiply high half (vector). Where <T> is 4H, 8H, 2S or 4S.

`UQRSHL Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer saturating rounding shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`SQRSHL Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer saturating rounding shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`UQSUB Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer saturating subtract (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`SQSUB Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer saturating subtract (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`URHADD Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer rounding halving add (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`SRHADD Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer rounding halving add (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

`URSHL Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer rounding shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`SRSHL Vd.<T>, Vn.<T>, Vm.<T>`

Signed integer rounding shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

`UQSHL Vd.<T>, Vn.<T>, Vm.<T>`

Unsigned integer saturating shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

SQSHL Vd.<T>, Vn.<T>, Vm.<T>

Signed integer saturating shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

USHL Vd.<T>, Vn.<T>, Vm.<T>

Unsigned integer shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

SSHL Vd.<T>, Vn.<T>, Vm.<T>

Signed integer shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

SUB Vd.<T>, Vn.<T>, Vm.<T>

Integer subtract (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D

FSUB Vd.<T>, Vn.<T>, Vm.<T>

Floating-point subtract (vector). Where <T> is 2S, 4S or 2D.

FRECPS Vd.<T>, Vn.<T>, Vm.<T>

Floating-point reciprocal step (vector). Where <T> is 2S, 4S or 2D. The embedded multiply-accumulate is fused in AArch64 FRECPS, whilst in AArch32 VRECPS it remains chained.

FRSQRTS Vd.<T>, Vn.<T>, Vm.<T>

Floating-point reciprocal square root step (vector). Where <T> is 2S, 4S or 2D. The embedded multiply-accumulate is fused in AArch64 FRSQRTS, whilst in AArch32 VRSQRTS it remains chained.

5.7.5 Scalar Arithmetic

FABD <V>d, <V>n, <V>m

Floating-point absolute difference (scalar). Subtracts <V>m from <V>n, and places the absolute value of the result in <V>d. Where <V> is S or D.

ADD Dd, Dn, Dm

Integer add (scalar).

CMEQ Dd, Dn, Dm

Integer compare mask equal (scalar).

CMEQ Dd, Dn, #0

Integer compare mask equal to zero (scalar).

CMHS Dd, Dn, Dm

Unsigned integer compare mask higher or same (scalar).

CMGE Dd, Dn, Dm

Signed integer compare mask greater than or equal (scalar).

CMGE Dd, Dn, #0

Signed integer compare mask greater than or equal to zero (scalar).

CMHI Dd, Dn, Dm

Unsigned integer compare mask higher (scalar).

CMGT Dd, Dn, Dm

Signed integer compare mask greater than (scalar).

CMGT Dd, Dn, #0

Signed integer compare mask greater than zero (scalar).

CMLS Dd, Dn, Dm

Unsigned integer compare mask lower or same (scalar).
Alias for CMHS with operands reversed.

```
CMLE Dd, Dn, Dm
```

Signed integer compare mask less than or equal (scalar).
Alias for CMGE with operands reversed.

```
CMLE Dd, Dn, #0
```

Signed integer compare mask less than or equal to zero (scalar).

```
CMLO Dd, Dn, Dm
```

Unsigned integer compare mask lower (scalar).
Alias for CMHI with operands reversed.

```
CMLT Dd, Dn, Dm
```

Signed integer compare mask less than (scalar).
Alias for CMGT with operands reversed.

```
CMLT Dd, Dn, #0
```

Signed integer compare mask less than zero (scalar).

```
CMTST Dd, Dn, Dm
```

Integer compare mask bitwise test (scalar).

```
FCMEQ <V>d, <V>n, <V>m
```

Floating-point compare mask equal (scalar). Where <V>is S or D.

```
FCMEQ <V>d, <V>n, #0
```

Floating-point compare mask equal to zero (scalar). Where <V>is S or D.

```
FCMGE <V>d, <V>n, <V>m
```

Floating-point compare mask greater than or equal (scalar). Where <V>is S or D.

```
FCMGE <V>d, <V>n, #0
```

Floating-point compare mask greater than or equal to zero (scalar). Where <V>is S or D.

```
FCMGT <V>d, <V>n, <V>m
```

Floating-point compare mask greater than (scalar). Where <V>is S or D.

```
FCMGT <V>d, <V>n, #0
```

Floating-point compare mask greater than zero (scalar). Where <V>is S or D.

```
FCMLE <V>d, <V>n, <V>m
```

Floating-point compare mask less than or equal (scalar). Where <V>is S or D.
Alias for FCMGE with operands reversed.

```
FCMLE <V>d, <V>n, #0
```

Floating-point compare mask less than or equal to zero (scalar). Where <V>is S or D.

```
FCMLT <V>d, <V>n, <V>m
```

Floating-point compare mask less than (scalar). Where <V>is S or D.
Alias for FCMGT with operands reversed.

```
FCMLT <V>d, <V>n, #0
```

Floating-point compare mask less than zero (scalar). Where <V>is S or D.

```
FACGE <V>d, <V>n, <V>m
```

Floating-point absolute compare mask greater than or equal (scalar). Where <V>is S or D.

```
FACGT <V>d, <V>n, <V>m
```

Floating-point absolute compare mask greater than (scalar). Where <V>is S or D.

`FACLE <V>d, <V>n, <V>m`

Floating-point absolute compare mask less than or equal (scalar). Where <V>is S or D.
Alias for FACGE with operands reversed.

`FACLT <V>d, <V>n, <V>m`

Floating-point absolute compare mask less than (scalar). Where <V>is S or D.
Alias for FACGT with operands reversed.

`SQADD <V>d, <V>n, <V>m`

Signed integer saturating add (scalar). Where <V> is B, H, S or D.

`UQADD <V>d, <V>n, <V>m`

Unsigned integer saturating add (scalar). Where <V> is B, H, S or D.

`SQDMULH <V>d, <V>n, <V>m`

Signed integer saturating doubling multiply high half (scalar). Where <V> is H or S.

`SQRDMULH <V>d, <V>n, <V>m`

Signed integer saturating rounding doubling multiply high half (scalar). Where <V> is H or S.

`UQRSHL <V>d, <V>n, <V>m`

Unsigned integer saturating rounding shift left (scalar). Where <V> is B, H, S or D.

`SQRSHL <V>d, <V>n, <V>m`

Signed integer saturating rounding shift left (scalar). Where <V> is B, H, S or D.

`UQSUB <V>d, <V>n, <V>m`

Unsigned integer saturating subtract (scalar). Where <V> is B, H, S or D.

`SQSUB <V>d, <V>n, <V>m`

Signed integer saturating subtract (scalar). Where <V> is B, H, S or D.

`UQSHL <V>d, <V>n, <V>m`

Unsigned integer saturating shift left (scalar). Where <V> is B, H, S or D.

`SQSHL <V>d, <V>n, <V>m`

Signed integer saturating shift left (scalar). Where <V> is B, H, S or D.

`URSHL Dd, Dn, Dm`

Unsigned integer rounding shift left (scalar).

`SRSHL Dd, Dn, Dm`

Signed integer rounding shift left (scalar).

`USHL Dd, Dn, Dm`

Unsigned integer shift left (scalar).

`SSHL Dd, Dn, Dm`

Signed integer shift left (scalar).

`SUB Dd, Dn, Dm`

Integer subtract (scalar).

`FMULX <V>d, <V>n, <V>m`

Floating-point multiply extended, like `FMUL` but $0 \times \pm\infty \to \pm 2$ (scalar). Where <V> is S or D.

`FRECPS <V>d, <V>n, <V>m`

Floating-point reciprocal step (scalar). Where <V>is S or D. The embedded multiply-accumulate is fused in AArch64 `FRECPS`, whilst in AArch32 `VRECPS` it remains chained.

```
FRSQRTS <V>d, <V>n, <V>m
```

Floating-point reciprocal square root step (scalar). Where <V>is S or D. The embedded multiply-accumulate is fused in AArch64 `FRSQRTS`, whilst in AArch32 `VRSQRTS` it remains chained.

5.7.6 Vector Widening/Narrowing Arithmetic

```
UABAL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer absolute difference and accumulate long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
UABAL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer absolute difference and accumulate long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
SABAL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer absolute difference and accumulate long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
SABAL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer absolute difference and accumulate long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
UABDL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer absolute difference long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
UABDL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer absolute difference long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
SABDL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer absolute difference long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
SABDL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer absolute difference long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
UADDL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer add long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
UADDL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer add long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
SADDL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer add long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
SADDL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer add long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
USUBL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer subtract long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
USUBL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Unsigned integer subtract long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
SSUBL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer subtract long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
SSUBL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer subtract long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

`UMLAL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Unsigned integer multiply-accumulate long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

`UMLAL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Unsigned integer multiply-accumulate long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

`SMLAL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer multiply-accumulate long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

`SMLAL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer multiply-accumulate long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

`UMLSL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Unsigned integer multiply-subtract from accumulator long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

`UMLSL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Unsigned integer multiply-subtract from accumulator long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

`SMLSL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer multiply-subtract from accumulator long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

`SMLSL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer multiply-subtract from accumulator long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

`UMULL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Unsigned integer multiply long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

`UMULL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Unsigned integer multiply long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

`SMULL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer multiply long (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

`SMULL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer multiply long (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

`PMULL Vd.8H, Vn.8B, Vm.8B`

Polynomial multiply long (vector).

`PMULL2 Vd.8H, Vn.16B, Vm.16B`

Polynomial multiply long (vector, second part).

`SQDMLAL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer saturating doubling multiply accumulate long (vector). Where the <Td>/<Ts> is 4S/4H or 2D/2S.

`SQDMLAL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer saturating doubling multiply accumulate long (vector, second part). Where the <Td>/<Ts> is 4S/8H or 2D/4S.

`SQDMLSL Vd.<Td>, Vn.<Ts>, Vm.<Ts>`

Signed integer saturating doubling multiply subtract from accumulator long (vector). Where the <Td>/<Ts> is 4S/4H or 2D/2S.

```
SQDMLSL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer saturating doubling multiply subtract from accumulator long (vector, second part). Where the <Td>/<Ts> is 4S/8H or 2D/4S.

```
SQDMULL Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer saturating doubling multiply long (vector). Where the <Td>/<Ts> is 4S/4H or 2D/2S.

```
SQDMULL2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Signed integer saturating doubling multiply long (vector, second part). Where the <Td>/<Ts> is 4S/8H or 2D/4S.

```
UADDW Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Unsigned integer add wide (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
UADDW2 Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Unsigned integer add wide (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
SADDW Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Signed integer add wide (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
SADDW2 Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Signed integer add wide (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
USUBW Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Unsigned integer subtract wide (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
USUBW2 Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Unsigned integer subtract wide (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
SSUBW Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Signed integer subtract wide (vector). Where the <Td>/<Ts> is 8H/8B, 4S/4H or 2D/2S.

```
SSUBW2 Vd.<Td>, Vn.<Td>, Vm.<Ts>
```

Signed integer subtract wide (vector, second part). Where the <Td>/<Ts> is 8H/16B, 4S/8H or 2D/4S.

```
RADDHN Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer rounding add and narrow high half (vector). Where the <Td>/<Ts> is 8B/8H, 4H/4S or 2S/2D.

```
RADDHN2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer rounding add and narrow high half (vector, second part). Where the <Td>/<Ts> is 16B/8H, 8H/4S or 4S/2D.

```
RSUBHN Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer rounding subtract and narrow high half (vector). Where the <Td>/<Ts> is 8B/8H, 4H/4S or 2S/2D.

```
RSUBHN2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer rounding subtract and narrow high half (vector, second part). Where the <Td>/<Ts> is 16B/8H, 8H/4S or 4S/2D.

```
ADDHN Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer add and narrow high half (vector). Where the <Td>/<Ts> is 8B/8H, 4H/4S or 2S/2D.

```
ADDHN2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer add and narrow high half (vector, second part). Where the <Td>/<Ts> is 16B/8H, 8H/4S or 4S/2D.

```
SUBHN Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer subtract and narrow high half (vector). Where the <Td>/<Ts> is 8B/8H, 4H/4S or 2S/2D.

```
SUBHN2 Vd.<Td>, Vn.<Ts>, Vm.<Ts>
```

Integer subtract and narrow high half (vector, second part). Where the <Td>/<Ts> is 16B/8H, 8H/4S or 4S/2D.

5.7.7 Scalar Widening/Narrowing Arithmetic

SQDMLAL <Vd>d, <Vs>n, <Vs>m

> Signed integer saturating doubling multiply accumulate long (scalar). Where the <Vd>/<Vs> is H/B, S/H or D/S.

SQDMLSL <Vd>d, <Vs>n, <Vs>m

> Signed integer saturating doubling multiply subtract from accumulator long (scalar). Where the <Vd>/<Vs> is S/H or D/S.

SQDMULL <Vd>d, <Vs>n, <Vs>m

> Signed integer saturating doubling multiply long (scalar). Where the <Vd>/<Vs> is S/H or D/S.

5.7.8 Vector Unary Arithmetic

ABS Vd.<T>, Vn.<T>

> Integer absolute value (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

SQABS Vd.<T>, Vn.<T>

> Signed integer saturating absolute (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

FABS Vd.<T>, Vn.<T>

> Floating-point absolute value (vector). Where <T> is 2S, 4S or 2D.

NEG Vd.<T>, Vn.<T>

> Integer negate (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

SQNEG Vd.<T>, Vn.<T>

> Signed integer saturating negate (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

FNEG Vd.<T>, Vn.<T>

> Floating-point negate (vector). Where <T> is 2S, 4S or 2D.

CLS Vd.<T>, Vn.<T>

> Signed integer count leading sign bits (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

CLZ Vd.<T>, Vn.<T>

> Integer count leading zero bits (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

CNT Vd.<T>, Vn.<T>

> Count non-zero bits (vector). Where <T> is 8B or 16B.

NOT Vd.<T>, Vn.<T>

> Bitwise invert (vector). Where <T> is 8B or 16B (an assembler should accept any valid arrangement). Normally disassembled as MVN.

MVN Vd.<T>, Vn.<T>

> Bitwise invert (vector). Where <T> is 8B or 16B (an assembler should accept any valid arrangement). Alias for NOT Vd.<T>,Vn.<T>

SUQADD Vd.<T>, Vn.<T>

> Signed integer saturating accumulate of unsigned value (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

USQADD Vd.<T>, Vn.<T>

> Unsigned integer saturating accumulate of signed value (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

UADALP Vd.<Td>, Vn.<Ts>

> Unsigned integer add and accumulate long pairwise (vector). Where <Td>/<Ts> is 4H/8B, 8H/16B, 2S/4H, 4S/8H, 1D/2S or 2D/4S.

`SADALP Vd.<Td>, Vn.<Ts>`

Signed integer add and accumulate long pairwise (vector). Where <Td>/<Ts> is 4H/8B, 8H/16B, 2S/4H, 4S/8H, 1D/2S or 2D/4S.

`UADDLP Vd.<Td>, Vn.<Ts>`

Unsigned integer add long pair (vector). Where <Td>/<Ts> is 4H/8B, 8H/16B, 2S/4H, 4S/8H, 1D/2S or 2D/4S.

`SADDLP Vd.<Td>, Vn.<Ts>`

Signed integer add long pair (vector). Where <Td>/<Ts> is 4H/8B, 8H/16B, 2S/4H, 4S/8H, 1D/2S or 2D/4S.

`FCVTL Vd.<Td>, Vn.<Ts>`

Floating-point convert long half-precision to single-precision, or single-precision to double-precision (vector). Where <Td>/<Ts> is 4S/4H or 2D/2S

`FCVTL2 Vd.<Td>, Vn.<Ts>`

Floating-point convert long half-precision to single-precision, or single-precision to double-precision (vector, second part). Where <Td>/<Ts> is 4S/8H or 2D/4S

`XTN Vd.<Td>, Vn.<Ts>`

Integer narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D.

`XTN2 Vd.<Td>, Vn.<Ts>`

Integer narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D.

`SQXTUN Vd.<Td>, Vn.<Ts>`

Signed integer saturating and unsigned narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D.

`SQXTUN2 Vd.<Td>, Vn.<Ts>`

Signed integer saturating and unsigned narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D.

`UQXTN Vd.<Td>, Vn.<Ts>`

Unsigned integer saturating narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D.

`UQXTN2 Vd.<Td>, Vn.<Ts>`

Unsigned integer saturating narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D.

`SQXTN Vd.<Td>, Vn.<Ts>`

Signed integer saturating narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D.

`SQXTN2 Vd.<Td>, Vn.<Ts>`

Signed integer saturating narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D.

`FCVTN Vd.<Td>, Vn.<Ts>`

Floating-point convert narrow single-precision to half-precision, or double-precision to single-precision (vector). Where <Td>/<Ts> is 4H/4S or 2S/2D.

`FCVTN2 Vd.<Td>, Vn.<Ts>`

Floating-point convert narrow single-precision to half-precision, or double-precision to single-precision (vector, second part). Where <Td>/<Ts> is 8H/4S or 4S/2D.

`FCVTXN Vd.2S, Vn.2D`

Floating-point convert narrow double-precision to single-precision with "exact" rounding (vector). The result is only suitable for further narrowing to half-precision without losing precision due to rounding twice.

`FCVTXN2 Vd.4S, Vn.2D`

Floating-point convert narrow double-precision to single-precision with "exact" rounding (vector, second part). The result is only suitable for further narrowing to half-precision without losing precision due to rounding twice.

`FRINTx Vd.<T>, Vn.<T>`

> Floating-point round to integral (vector). Where <T> is 2S, 4S or 2D. The letter x selects the rounding mode: N (nearest, ties to even); A (nearest, ties away from zero), P (towards +Inf); M (towards –Inf), Z (towards zero), I (using FPCR rounding mode) and X (using FPCR rounding mode, with exactness test).

`FSQRT Vd.<T>, Vn.<T>`

> Floating-point square root (vector). Where <T> is 2S, 4S or 2D.

`URECPE Vd.<T>, Vn.<T>`

> Unsigned integer reciprocal estimate (vector). Where <T> is 2S or 4S.

`FRECPE Vd.<T>, Vn.<T>`

> Floating-point reciprocal estimate (vector). Where <T> is 2S, 4S or 2D.

`URSQRTE Vd.<T>, Vn.<T>`

> Unsigned integer reciprocal square root estimate (vector). Where <T> is 2S or 4S.

`FRSQRTE Vd.<T>, Vn.<T>`

> Floating-point reciprocal square root estimate (vector). Where <T> is 2S, 4S or 2D.

`RBIT Vd.<T>, Vn.<T>`

> Bit reverse (vector): reverses the bits within each byte vector element. Where <T> is 8B or 16B.

`REV16  Vd.<T>, Vn.<T>`

> Element reverse in 16-bit halfwords (vector). Where <T> is 8B or 16B.

`REV32 Vd.<T>, Vn.<T>`

> Element reverse in 32-bit words (vector). Where <T> is 8B, 16B, 4H, or 8H.

`REV64 Vd.<T>, Vn.<T>`

> Element reverse in 64-bit doublewords (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

5.7.9 Scalar Unary Arithmetic

`ABS Dd, Dn`

> Signed 64-bit integer absolute (scalar).

`SQABS <V>d, <V>n`

> Signed integer saturating absolute (scalar). Where <V> is B, H, S or D.

`NEG Dd, Dn`

> Signed 64-bit integer negate (scalar).

`SQNEG <V>d, <V>n`

> Signed integer saturating negate (scalar). Where <V> is B, H, S or D.

`SUQADD <V>d, <V>n`

> Signed integer saturating accumulate of unsigned value (scalar). Where <V> is B, H, S or D.

`USQADD <V>d, <V>n`

> Unsigned integer saturating accumulate of signed value Where <V> is B, H, S or D.

`SQXTUN <Vd>d, <Vs>n`

> Signed integer saturating and unsigned narrow (scalar). Where <Vd>/<Vs> is B/H, H/S or S/D.

`UQXTN <Vd>d, <Vs>n`

> Unsigned integer saturating narrow (scalar). Where <Vd>/<Vs> is B/H, H/S or S/D.

`SQXTN <Vd>d, <Vs>n`

> Signed integer saturating narrow (scalar). Where <Vd>/<Vs> is B/H, H/S or S/D.

```
FCVTXN Sd, Dn
```

Floating-point convert narrow double-precision to single-precision with "exact" rounding (scalar).

```
FRECPE <V>d, <V>n
```

Floating-point reciprocal estimate (scalar). Where <V> is S or D.

```
FRECPX <V>d, <V>n
```

Floating-point reciprocal exponent (scalar). Where <V> is S or D.

```
FRSQRTE <V>d, <V>n
```

Floating-point reciprocal square root estimate (scalar). Where <V> is S or D.

5.7.10 Vector-by-element Arithmetic

In all cases the immediate `index` is a constant in the range 0 to nelem(<Ts>)–1.

```
FMLA Vd.<T>, Vn.<T>, Vm.<Ts>[index]
```

Floating-point fused multiply add (vector, by element). Where <T>/<Ts> is 2S/S, 4S/S or 2D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
FMLS Vd.<T>, Vn.<T>, Vm.<Ts>[index]
```

Floating-point fused multiply subtract (vector, by element). Where <T>/<Ts> is 2S/S, 4S/S or 2D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
FMUL Vd.<T>, Vn.<T>, Vm.<Ts>[index]
```

Floating-point multiply (vector, by element). Where <Td>/<Ts> is 2S/S 4S/S or 2D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
FMULX Vd.<T>, Vn.<T>, Vm.<Ts>[index]
```

Floating-point multiply extended (vector, by element): like FMUL but $0 \times \pm\infty \rightarrow \pm2$. Where <Td>/<Ts> is 2S/S, 4S/S or 2D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
MLA Vd.<T>, Vn.<T>, Vm.<Ts>[index]
```

Integer multiply accumulate (vector, by element). Where <T>/<Ts> is 4H/H, 8H/H, 2S/S or 4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
MLS Vd.<T>, Vn.<T>, Vm.<Ts>[index]
```

Integer multiply subtract (vector, by element). Where <T>/<Ts> is 4H/H, 8H/H, 2S/S or 4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
MUL Vd.<T>, Vn.<T>, Vm.<Ts>[index]
```

Integer multiply (vector, by element). Where <T>/<Ts> is 4H/H, 8H/H, 2S/S or 4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SMLAL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```

Signed integer multiply accumulate long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SMLAL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```

Signed integer multiply accumulate long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SMLSL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```

Signed integer multiply subtract long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SMLSL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```

Signed integer multiply subtract long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15. If <Ts> is H, then Vm must be in the range V0-V15.

```
SMULL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer multiply long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SMULL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer multiply long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
UMLAL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Unsigned integer multiply accumulate long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
UMLAL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Unsigned integer multiply accumulate long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
UMLSL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Unsigned integer multiply subtract long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
UMLSL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Unsigned integer multiply subtract long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
UMULL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Unsigned integer multiply long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
UMULL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Unsigned integer multiply long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMLAL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer saturating doubling multiply accumulate long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMLAL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer saturating doubling multiply accumulate long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMLSL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer saturating doubling multiply subtract long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMLSL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer saturating doubling multiply subtract long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMULL Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer saturating doubling multiply long (vector, by element). Where <Ta>/<Tb>/<Ts> is 4S/4H/H or 2D/2S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMULL2 Vd.<Ta>, Vn.<Tb>, Vm.<Ts>[index]
```
Signed integer saturating doubling multiply long (vector, by element, second part). Where <Ta>/<Tb>/<Ts> is 4S/8H/H or 2D/4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMULH Vd.<Td>, Vn.<Td>, Vm.<Ts>[index]
```
Signed integer saturating doubling multiply returning high half (vector, by element). Where <Td>/<Ts> is 4H/H, 8H/H, 2S/S or 4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQRDMULH Vd.<Td>, Vn.<Td>, Vm.<Ts>[index]
```

Signed integer saturating rounding doubling multiply returning high half (vector, by element). Where <Td>/<Ts> is 4H/H, 8H/H, 2S/S or 4S/S. If <Ts> is H, then Vm must be in the range V0-V15.

5.7.11 Scalar-by-element Arithmetic

In all cases the immediate `index` is a constant in the range 0 to nelem(<Ts>)–1.

```
FMLA <V>d, <V>n, Vm.<Ts>[index]
```

Floating-point fused multiply add (scalar, by element). Where <V>/<Ts> is S/S or D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
FMLS <V>d, <V>n, Vm.<Ts>[index]
```

Floating-point fused multiply subtract (scalar, by element). Where <V>/<Ts> is S/S or D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
FMUL <V>d, <V>n, Vm.<Ts>[index]
```

Floating-point multiply (scalar, by element). Where <V>/<Ts> is S/S or D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
FMULX <V>d, <V>n, Vm.<Ts>[index]
```

Floating-point multiply extended (scalar, by element): like FMUL but $0 \times \pm\infty \rightarrow \pm 2$. Where <V>/<Ts> is S/S, or D/D. If <Ts> is S, then Vm must be in the range V0-V15.

```
SQDMLAL <Va>d, <Vb>n, Vm.<Ts>[index]
```

Signed integer saturating doubling multiply accumulate long (scalar, by element). Where <Va>/<Vb>/<Ts> is S/H/H or D/S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMLSL <Va>d, <Vb>n, Vm.<Ts>[index]
```

Signed integer saturating doubling multiply subtract long (scalar, by element). Where <Va>/<Vb>/<Ts> is S/H/H or D/S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMULL <Va>d, <Vb>n, Vm.<Ts>[index]
```

Signed integer saturating doubling multiply long (scalar, by element). Where <Va>/<Vb>/<Ts> is S/H/H or D/S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQDMULH <V>d, <V>n, Vm.<Ts>[index]
```

Signed integer saturating doubling multiply returning high half (scalar, by element). Where <V>/<Ts> is H/H or S/S. If <Ts> is H, then Vm must be in the range V0-V15.

```
SQRDMULH <V>d, <V>n, Vm.<Ts>[index]
```

Signed integer saturating rounding doubling multiply returning high half (scalar, by element). Where <V>/<Ts> is H/H or S/S. If <Ts> is H, then Vm must be in the range V0-V15.

5.7.12 Vector Permute

```
EXT Vd.<T>, Vn.<T>, Vm.<T>, #index
```

Bitwise extract (vector). Where <T> is either 8B or 16B. The index is an immediate value in the range 0 to nelem(<T>)-1.

The following are replacements for the ARMv7 VTRN, VUZP and VZIP instructions which had two destination registers. Semantically these are identical to the ARMv7 instruction except that UZP1/TRN1/ZIP1 produce what would have been the Dn/Qn output of the ARMv7 instruction, whilst UZP2/TRN2/ZIP2 produce what would have been the Dm/Qm output.

```
TRN1 Vd.<T>, Vn.<T>, Vm.<T>
```

Vector element transpose (first part). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
TRN2 Vd.<T>, Vn.<T>, Vm.<T>
```

Vector element transpose (second part). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
UZP1 Vd.<T>, Vn.<T>, Vm.<T>
```

Vector element unzip (first part). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
UZP2 Vd.<T>, Vn.<T>, Vm.<T>
```

Vector element unzip (second part). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
ZIP1 Vd.<T>, Vn.<T>, Vm.<T>
```

Vector element zip (first part). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
ZIP2 Vd.<T>, Vn.<T>, Vm.<T>
```

Vector element zip (second part). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

5.7.13 Vector Immediate

```
MOVI Vn.<T>, #uimm8{, LSL #shift}
```

Move immediate (vector, shifted): replicates LSL(uimm8,shift) into each 32-bit element. Where <T> is 2S or 4S, and shift is 0, 8, 16 or 24 (default 0).

```
MOVI Vn.<T>, #uimm8, MSL #shift
```

Move immediate (vector, masked): replicates MSL(uimm8,shift) into each 32-bit element. Where <T> is 2S or 4S, and shift is 8 or 16. The MSL operator is a left shift, but filling the low order bits with ones instead of zeros.

```
MOVI Vn.<T>, #uimm8{, LSL #shift}
```

Move immediate (vector, shifted): replicates LSL(uimm8,shift) into each 16-bit element. Where <T> is 4H or 8H, and shift is 0 or 8 (default 0).

```
MOVI Vn.<T>, #uimm8
```

Move immediate (vector) : replicates uimm8 into each 8-bit element. Where <T> is 8B or 16B.

```
MOVI Vn.2D, #uimm64
```

Move immediate (vector) : replicates a "byte mask immediate" consisting of 8 bytes, each byte having only the value 0x00 or 0xff, into each 64-bit element.

```
MOVI Dn, #uimm64
```

Move immediate (scalar) : moves a "byte mask" immediate consisting of 8 bytes, each byte having only the value 0x00 or 0xff, into a 64-bit vector register.

```
MVNI Vn.<T>, #uimm8{, LSL #shift}
```

Move inverted immediate (vector, shifted): replicates NOT(LSL(uimm8,shift)) into each 32-bit element. Where <T> is 2S or 4S, and shift is 0, 8, 16 or 24 (default 0).

```
MVNI Vn.<T>, #uimm8, MSL #shift
```

Move inverted immediate (vector, masked): replicates NOT(MSL(uimm8,shift)) into each 32-bit element. Where <T> is 2S or 4S, and shift is 8 or 16. The MSL operator is a left shift, but filling the low order bits with ones instead of zeros.

```
MVNI Vn.<T>, #uimm8{, LSL #shift}
```

Move inverted immediate (vector, shifted): replicates NOT(LSL(uimm8,shift)) into each 16-bit element. Where <T> is 4H or 8H, and shift is 0 or 8 (default 0).

```
FMOV Vn.<T>, #fpimm
```

Floating point move immediate (vector). Where <T> is 2S, 4S or 2D, and fpimm is a floating point constant replicated into each vector element. The constant may be specified either in decimal notation (e.g. "12.0" or "-1.2e1"), or as a string beginning "0x" followed by the hexadecimal representation of its IEEE754 encoding. A disassembler should prefer the decimal notation, so long as the value can be displayed precisely. The floating point value must be expressable as $\pm n \div 16 \times 2^r$, where n and r are integers such that $16 \le n \le 31$ and $-3 \le r \le 4$, i.e. a normalized binary floating point encoding with sign, 4 bits of fraction and a 3-bit exponent.

`BIC Vn.<T>, #uimm8{, LSL #shift}`

> Bitwise bit clear immediate (vector): bitwise AND of NOT(LSL(uimm8,shift)) with each 32-bit element. Where <T> is 2S or 4S, and shift is 0, 8, 16 or 24 (default 0).

`BIC Vn.<T>, #uimm8{, LSL #shift}`

> Bitwise bit clear immediate (vector): bitwise AND of NOT(LSL(uimm8,shift)) with each 16-bit element. Where <T> is 4H or 8H, and shift is 0 or 8 (default 0).

`ORR Vn.<T>, #uimm8{, LSL #shift}`

> Bitwise OR immediate (vector): bitwise OR of LSL(uimm8,shift) with each 32-bit element. Where <T> is 2S or 4S, and shift is 0, 8, 16 or 24 (default 0).

`ORR Vn.<T>, #uimm8{, LSL #shift}`

> Bitwise OR immediate (vector): bitwise OR of LSL(uimm8,shift) with each 16-bit element. Where <T> is 4H or 8H, and shift is 0 or 8 (default 0).

5.7.14 Vector Shift (immediate)

`USHR Vd.<T>, Vn.<T>, #shift`

> Unsigned integer shift right (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`SSHR Vd.<T>, Vn.<T>, #shift`

> Signed integer shift right (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`URSHR Vd.<T>, Vn.<T>, #shift`

> Unsigned integer rounding shift right (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`SRSHR Vd.<T>, Vn.<T>, #shift`

> Signed integer rounding shift right (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`USRA Vd.<T>, Vn.<T>, #shift`

> Unsigned integer shift right and accumulate (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`SSRA Vd.<T>, Vn.<T>, #shift`

> Signed integer shift right and accumulate (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`URSRA Vd.<T>, Vn.<T>, #shift`

> Unsigned integer rounding shift right and accumulate (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`SRSRA Vd.<T>, Vn.<T>, #shift`

> Signed integer rounding shift right and accumulate (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`SRI Vd.<T>, Vn.<T>, #shift`

> Integer shift right and insert (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 1 to elsize(<T>).

`SHRN Vd.<Td>, Vn.<Ts>, #shift`

> Integer shift right narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`SHRN2 Vd.<Td>, Vn.<Ts>, #shift`

Integer shift right narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`UQSHRN Vd.<Td>, Vn.<Ts>, #shift`

Unsigned integer saturating shift right narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`UQSHRN2 Vd.<Td>, Vn.<Ts>, #shift`

Unsigned integer saturating shift right narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`SQSHRN Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating shift right narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`SQSHRN2 Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating shift right narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`RSHRN Vd.<Td>, Vn.<Ts>, #shift`

Integer rounding shift right narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`RSHRN2 Vd.<Td>, Vn.<Ts>, #shift`

Integer rounding shift right narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`UQRSHRN Vd.<Td>, Vn.<Ts>, #shift`

Unsigned integer saturating rounding shift right narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`UQRSHRN2 Vd.<Td>, Vn.<Ts>, #shift`

Unsigned integer saturating rounding shift right narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`SQRSHRN Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating rounding shift right narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`SQRSHRN2 Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating rounding shift right narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`SQSHRUN Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating shift right unsigned narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`SQSHRUN2 Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating shift right unsigned narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`SQRSHRUN Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating rounding shift right unsigned narrow (vector). Where <Td>/<Ts> is 8B/8H, 4H/4S, or 2S/2D; and shift is in the range 1 to elsize(<Td>).

`SQRSHRUN2 Vd.<Td>, Vn.<Ts>, #shift`

Signed integer saturating rounding shift right unsigned narrow (vector, second part). Where <Td>/<Ts> is 16B/8H, 8H/4S, or 4S/2D; and shift is in the range 1 to elsize(<Td>).

`SHL Vd.<T>, Vn.<T>, #shift`

Unsigned integer shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<T>)-1.

`UQSHL Vd.<T>, Vn.<T>, #shift`

Unsigned integer saturating shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<T>)-1.

`SQSHL Vd.<T>, Vn.<T>, #shift`

Signed integer saturating shift left (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<T>)-1.

`SQSHLU Vd.<T>, Vn.<T>, #shift`

Signed integer saturating shift left unsigned (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<T>)-1.

`SLI Vd.<T>, Vn.<T>, #shift`

Integer shift left and insert (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<T>)-1.

`USHLL Vd.<Td>, Vn.<Ts>, #shift`

Unsigned integer shift left long (vector). Where <Td>/<Ts> is 8H/8B, 4S/4H, or 2D/2S; and shift is in the range 0 to elsize(<Ts>)-1.

`USHLL2 Vd.<Td>, Vn.<Ts>, #shift`

Unsigned integer shift left long (vector, second part). Where <Td>/<Ts> is 8H/16B, 4S/8H, or 2D/4S; and shift is in the range 0 to elsize(<Ts>)-1.

`UXTL Vd.<Td>, Vn.<Ts>`

Unsigned integer lengthen (vector). Where <Td>/<Ts> is 8H/8B, 4S/4H, or 2D/2S.
Alias for USHLL Vd.<Td>,Vn.<Ts>,#0.

`UXTL2 Vd.<Td>, Vn.<Ts>`

Unsigned integer lengthen (vector, second part). Where <Td>/<Ts> is 8H/16B, 4S/8H, or 2D/4S.
Alias for USHLL2 Vd.<Td>,Vn.<Ts>,#0.

`SSHLL Vd.<Td>, Vn.<Ts>, #shift`

Signed integer shift left long (vector). Where <Td>/<Ts> is 8H/8B, 4S/4H, or 2D/2S; and shift is in the range 0 to elsize(<Ts>)-1.

`SSHLL2 Vd.<Td>, Vn.<Ts>, #shift`

Signed integer shift left long (vector, second part). Where <Td>/<Ts> is 8H/16B, 4S/8H, or 2D/4S; and shift is in the range 0 to elsize(<Ts>)-1.

`SXTL Vd.<Td>, Vn.<Ts>`

Signed integer lengthen (vector). Where <Td>/<Ts> is 8H/8B, 4S/4H, or 2D/2S.
Alias for SSHLL Vd.<Td>,Vn.<Ts>,#0.

`SXTL2 Vd.<Td>, Vn.<Ts>`

Signed integer lengthen (vector, second part). Where <Td>/<Ts> is 8H/16B, 4S/8H, or 2D/4S.
Alias for SSHLL2 Vd.<Td>,Vn.<Ts>,#0.

5.7.15 Scalar Shift (immediate)

`USHR Dd, Dn, #shift`

Unsigned integer shift right (scalar). Where shift is in the range 1 to 64.

`SSHR Dd, Dn, #shift`

Signed integer shift right (scalar). Where shift is in the range 1 to 64.

`URSHR Dd, Dn, #shift`

Unsigned integer rounding shift right (scalar). Where shift is in the range 1 to 64.

`SRSHR Dd, Dn, #shift`

Signed integer rounding shift right (scalar). Where shift is in the range 1 to 64.

`USRA Dd, Dn, #shift`

Unsigned integer shift right and accumulate (scalar). Where shift is in the range 1 to 64.

`SSRA Dd, Dn, #shift`

Signed integer shift right and accumulate (scalar). Where shift is in the range 1 to 64.

`URSRA Dd, Dn, #shift`

Unsigned integer rounding shift right and accumulate (scalar). Where shift is in the range 1 to 64.

`SRSRA Dd, Dn, #shift`

Signed integer rounding shift right and accumulate (scalar). Where shift is in the range 1 to 64.

`SRI Dd, Dn, #shift`

Integer shift right and insert (scalar). Where shift is in the range 1 to 64.

`UQSHRN <Vd>d, <Vs>n, #shift`

Unsigned integer saturating shift right narrow (scalar). Where <Vd>/<Vs> is B/H, H/S, or S/D; and shift is in the range 1 to elsize(<Vd>).

`SQSHRN <Vd>d, <Vs>n, #shift`

Signed integer saturating shift right narrow (scalar). Where <Vd>/<Vs> is B/H, H/S, or S/D; and shift is in the range 1 to elsize(<Vd>).

`UQRSHRN <Vd>d, <Vs>n, #shift`

Unsigned integer saturating rounding shift right narrow (scalar). Where <Vd>/<Vs> is B/H, H/S, or S/D; and shift is in the range 1 to elsize(<Vd>).

`SQRSHRN <Vd>d, <Vs>n, #shift`

Signed integer saturating rounding shift right narrow (scalar). Where <Vd>/<Vs> is B/H, H/S, or S/D; and shift is in the range 1 to elsize(<Vd>).

`SQSHRUN <Vd>d, <Vs>n, #shift`

Signed integer saturating shift right unsigned narrow (scalar). Where <Vd>/<Vs> is B/H, H/S, or S/D; and shift is in the range 1 to elsize(<Vd>).

`SQRSHRUN <Vd>d, <Vs>n, #shift`

Signed integer saturating rounding shift right unsigned narrow (scalar). Where <Vd>/<Vs> is B/H, H/S, or S/D; and shift is in the range 1 to elsize(<Vd>).

`SHL Dd, Dn, #shift`

Unsigned integer shift left (scalar). Where shift is in the range 0 to 63.

`UQSHL <V>d, <V>n, #shift`

Unsigned integer saturating shift left (scalar). Where <V> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<V>)-1.

`SQSHL <V>d, <V>n, #shift`

Signed integer saturating shift left (scalar). Where <V> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<V>)-1.

`SQSHLU <V>d, <V>n, #shift`

Signed integer saturating shift left unsigned (scalar). Where <V> is 8B, 16B, 4H, 8H, 2S, 4S or 2D; and shift is in the range 0 to elsize(<V>)-1.

```
SLI Dd, Dn, #shift
```

        Integer shift left and insert (scalar). Where shift is in the range 0 to 63.

5.7.16 Vector Floating Point / Integer Convert

These instructions raise the Invalid Operation exception (`FPSR.IOC`) in response to a floating point input of NaN, Infinity, or a numerical value that cannot be represented within the destination register. An out of range integer or fixed-point result will also be saturated to the destination size. A numeric result which differs from the input will raise the Inexact exception (`FPSR.IXC`).

```
FCVTxS Vd.<T>, Vn.<T>
```

        Floating-point convert to signed integer of same size (vector). Where <T> is 2S, 4S or 2D. The letter x selects the rounding mode: N (nearest, ties to even); A (nearest, ties away from zero), P (towards +Inf); M (towards –Inf), Z (towards zero).

```
FCVTZS Vd.<T>, Vn.<T>, #fbits
```

        Floating-point convert to signed fixed-point of same size (vector) with rounding towards zero. Where <T> is 2S, 4S or 2D. The number of fractional bits is represented by fbits in the range 1 to 64.

```
FCVTxU Vd.<T>, Vn.<T>
```

        Floating-point convert to unsigned integer of same size (vector). Where <T> is 2S, 4S or 2D. The letter x selects the rounding mode: N (nearest, ties to even); A (nearest, ties away from zero), P (towards +Inf); M (towards –Inf), Z (towards zero).

```
FCVTZU Vd.<T>, Vn.<T>, #fbits
```

        Floating-point convert to unsigned fixed-point of same size (vector) with rounding towards zero. Where <T> is 2S, 4S or 2D. The number of fractional bits is represented by fbits in the range 1 to 64.

```
SCVTF Vd.<T>, Vn.<T>
```

        Signed integer convert to floating-point of same size (vector). Where <T> is 2S, 4S or 2D.

```
SCVTF Vd.<T>, Vn.<T>, #fbits
```

        Signed fixed-point convert to floating-point of same size (vector). Where <T> is 2S, 4S or 2D. The number of fractional bits is represented by fbits in the range 1 to 64.

```
UCVTF Vd.<T>, Vn.<T>
```

        Unsigned integer convert to floating-point of same size (vector). Where <T> is 2S, 4S or 2D.

```
UCVTF Vd.<T>, Vn.<T>, #fbits
```

        Unsigned fixed-point convert to floating-point of same size (vector). Where <T> is 2S, 4S or 2D. The number of fractional bits is represented by fbits in the range 1 to 64.

5.7.17 Scalar Floating Point / Integer Convert

These instructions raise the Invalid Operation exception (`FPSR.IOC`) in response to a floating point input of NaN, Infinity, or a numerical value that cannot be represented within the destination register. An out of range integer or fixed-point result will also be saturated to the destination size. A numeric result which differs from the input will raise the Inexact exception (`FPSR.IXC`).

```
FCVTxS <V>d, <V>n
```

        Floating-point convert to signed integer of same size (scalar). Where <V> is S or D. The letter x selects the rounding mode: N (nearest, ties to even); A (nearest, ties away from zero), P (towards +Inf); M (towards –Inf), Z (towards zero).

```
FCVTZS <V>d, <V>n, #fbits
```

        Floating-point convert to signed fixed-point of same size (scalar) with rounding towards zero. Where <V> is S or D. The number of fractional bits is represented by fbits in the range 1 to 64.

`FCVTxU <V>d, <V>n`

Floating-point convert to unsigned integer of same size (scalar). Where <V> is S or D. The letter x selects the rounding mode: N (nearest, ties to even); A (nearest, ties away from zero), P (towards +Inf); M (towards –Inf), Z (towards zero).

`FCVTZU <V>d, <V>n, #fbits`

Floating-point convert to unsigned fixed-point of same size (scalar) with rounding towards zero. Where <V> is S or D. The number of fractional bits is represented by fbits in the range 1 to 64.

`SCVTF <V>d, <V>n`

Signed integer convert to floating-point of same size (scalar). Where <V> is S or D.

`SCVTF <V>d, <V>n, #fbits`

Signed fixed-point convert to floating-point of same size (scalar). Where <V> is S or D. The number of fractional bits is represented by fbits in the range 1 to 64.

`UCVTF <V>d, <V>n`

Unsigned integer convert to floating-point of same size (scalar). Where <V> is S or D.

`UCVTF <V>d, <V>n, #fbits`

Unsigned fixed-point convert to floating-point of same size (scalar). Where <V> is S or D. The number of fractional bits is represented by fbits in the range 1 to 64.

5.7.18 Vector Reduce (across lanes)

`ADDV <V>d, Vn.<T>`

Integer sum elements to scalar (vector). Where <V>/<T> is B/8B, B/16B, H/4H, H/8H, S/2S, or S/4S.

`SADDLV <V>d, Vn.<T>`

Signed integer sum elements to scalar long (vector). Where <V>/<T> is H/8B, H/16B, S/4H, S/8H, D/2S, or D/4S.

`UADDLV <V>d, Vn.<T>`

Unsigned integer sum elements to scalar long (vector). Where <V>/<T> is H/8B, H/16B, S/4H, S/8H, D/2S, or D/4S.

`SMAXV <V>d, Vn.<T>`

Signed integer maximum element to scalar (vector). Where <V>/<T> is B/8B, B/16B, H/4H, H/8H, S/2S, or S/4S.

`SMINV <V>d, Vn.<T>`

Signed integer minimum element to scalar (vector). Where <V>/<T> is B/8B, B/16B, H/4H, H/8H, S/2S, or S/4S.

`UMAXV <V>d, Vn.<T>`

Unsigned integer maximum element to scalar (vector). Where <V>/<T> is B/8B, B/16B, H/4H, H/8H, S/2S, or S/4S.

`UMINV <V>d, Vn.<T>`

Unsigned integer minimum element to scalar (vector). Where <V>/<T> is B/8B, B/16B, H/4H, H/8H, S/2S, or S/4S.

`FMAXV Sd, Vn.4S`

Floating-point maximum element to scalar (vector), equivalent to a sequence of pairwise reductions.

`FMAXNMV Sd, Vn.4S`

Floating-point maxNum element to scalar (vector), equivalent to a sequence of pairwise reductions.

`FMINV Sd, Vn.4S`

Floating-point minimum element to scalar (vector), equivalent to a sequence of pairwise reductions.

```
FMINNMV Sd, Vn.4S
```

Floating-point minNum element to scalar (vector), equivalent to a sequence of pairwise reductions.

5.7.19 Vector Pairwise Arithmetic

```
ADDP Vd.<T>, Vn.<T>, Vm.<T>
```

Integer add pair (vector). Where <T> is 8B, 16B, 4H, 8H, 2S, 4S or 2D.

```
FADDP Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point add pair (vector). Where <T> is 2S, 4S or 2D.

```
SMAXP Vd.<T>, Vn.<T>, Vm.<T>
```

Signed integer maximum pair (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

```
UMAXP Vd.<T>, Vn.<T>, Vm.<T>
```

Unsigned integer maximum pair (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

```
FMAXP Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point maximum pair (vector). Where <T> is 2S, 4S or 2D.

```
FMAXNMP Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point maxNum pair (vector). Where <T> is 2S, 4S or 2D.

```
SMINP Vd.<T>, Vn.<T>, Vm.<T>
```

Signed integer minimum pair (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

```
UMINP Vd.<T>, Vn.<T>, Vm.<T>
```

Unsigned integer minimum pair (vector). Where <T> is 8B, 16B, 4H, 8H, 2S or 4S.

```
FMINP Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point minimum pair (vector). Where <T> is 2S, 4S or 2D.

```
FMINNMP Vd.<T>, Vn.<T>, Vm.<T>
```

Floating-point minNum pair (vector). Where <T> is 2S, 4S or 2D.

5.7.20 Scalar Reduce (pairwise)

```
ADDP Dd, Vn.2D
```

Integer pairwise sum (scalar).

```
FADDP <V>d, Vn.<T>
```

Floating-point pairwise sum (scalar). Where <V>/<T> is S/2S or D/2D.

```
FMAXP <V>d, Vn.<T>
```

Floating-point pairwise maximum (scalar). Where <V>/<T> is S/2S or D/2D.

```
FMAXNMP <V>d, Vn.<T>
```

Floating-point pairwise maxNum (scalar). Where <V>/<T> is S/2S or D/2D.

```
FMINP <V>d, Vn.<T>
```

Floating-point pairwise minimum (scalar). Where <V>/<T> is S/2S or D/2D.

```
FMINNMP <V>d, Vn.<T>
```

Floating-point pairwise minNum (scalar). Where <V>/<T> is S/2S or D/2D.

5.7.21 Vector Table Lookup

`TBL Vd.<T>, {Vn*.16B}, Vm.<T>`

Table lookup (vector). Where <T> may be 8B or 16B, and Vn* is a list of between one and four consecutively numbered vector registers each holding sixteen 8-bit table elements. The list braces "{ }" are concrete symbols, and do not indicate an optional field as elsewhere in this manual.

`TBX Vd.<T>, {Vn*.16B}, Vm.<T>`

Table lookup extension (vector). Where <T> may be 8B or 16B, and Vn* is a list of between one and four consecutively numbered vector registers each holding sixteen 8-bit table elements. The list braces "{ }" are concrete symbols, and do not indicate an optional field as elsewhere in this manual.

5.7.22 Vector Load-Store Structure

All SIMD load-store structure instructions use the syntax term `vaddr` as shorthand for the following addressing modes:

`[base]`

> Memory addressed by base register `Xn` or `SP`.

`[base],Xm`

> Memory addressed by base register `Xn` or `SP`, post-incremented by 64-bit index register `Xm`.

`[base],#imm`

> Memory addressed by `Xn` or `SP`, post-incremented by an immediate value which must equal the total number of bytes transferred to/from memory.

Register notation of the form `Vt+n` in the register lists below indicates that the register number is required to be equal to $(t + n)$ MOD 32. Furthemore the list braces "{ }" are concrete symbols, and do not indicate an optional field as elsewhere in this manual.

Like other load-store instructions they permit arbitrary address alignment, unless strict alignment checking is enabled, in which case alignment to the size of the element is checked. However unlike the general-purpose load-store instructions, the vector load-store instructions make no guarantee of atomicity, even when the address is naturally aligned to the size of element.

5.7.22.1 Load-Store Multiple Structures

In all of these instructions `<T>` is one of `8B`, `16B`, `4H`, `8H`, `2S`, `4S`, `2D` and additionally the LD1 and ST1 instructions support the `1D` format. The post-increment immediate offset, if present, must be 8, 16, 24, 32, 48 or 64, depending on the number of elements transferred.

`LD1 {Vt.<T>}, vaddr`

> Load multiple 1-element structures (to one register)

`LD1 {Vt.<T>, Vt+1.<T>}, vaddr`

> Load multiple 1-element structures (to two consecutive registers)

`LD1 {Vt.<T>, Vt+1.<T>, Vt+2.<T>}, vaddr`

> Load multiple 1-element structures (to three consecutive registers)

`LD1 {Vt.<T>, Vt+1.<T>, Vt+2.<T>, Vt+3.<T>}, vaddr`

> Load multiple 1-element structures (to four consecutive registers)

`LD2 {Vt.<T>, Vt+1.<T>}, vaddr`

> Load multiple 2-element structures (to two consecutive registers)

`LD2 {Vt.<T>, Vt+2.<T>}, vaddr`

> Load multiple 2-element structures (to two alternating registers)

`LD3 {Vt.<T>, Vt+1.<T>, Vt+2.<T>}, vaddr`

> Load multiple 3-element structures (to three consecutive registers)

`LD3 {Vt.<T>, Vt+2.<T>, Vt+4.<T>}, vaddr`

> Load multiple 3-element structures (to three alternating registers)

`LD4 {Vt.<T>, Vt+1.<T>, Vt+2.<T>, Vt+3.<T>}, vaddr`

> Load multiple 4-element structures (to four consecutive registers)

`LD4 {Vt.<T>, Vt+2.<T>, Vt+4.<T>, Vt+6.<T>}, vaddr`

> Load multiple 4-element structures (to four alternating registers)

```
ST1 {Vt.<T>}, vaddr
```
  Store multiple 1-element structures (from one register)
```
ST1 {Vt.<T>, Vt+1.<T>}, vaddr
```
  Store multiple 1-element structures (from two consecutive registers)
```
ST1 {Vt.<T>, Vt+1.<T>, Vt+2.<T>}, vaddr
```
  Store multiple 1-element structures (from three consecutive registers)
```
ST1 {Vt.<T>, Vt+1.<T>, Vt+2.<T>, Vt+3.<T>}, vaddr
```
  Store multiple 1-element structures (from four consecutive registers)
```
ST2 {Vt.<T>, Vt+1.<T>}, vaddr
```
  Store multiple 2-element structures (from two consecutive registers)
```
ST2 {Vt.<T>, Vt+2.<T>}, vaddr
```
  Store multiple 2-element structures (from two alternating registers)
```
ST3 {Vt.<T>, Vt+1.<T>, Vt+2.<T>}, vaddr
```
  Store multiple 3-element structures (from three consecutive registers)
```
ST3 {Vt.<T>, Vt+2.<T>, Vt+4.<T>}, vaddr
```
  Store multiple 3-element structures (from three alternating registers)
```
ST4 {Vt.<T>, Vt+1.<T>, Vt+2.<T>, Vt+3.<T>}, vaddr
```
  Store multiple 4-element structures (from four consecutive registers)
```
ST4 {Vt.<T>, Vt+2.<T>, Vt+4.<T>, Vt+6.<T>}, vaddr
```
  Store multiple 4-element structures (from four alternating registers)

5.7.22.2 Load-Store Single Structure

In all of these instructions <T> is one of B, H, S or D, except that type B is not available in conjunction with the alternate register variant . The post-increment immediate offset, if present, must be 1, 2, 3, 4, 6, 8, 12, 16, 24 or 32, depending on the number of elements transferred.
```
LD1 {Vt.<T>}[index], vaddr
```
  Load single 1-element structure to one lane (of one register)
```
LD2 {Vt.<T>, Vt+1.<T>}[index], vaddr
```
  Load single 2-element structure to one lane (of two consecutive registers)
```
LD2 {Vt.<T>, Vt+2.<T>}[index], vaddr
```
  Load single 2-element structure to one lane (of two alternating registers)
```
LD3 {Vt.<T>, Vt+1.<T>, Vt+2.<T>}[index], vaddr
```
  Load single 3-element structure to one lane (of three consecutive registers)
```
LD3 {Vt.<T>, Vt+2.<T>, Vt+4.<T>}[index], vaddr
```
  Load single 3-element structure to one lane (of three alternating registers)
```
LD4 {Vt.<T>, Vt+1.<T>, Vt+2.<T>, Vt+3.<T>}[index], vaddr
```
  Load single 4-element structure to one lane (of four consecutive registers)
```
LD4 {Vt.<T>, Vt+2.<T>, Vt+4.<T>, Vt+6.<T>}[index], vaddr
```
  Load single 4-element structure to one lane (of four alternating registers)
```
ST1 {Vt.<T>}[index], vaddr
```
  Store single 1-element structure from one lane (of one register)
```
ST2 {Vt.<T>, Vt+1.<T>}[index], vaddr
```
  Store single 2-element structure from one lane (of two consecutive registers)

ST2 {Vt.<T>, Vt+2.<T>}[index], vaddr

    Store single 2-element structure from one lane (of two alternating registers)

ST3 {Vt.<T>, Vt+1.<T>, Vt+2.<T>}[index], vaddr

    Store single 3-element structure from one lane (of three consecutive registers)

ST3 {Vt.<T>, Vt+2.<T>, Vt+4.<T>}[index], vaddr

    Store single 3-element structure from one lane (of three alternating registers)

ST4 {Vt.<T>, Vt+1.<T>, Vt+2.<T>, Vt+3.<T>}[index], vaddr

    Store single 4-element structure from one lane (of four consecutive registers)

ST4 {Vt.<T>, Vt+2.<T>, Vt+4.<T>, Vt+6.<T>}[index], vaddr

    Store single 4-element structure from one lane (of four alternating registers)

5.7.22.3 Load Single Structure and Replicate

In all of these instructions <T> is one of 8B, 16B, 4H, 8H, 2S, 4S, 1D or 2D. The post-increment immediate offset, if present, must be 1, 2, 3, 4, 6, 8, 12, 16, 24 or 32, depending on the number of elements transferred.

LD1R {Vt.<T>}, vaddr

    Load single 1-element structure to all lanes (of one register)

LD1R {Vt.<T>, Vt+1.<T>}, vaddr

    Load single 1-element structure to all lanes (of two consecutive registers)

LD2R {Vt.<T>, Vt+1.<T>}, vaddr

    Load single 2-element structure to all lanes (of two consecutive registers)

LD2R {Vt.<T>, Vt+2.<T>}, vaddr

    Load single 2-element structure to all lanes (of two alternating registers)

LD3R {Vt.<T>, Vt+1.<T>, Vt+2.<T>}, vaddr

    Load single 3-element structure to all lanes (of three consecutive registers)

LD3R {Vt.<T>, Vt+2.<T>, Vt+4.<T>}, vaddr

    Load single 3-element structure to all lanes (of three alternating registers)

LD4R {Vt.<T>, Vt+1.<T>, Vt+2.<T>, Vt+3.<T>}, vaddr

    Load single 4-element structure to all lanes (of four consecutive registers)

LD4R {Vt.<T>, Vt+2.<T>, Vt+4.<T>, Vt+6.<T>}, vaddr

    Load single 4-element structure to all lanes (of four alternating registers)

5.7.23 AArch32 Equivalent Advanced SIMD Mnemonics

New or changed functionality is highlighted.

| AArch32 | AArch64 | | | | | Description |
| | Integer | | | Floating point | Poly | |
	Agnostic	Unsigned	Signed			
VABA		UABA	SABA			Integer vector absolute difference and accumulate
VABAL		UABAL UABAL2	SABAL SABAL2			Integer vector absolute difference and accumulate long
VABD		UABD	SABD	FABD		Vector absolute difference
VABDL		UABDL UABDL2	SABDL SABDL2			Integer vector absolute difference long
VABS			ABS	FABS		Vector absolute value
VACGE				FACGE		Floating-point vector absolute compare greater than or equal
VACGT				FACGT		Floating-point vector absolute compare greater than
VACLE				FACLE		Floating-point vector absolute compare less than or equal
VACLT				FACLT		Floating-point vector absolute compare less than
VADD	ADD			FADD		Vector add
VADDHN	ADDHN ADDHN2					Integer vector add and narrow high half
VADDL		UADDL UADDL2	SADDL SADDL2			Integer vector add long
VADDW		UADDW UADDW2	SADDW SADDW2			Integer vector add wide

VAND	AND					Bitwise vector AND
VBIC	BIC					Bitwise vector bit clear
VBIF	BIF					Bitwise vector insert if false
VBIT	BIT					Bitwise vector insert if true
VBSL	BSL					Bitwise vector select
VCEQ	CMEQ			FCMEQ		Vector compare equal
VCGE		CMHS	CMGE	FCMGE		Vector compare greater than or equal
VCGT		CMHI	CMGT	FCMGT		Vector compare greater than
VCLE		CMLS	CMLE	FCMLE		Vector compare less than or equal
VCLS	CLS					Integer vector count leading sign bits
VCLT		CMLO	CMLT	FCMLT		Vector compare less than
VCLZ	CLZ					Integer vector count leading zero bits
VCMP				FCMP		Floating-point compare
VCMPE				FCMPE		Floating-point compare (exceptions on quiet NaNs)
VCNT	CNT					Vector count non-zero bits
VCVT.s32.f32				FCVTZS		Vector floating-point convert to signed integer (round to zero)
new				FCVTxS		Vector floating-point convert to signed integer (round to x)
VCVT.u32.f32				FCVTZU		Vector floating-point convert to unsigned integer (round to zero)

new				FCVTxU		Vector floating-point convert to unsigned integer (round to x)
VCVT.f32.i32		UCVTF	SCVTF			Vector integer convert to floating-point
VCVT.f*.f*				FCVTN FCVTL		Vector convert floating-point precision
new				FCVTXN		Vector convert double to single-precision (inexact)
new				FRINTx		Vector floating-point round to integral f-p value (towards x)
new				FDIV		Vector floating-point divide
VDUP	DUP					Duplicate single vector element to all elements
new	INS					Insert single element in another element
VEOR	EOR					Bitwise vector exclusive OR
VEXT	EXT					Bitwise vector extract
VHADD		UHADD	SHADD			Integer vector halving add
VHSUB		UHSUB	SHSUB			Integer vector halving subtract
VLD1..4	LD1..4					Vector structure /element load
VLD1..4	LD1..4R					Vector replicated element load
VLDM/VLDR	LDP/LDR					Vector load pair/register
VMAX		UMAX	SMAX	FMAX		Vector maximum
new				FMAXNM		Floating-point vector maxNum
VMIN		UMIN	SMIN	FMIN		Vector minimum
new				FMINNM		Floating-point vector minNum

VMLA	MLA			n/a		Vector chained multiply-accumulate
VFMA				FMLA		Vector fused multiply-accumulate
VMLAL		UMLAL UMLAL2	SMLAL SMLAL2			Integer vector multiply-accumulate long
VMLS	MLS			n/a		Vector chained multiply-subtract
VFMS				FMLS		Vector fused multiply-subtract
VMLSL		UMLSL UMLSL2	SMLSL SMLSL2			Integer vector multiply-subtract long
VMOV	MOV	UMOV	SMOV	FMOV		Vector move
VMOVL		UXTL UXTL2	SXTL SXTL2			Integer vector lengthen (pseudo for USHLL/SSHLL # 0)
VMOVN	XTN					Integer vector narrow
VMUL	MUL			FMUL	PMUL	Vector multiply
new				FMULX		Floating-point vector multiply extended (0xINF→2)
VMULL		UMULL UMULL2	SMULL SMULL2		PMULL	Vector multiply long
VMVN	MVN					Bitwise vector NOT
VNEG			NEG	FNEG		Vector negate
VORN	ORN					Bitwise vector OR NOT
VORR	ORR					Bitwise vector OR
VPADAL		UADALP	SADALP			Integer vector add and accumulate long pair
VPADD	ADDP			FADDP		Vector add pair
VPADDL		UADDLP	SADDLP			Integer vector add long pair
VPMAX		UMAXP	SMAXP	FMAXP		Vector max pair
new				FMAXNMP		Floating-point vector maxNum pair
VPMIN		UMINP	SMINP	FMINP		Vector min pair

new			FMINNMP		Floating-point vector minNum pair
VQABS			SQABS		Signed integer saturating vector absolute
VQADD		SQADD	UQADD		Integer saturating vector add
new		SUQADD			Signed integer saturating vector accumulate of unsigned value
new			USQADD		Unsigned integer saturating vector accumulate of signed value
VQDMLAL			SQDMLAL SQDMLAL2		Signed integer saturating vector doubling multiply accumulate long
VQDMLSL			SQDMLSL SQDMLSL2		Signed integer saturating vector doubling multiply subtract from accumulator long
VQDMULH			SQDMULH		Signed integer saturating vector doubling multiply high half
VQDMULL			SQDMULL SQDMULL2		Signed integer saturating vector doubling multiply long
VQMOVN		UQXTN UQXTN2	SQXTN SQXTN2		Integer saturating vector narrow
VQMOVUN			SQXTUN SQXTUN2		Signed integer saturating vector and unsigned narrow
VQNEG			SQNEG		Signed integer saturating vector negate
VQRDMULH			SQRDMULH		Signed integer vector saturating rounding doubling multiply high half

VQRSHL		UQRSHL	SQRSHL			Integer saturating vector rounding shift left
VQRSHRN		UQRSHRN	SQRSHRN			Integer saturating vector shift right rounded narrow
VQRSHRUN			SQRSHRUN			Signed integer saturating vector shift right rounded unsigned narrow
VQSHL		UQSHL	SQSHL			Integer saturating vector shift left
VQSHLU			SQSHLU			Signed integer saturating vector shift left unsigned
VQSHRN		UQSHRN	SQSHRN			Integer saturating vector shift right narrow
VQSHRUN			SQSHRUN			Signed integer saturating vector shift right unsigned narrow
VQSUB		UQSUB	SQSUB			Integer saturating vector subtract
VRADDHN	RADDHN					Integer vector rounding add and narrow high half
VRECPE		URECPE		FRECPE		Vector reciprocal estimate
VRECPS				FRECPS		Floating-point vector reciprocal step (FRECPS uses fused mac; VRECPS remains non-fused)
new				FRECPX		Floating-point reciprocal exponent
new	RBIT					Vector reverse bits in bytes
VREV16 **VREV32** **VREV64**	REV16 REV32 REV64					Vector reverse elements
VRHADD		URHADD	SRHADD			Integer rounding vector halving add
VRSHL		URSHL	SRSHL			Integer rounding vector shift left

VRSHR		URSHR	SRSHR			Integer rounding vector shift right
VRSHRN	RSHRN RSHRN2					Integer rounding vector shift right narrow
VRSQRTE		URSQRTE		FRSQRTE		Vector reciprocal square root estimate
VRSQRTS				FRSQRTS		Floating-point reciprocal square root step (FRSQRTS uses fused mac; VRSQRTS remains non-fused)
VRSRA		URSRA	SRSRA			Integer rounding vector shift right and accumulate
VRSUBHN	RSUBHN RSUBHN2					Integer rounding vector subtract and narrow high half
VSHL	SHL					Integer vector shift left
VSHLL		USHLL	SSHLL			Integer vector shift left long
VSHR		USHR	SSHR			Integer vector shift right
VSHRN	SHRN SHRN2					Integer vector shift right narrow
VSLI	SLI					Integer vector shift left and insert
new				FSQRT		Floating-point vector square root
VSRA		USRA	SSRA			Integer vector shift right and accumulate
VSRI	SRI					Integer vector shift right and insert
VST1..4	ST1..4					Vector structure store
VSTM/VSTR	STP/STR					Vector store pair/register
VSUB	SUB			FSUB		Vector subtract
VSUBHN	SUBHN SUBHN2					Integer vector subtract and narrow high half

VSUBL		USUBL USUBL2	SSUBL SSUBL2			Integer vector subtract long
VSUBW		USUBW USUBW2	SSUBW SSUBW2			Integer vector subtract wide
VSWP	n/a					Vector swap
VTBL	TBL					Vector table lookup
VTBX	TBX					Vector table extension
VTRN	TRN1 TRN2					Vector element transpose
VTST	CMTST					Vector test bits
VUZP	UZP1 UZP2					Vector element unzip
VZIP	ZIP ZIP2					Vector element zip
new	ADDV					Integer sum elements in vector
new		SADDLV	UADDLV			Integer sum elements in vector long
new		SMAXV	UMAXV	FMAXV		Maximum element in vector
new				FMAXNMV		Floating-point maxNum element in vector
new		SMINV	UMINV	FMINV		Minimum element in vector
new				FMINNMV		Floating-point minNum element in vector

5.7.24 Crypto Extension

The optional Crypto extension shares the FP/SIMD register file. For more information see [AES], [GCM] and [SHA].

```
PMULL Vd.1Q, Vn.1D, Vm.1D
```
Polynomial multiply long (vector): AES-GCM acceleration 64x64 to 128-bit.

```
PMULL2 Vd.1Q, Vn.2D, Vm.2D
```
Polynomial multiply long (vector, second part). Upper lanes AES-GCM acceleration 64x64 to 128-bit.

```
AESE Vd.16B, Vn.16B
```
AES single round encryption.

```
AESD Vd.16B, Vn.16B
```
AES single round decryption.

```
AESMC Vd.16B, Vn.16B
```
AES mix columns.

```
AESIMC Vd.16B, Vn.16B
```
AES inverse mix columns.

```
SHA256H Qd, Qn, Vm.4S
```
SHA256 hash update accelerator.

```
SHA256H2 Qd, Qn, Vm.4S
```
SHA256 hash update accelerator, upper part.

```
SHA256SU0 Vd.4S, Vn.4S
```
SHA256 schedule update accelerator, first part

```
SHA256SU1 Vd.4S, Vn.4S, Vm.4S
```
SHA256 schedule update accelerator, second part

```
SHA1C Qd, Sn, Vm.4S
```
SHA1 hash update accelerator (choose).

```
SHA1P Qd, Sn, Vm.4S
```
SHA1 hash update accelerator (parity).

```
SHA1M Qd, Sn, Vm.4S
```
SHA1 hash update accelerator (majority).

```
SHA1H Sd, Sn
```
SHA1 hash update accelerator (rotate left by 30).

```
SHA1SU0 Vd.4S, Vn.4S, Vm.4S
```
SHA1 schedule update accelerator, first part

```
SHA1SU1 Vd.4S, Vn.4S
```
SHA1 schedule update accelerator, second part

5.8 System Instructions

The following instruction groups are supported:

- Exception generating instructions
- System register access
- System management
- Architectural hints
- Barriers and CLREX

In several of the system instructions described in this section, the following terms are used to describe operands:

op0

>A 2-bit opcode field with an immediate value 2 or 3.

op1, op2

>A 3-bit opcode field with an immediate value in the range 0 to 7.

Cn

>A 4-bit opcode field named for historical reasons C0 − C15.

Cm

>A 4-bit opcode field named for historical reasons C0 − C15.

5.8.1 Exception Generation and Return

5.8.1.1 Non-debug exceptions

SVC #uimm16

>Generate exception targeted at exception level 1 (system), with 16-bit payload in uimm16.

HVC #uimm16

>Generate exception targeted at exception level 2 (hypervisor) , with 16-bit payload in uimm16.

SMC #uimm16

>Generate exception targeted at exception level 3 (secure monitor), with 16-bit payload in uimm16.

ERET

>Exception return: reconstructs the processor state from the current exception level's SPSR_ELn register, and branches to the address in ELR_ELn.

5.8.1.2 Debug exceptions

BRK #uimm16

>Monitor mode software breakpoint: exception routed to a debug monitor executing in EL1 or EL2, with 16-bit payload in uimm16.

HLT #uimm16

>Halting mode software breakpoint: enters halting mode debug state if enabled, else treated as UNALLOCATED. With 16-bit payload in uimm16.

DCPS1 {#uimm16}

>Debug Change Processor State to EL1 (valid in halting mode debug state only), the optional 16-bit immediate uimm16 defaults to zero and is ignored by the hardware.

```
DCPS2 {#uimm16}
```

> Debug Change Processor State to EL2 (valid in halting mode debug state only), the optional 16-bit immediate `uimm16` defaults to zero and is ignored by the hardware.

```
DCPS3 {#uimm16}
```

> Debug Change Processor State to EL3 (valid in halting mode debug state only), the optional 16-bit immediate `uimm16` defaults to zero and is ignored by the hardware.

```
DRPS
```

> Debug Restore Processor State: restores the processor to the exception level and mode recorded in the current exception level's `SPSR_ELn` register (valid in halting mode debug state only).

5.8.2 System Register Access

```
MRS Xt, <system_register>
```

> Move `<system_register>` to `Xt`, where `<system_register>` is a system register name, or for implementation-defined registers a name of the form "S<op0>_<op1>_<Cn>_<Cm>_<op2>", e.g. "S3_4_c13_c9_7".

```
MSR <system_register>, Xt
```

> Move `Xt` to `<system_register>`, where `<system_register>` is a system register name, or for implementation-defined registers a name of the form "S<op0>_<op1>_<Cn>_<Cm>_<op2>", e.g. "S3_4_c13_c9_7"..

```
MSR DAIFClr, #uimm4
```

> Uses `uimm4` as a bitmask to select the clearing of one or more of the `DAIF` exception mask bits: bit 3 selects the `D` mask, bit 2 the `A` mask, bit 1 the `I` mask and bit 0 the `F` mask.

```
MSR DAIFSet, #uimm4
```

> Uses `uimm4` as a bitmask to select the setting of one or more of the `DAIF` exception mask bits: bit 3 selects the `D` mask, bit 2 the `A` mask, bit 1 the `I` mask and bit 0 the `F` mask.

```
MSR SPSel, #uimm4
```

> Uses `uimm4` as a control value to select the stack pointer: if bit 0 is set it selects the current exception level's stack pointer, if bit 0 is clear it selects shared EL0 stack pointer. Bits 1 to 3 of `uimm4` are reserved and should be zero.

5.8.3 System Management

Where the operands of a `SYS` instruction match an entry in the `<xx_op>` tables below, then the associated alias is the preferred disassembly. Otherwise the `SYS` or `SYSL` mnemonics shall be used, permitting generation and disassembly of arbitrary implementation-defined system instructions.

```
SYS #op1, Cn, Cm, #op2{, Xt}
```

> Perform system maintenance instruction with optional source register `Xt` (defaulting to `XZR`), with the operation selected by `op1`, `Cn`, `Cm`, and `op2`.

```
SYSL Xt, #op1, Cn, Cm, #op2
```

> Perform system maintenance instruction returning a result in destination register `Xt`, with the operation selected by `op1`, `Cn`, `Cm`, and `op2`.

```
IC <ic_op>{, Xt}
```

> Instruction cache maintenance instruction, where `Xt` is the address argument as required (defaulting to `XZR`) and `<ic_op>` is defined as:
> ```
> <ic_op>    ::= <function><type><point>{<domain>}
> <function> ::= "I" (invalidate)
> <type>     ::= "ALL" (entire cache) | "VA" (by virtual address)
> ```

```
<point>    ::= "U" (to point of unification)
<domain>   ::= "IS" (inner sharable)
```

This is the preferred alias for the SYS instruction with the following operand values:

<ic_op>	op1	Cn	Cm	op2	{Xt}
IALLUIS	0	C7	C1	0	
IALLU	0	C7	C5	0	
IVAU	3	C7	C5	1	✔

DC <dc_op>, Xt

Data cache maintenance instruction, where Xt is the address argument and <dc_op> is defined as:

```
<dc_op>     ::= <function><type>{<point>}
<function>  ::= "I" (invalidate) | "C" (clean) | "CI" (clean & invalidate)
                | "Z" (zero)
<type>      ::= "VA" (by virtual address) | "SW" (by set/way)
<point>     ::= "C" (to point of coherency) | "U" (to point of unification)
```

This is the preferred alias for the SYS instruction with the following operand values:

<dc_op>	op1	Cn	Cm	op2
ZVA	3	C7	C4	1
IVAC	0	C7	C6	1
ISW	0	C7	C6	2
CVAC	3	C7	C10	1
CSW	0	C7	C10	2
CVAU	3	C7	C11	1
CIVAC	3	C7	C14	1
CISW	0	C7	C14	2

AT <at_op>, Xt

Address Translation instruction, where Xt is the address argument and <at_op> is defined as:

```
<at_op>     ::= <type><level><readwrite>
<type>      ::= "S1" (stage 1 translation) | "S12" (stage 1 and 2 translation)
<level>     ::= "E0" (exception level 0) | "E1" (exception level 1)
                | "E2" (exception level 2) | "E3" (exception level 3)
<readwrite> ::= "R" (read) | "W" (write)
```

This is the preferred alias for the SYS instruction with the following operand values：

<at_op>	op1	Cn	Cm	op2
S1E1R	0	C7	C8	0
S1E2R	4	C7	C8	0
S1E3R	6	C7	C8	0
S1E1W	0	C7	C8	1
S1E2W	4	C7	C8	1
S1E3W	6	C7	C8	1
S1E0R	0	C7	C8	2
S1E0W	0	C7	C8	3
S12E1R	4	C7	C8	4
S12E1W	4	C7	C8	5
S12E0R	4	C7	C8	6
S12E0W	4	C7	C8	7

```
TLBI <tlbi_op>{, Xt}
```

TLB invalidation instruction, where `Xt` is the address argument if required (defaulting to `XZR`).

```
<tlbi_op>   ::= <type><level>{<domain>}
<type>      ::= "ALL" (all translations at level)
                | "VMALL" (all stage 1 translations, current VMID)
                | "VMALLS12" (all stage 1 & 2 translations, current VMID)
                | "ASID" (translations matching ASID)
                | "VA" (translations matching VA and ASID)
                | "VAL" (last-level translations matching VA and ASID)
                | "VAA" (translations matching VA, all ASIDs)
                | "VAAL" (last-level translations matching VA, all ASIDs)
                | "IPAS2" (stage 2 translations matching IPA, current VMID)
                | "IPAS2L" (last-level stage 2 translations matching IPA, current VMID)
<level>     ::= "E0" (exception level 0) | "E1" (exception level 1)
                | "E2" (exception level 2) | "E3" (exception level 3)
<domain>    ::= "IS" (inner sharable)
```

This is the preferred alias for the `SYS` instruction with the following operand values:

<tlbi_op>	op1	Cn	Cm	op2	{Xt}
IPAS2E1IS	4	C8	C0	1	✔
IPAS2LE1IS	4	C8	C0	5	✔
VMALLE1IS	0	C8	C3	0	
ALLE2IS	4	C8	C3	0	
ALLE3IS	6	C8	C3	0	
VAE1IS	0	C8	C3	1	✔
VAE2IS	4	C8	C3	1	✔
VAE3IS	6	C8	C3	1	✔
ASIDE1IS	0	C8	C3	2	✔
VAAE1IS	0	C8	C3	3	✔
ALLE1IS	4	C8	C3	4	
VALE1IS	0	C8	C3	5	✔
VAALE1IS	0	C8	C3	7	✔
VMALLE1	0	C8	C7	0	
ALLE2	4	C8	C7	0	
VALE2IS	4	C8	C3	5	✔
VALE3IS	6	C8	C3	5	✔
VMALLS12E1IS	4	C8	C3	6	
ALLE3	6	C8	C7	0	
IPAS2E1	4	C8	C4	1	✔
IPAS2LE1	4	C8	C4	5	✔
VAE1	0	C8	C7	1	✔
VAE2	4	C8	C7	1	✔
VAE3	6	C8	C7	1	✔
ASIDE1	0	C8	C7	2	✔
VAAE1	0	C8	C7	3	✔
ALLE1	4	C8	C7	4	
VALE1	0	C8	C7	5	✔
VALE2	4	C8	C7	5	✔
VALE3	6	C8	C7	5	✔
VMALLS12E1	4	C8	C7	6	
VAALE1	0	C8	C7	7	✔

5.8.4 Architectural Hints

NOP

> No Operation. May be used to enforce instruction alignment, but has no execution timing constraints and so may be safely deleted from the instruction stream.

YIELD

> Yield hint.

WFE

> Wait For Event.

WFI

> Wait For Interrupt.

SEV

> Send Event: send event globally. Note that in ARMv8 a DSB and SEV instruction are in most cases not required following a synchronization operation such as unlocking a spin-lock or releasing a semaphore. A memory transaction which clears a processor's global exclusive monitor also implicitly generates an event for that processor, as held in the Event register and used by the WFE instruction.

SEVL

> Send Event Local: send event locally, without being required to affect other processors, for example to prime a wait-loop which starts with a WFE instruction.

HINT #uimm7

> Unallocated hint, where uimm7 is in the range 6-127. The unallocated hint instructions behave as a NOP but might be allocated to other hint functionality in future revisions of the architecture.

5.8.5 Barriers and CLREX

CLREX {#uimm4}

> Clear Exclusive: clears the local record of the executing processor that an address has had a request for an exclusive access. The 4-bit immediate uimm4 defaults to 0xf if omitted, with all other values unallocated.

DSB <option>|#uimm4

> Data Synchronization Barrier, where <option> is any barrier option, as below, or a 4-bit immediate uimm4 for unallocated values of option:

DMB <option>|#uimm4

> Data Memory Barrier, where <option> is any barrier option, as below, or a 4-bit immediate uimm4 for unallocated values of option.

ISB {SY|#uimm4}

> Instruction Synchronization Barrier, where SY encoded as value 0xf is the default, or a 4-bit immediate uimm4 for other unallocated values of option.

The following table defines the allocated values of data barrier option. Unallocated values behave as SY but might be allocated to other barrier functionality in future revisions of the architecture.

<option>	Value	Shareability Domain	Ordered Accesses (before-after)
OSHLD	0x1		Load-Load, Load-Store
OSHST	0x2	Outer shareable	Store-Store
OSH	0x3		Any-Any
NSHLD	0x5		Load-Load, Load-Store
NSHST	0x6	Non-shareable	Store-Store
NSH	0x7		Any-Any
ISHLD	0x9		Load-Load, Load-Store
ISHST	0xa	Inner shareable	Store-Store
ISH	0xb		Any-Any
LD	0xd		Load-Load, Load-Store
ST	0xe	Full system	Store-Store
SY	0xf		Any-Any

6 A32 & T32 INSTRUCTION SETS

Some of the new functionality found in the A64 instruction set is independent of the general purpose register width, and is therefore equally applicable to AArch32 state, namely the enhanced barrier types and load-acquire/store-release, the new IEEE 754-2008 operations, and the cryptography extensions. These new functions are added as part of ARMv8 to the A32 and T32 instruction sets as described in this section.

Note that the A32 and T32 assembler syntax remains unchanged from ARMv7 UAL. The syntax term `<c>` where used below represents a standard ARM condition code – mnemonics which do not include `<c>` may not be conditionally executed.

6.1 Partial Deprecation of `IT`

In conjunction with the reduction of conditionality in the A64 instruction set, and to facilitate higher performance implementations of the architecture in the future, ARMv8 deprecates some uses of the T32 `IT` instruction. All uses of `IT` that apply to other than a single subsequent 16-bit instruction from a restricted set are deprecated, as are explicit references to R15 (i.e. PC) within that single 16-bit instruction. This permits the non-deprecated forms of `IT` and subsequent instruction to be treated by the processor as a single 32-bit conditional instruction. The restricted set of 16-bit instructions which are *not* deprecated when used in conjunction with `IT` are as follows:

Permitted 16-Bit Instructions	Class	But deprecated…
`MOV, MVN`	Move	when `Rm` or `Rd` is PC
`LDR, LDRB, LDRH, LDRSB, LDRSH`	Load	for PC-relative "load literal" forms
`STR, STRB, STRH`	Store	
`ADD, ADC, RSB, SBC, SUB`	Add/Subtract	`ADD/SUB SP,SP,#imm` or when `Rm`, `Rdn` or `Rdm` is PC
`CMP, CMN`	Compare	when `Rm` or `Rn` is PC
`MUL`	Multiply	
`ASR, LSL, LSR, ROR`	Shift	
`AND, BIC, EOR, ORR, TST`	Logical	
`BX, BLX`	Branch to register	when `Rm` is PC

The `IT` instruction remains fully available in order to execute ARMv7 T32 code, but to verify conformance with the deprecation a new control bit permits privileged software to disable the deprecated forms of the `IT` instruction, causing them to generate an Undefined Instruction exception.

6.2 Load-Acquire / Store-Release

These new instructions provide similar functionality to the A64 instructions described in section 5.2.8 above. Natural alignment is required in all cases, and to 8 bytes in the case of `LDRAEXD` and `STRLEXD`: an unaligned address will cause an alignment fault.

6.2.1 Non-Exclusive

`LDRA<c> Rt, [Rn{,#0}]`

> Load-Acquire Word: loads a word from memory addressed by `Rn` into `Rt`.

`LDRAB<c> Rt, [Rn{,#0}]`

> Load-Acquire Byte: loads a byte from memory addressed by `Rn` and zero-extends it into `Rt`.

`LDRAH<c> Rt, [Rn{,#0}]`

> Load-Acquire Halfword: loads a halfword from memory addressed by `Rn` and zero-extends it into `Rt`.

`STRL<c> Rt, [Rn{,#0}]`

> Store-Release Word: stores a word from `Rt` to memory addressed by `Rn`.

`STRLB<c> Rt, [Rn{,#0}]`

> Store-Release Byte: stores a byte from `Rt` to memory addressed by `Rn`.

`STRLH<c> Rt, [Rn{,#0}]`

> Store-Release Halfword: stores a halfword from `Rt` to memory addressed by `Rn`.

6.2.2 Exclusive

`LDRAEX<c> Rt, [Rn{,#0}]`

> Load-Acquire Exclusive Word: loads a word from memory addressed by `Rn` into `Rt`. Records the physical address as an exclusive access.

`LDRAEXB<c> Rt, [Rn{,#0}]`

> Load-Acquire Exclusive Byte: loads a byte from memory addressed by `Rn` and zero-extends it into `Rt`. Records the physical address as an exclusive access.

`LDRAEXH<c> Rt, [Rn{,#0}]`

> Load-Acquire Exclusive Halfword: loads a halfword from memory addressed by `Rn` and zero-extends it into `Rt`. Records the physical address as an exclusive access.

`LDRAEXD<c> Rt, Rt2, [Rn{,#0}]`

> Load-Acquire Exclusive Double: loads two words from memory addressed by `base` to `Rt` and `Rt2`. Records the physical address as an exclusive access. The register `Rt` must be an even-numbered register less than 14 and `Rt2` must be $R_{(t+1)}$.

`STRLEX<c> Rd, Rt, [Rn{,#0}]`

> Store-Release Exclusive: stores a word from `Rt` to memory addressed by `Rn`, and sets `Rd` to the returned exclusive access status.

`STRLEXB<c> Rd, Rt, [Rn{,#0}]`

> Store-Release Exclusive Byte: stores a byte from `Rt` to memory addressed by `Rn`, and sets `Rd` to the returned exclusive access status.

`STRLEXH<c> Rd, Rt, [Rn{,#0}]`

> Store-Release Exclusive Halfword: stores a halfword from `Rt` to memory addressed by `Rn`, and sets `Rd` to the returned exclusive access status.

`STRLEXD<c> Rd, Rt, Rt2, [Rn{,#0}]`

> Store-Release Exclusive Double: stores two words from `Rt` and `Rt2` to memory addressed by `Rn`, and sets `Rd` to the returned exclusive access status. The register `Rt` must be an even-numbered register less than 14 and `Rt2` must be $R_{(t+1)}$.

6.3 VFP Scalar Floating-point

6.3.1 Floating-point Conditional Select

The new `VSEL` instruction is equivalent of the A64 `FCSEL` instruction in section 5.6.11, For A32 it provides an alternative to a pair of conditional `VMOV` instructions, while for T32 as it does not use an `IT` prefix it compensates for the partial deprecation of `IT` described in §6.1 above. The condition code `<fc>` may be one of `GE`, `GT`, `EQ` and `VS` only; the effect of the inverted conditions `LT`, `LE`, `NE` and `VC` may be achieved by reversing the order of the source operands.

```
VSEL<fc>.F32   Sd, Sn, Sm
```

       Single-precision conditional select: `Sd = if <fc> then Sn else Sm`.

```
VSEL<fc>.F64   Dd, Dn, Dm
```

       Double-precision conditional select: `Dd = if <fc> then Dn else Dm`.

6.3.2 Floating-point minNum/maxNum

The new `VMAXNNM` and `VMINNM` instructions implement the `minNum(x,y)` and `maxNum(x,y)` operations defined by the IEEE 754-2008 standard, and are equivalent to A64's `FMAXNM` and `FMINNM` instructions. They return the numerical operand when one operand is numerical and the other is a quiet NaN, but otherwise the result is identical to VFP `VMAX` and `VMIN`. These instructions may not be conditional.

```
VMAXNM.F32    Sd, Sn, Sm
```

      Single-precision maxNum (scalar).

```
VMAXNM.F64    Dd, Dn, Dm
```

      Double-precision maxNum (scalar).

```
VMINNM.F32    Sd, Sn, Sm
```

      Single-precision minNum (scalar).

```
VMINNM.F64    Dd, Dn, Dm
```

      Double-precision minNum (scalar).

6.3.3 Floating-point Convert (floating-point to integer)

These new instructions extend the existing ARMv7 VFP `VCVT` instructions by providing four additional explicit rounding modes, where ARMv7 `VCVT` rounds towards zero, giving an equivalent set of options to the A64 `FCVTS` and `FCVTU` instructions described in section 5.6.4.2. The syntax term `<r>` selects the rounding direction as follows: `N` (nearest, ties to even), `A` (nearest, ties away from zero), `P` (towards +Inf) or `M` (towards –Inf). These instructions may not be conditional.

```
VCVT<r>.S32.F64      Sd, Dm
```

      Convert double-precision floating-point to signed 32-bit integer with explicit rounding direction (scalar).

```
VCVT<r>.S32.F32      Sd, Sm
```

      Convert single-precision floating-point to signed 32-bit integer with explicit rounding direction (scalar).

```
VCVT<r>.U32.F64      Sd, Dm
```

      Convert double-precision floating-point to unsigned 32-bit integer with explicit rounding direction (scalar).

```
VCVT<r>.U32.F32      Sd, Sm
```

      Convert single-precision floating-point to unsigned 32-bit integer with explicit rounding direction (scalar).

6.3.4 Floating-point Convert (half-precision to/from double-precision)

The VFP `VCVTT` and `VCVTB` instructions are extended to permit direct conversion between half-precision and double-precision floating-point as a single operation, preventing double rounding errors. The syntax term `<y>` below is either `T` (top half) or `B` (bottom half).

```
VCVT<y><c>.F64.F16        Dd, Sm
```

> Convert from half-precision value in top or bottom of `Sm` to double-precision in `Dd` (scalar).

```
VCVT<y><c>.F16.F64        Sd, Dm
```

> Convert from double-precision value in `Dm` to in half-precision value in top or bottom of `Sd` (scalar).

6.3.5 Floating-point Round to Integral

The new "round to integral" instructions round a floating-point value to the nearest integral floating-point value of the same size, equivalent to the A64 `FRINT*` instructions in section 5.6.5. The only floating-point exceptions that can be raised by these instructions are `FPSCR.IOC` (Invalid Operation) for a Signaling NaN input, or `FPSCR.IDC` (Input Denormal) for a denormal input when flush-to-zero mode is enabled. For `VRINTX` only the `FPSCR.IXC` (Inexact) exception may be raised if the result is numeric and does not have the same numerical value as the source. A zero input gives a zero result with the same sign, an infinite input gives an infinite result with the same sign, and a NaN is propagated as for normal arithmetic.

A subset of the rounding instructions may be conditional when the syntax term `<x>` selects the rounding direction as follows: `Z` (towards zero), `R` (FPSCR rounding mode), or `X` (FPSCR rounding mode and signal inexactness).

```
VRINT<x><c>.F64.F64  Dd, Dm
```

> Round a double-precision value to nearest integral double-precision value (scalar), with half-way cases rounding according to `<x>`.

```
VRINT<x><c>.F32.F32  Sd, Sm
```

> Round a single-precision value to nearest integral single-precision value (scalar), with half-way cases rounding according to `<x>`.

The remaing rounding instructions are not conditional when syntax term `<r>` selects the rounding direction as follows: `N` (nearest, ties to even), `A` (nearest, ties away from zero), `P` (towards +Inf) or `M` (towards –Inf).

```
VRINT<r>.F64.F64        Dd, Dm
```

> Round a double-precision value to nearest integral double-precision value (scalar), with half-way cases rounding according to `<r>`.

```
VRINT<r>.F32.F32        Sd, Sm
```

> Round a single-precision value to nearest integral single-precision value (scalar), with half-way cases rounding according to `<r>`.

6.4 Advanced SIMD Floating-Point

The AArch32 Advanced SIMD extension continues to support only single-precision (32-bit) floating-point data types, with fixed operating modes of Round to Nearest, Default NaN and Flush-to-Zero. However it is extended with the addition of the following new instructions.

6.4.1 Floating-point minNum/maxNum

Vector forms of the new `VMAXNM` and `VMINNM` instructions described in section 6.3.2 above.

```
VMAXNM.F32      Dd, Dn, Dm
VMAXNM.F32      Qd, Qn, Qm
```
> Single-precision maxNum (vector).

```
VMINNM.F32      Dd, Dn, Dm
VMINNM.F32      Qd, Qn, Qm
```
> Single-precision minNum (vector).

6.4.2 Floating-point Convert

Vector forms of the floating-point to integer convert instructions described in section 6.3.3 above. The syntax term `<r>` selects the rounding direction: `N` (nearest, ties to even), `A` (nearest, ties away from zero), `P` (towards +Inf) or `M` (towards –Inf).

```
VCVT<r>.S32.F32        Dd, Dm
VCVT<r>.S32.F32        Qd, Qm
```
> Convert single-precision floating-point to signed 32-bit integer with explicit rounding direction (vector).

```
VCVT<r>.U32.F32        Dd, Dm
VCVT<r>.U32.F32        Qd, Qm
```
> Convert single-precision floating-point to unsigned 32-bit integer with explicit rounding direction (vector).

6.4.3 Floating-point Round to Integral

Vector forms of the floating-point rounding instructions described in section 6.3.5 above. The syntax term `<rx>` selects the rounding direction as follows: `N` (nearest, ties to even), `A` (nearest, ties away from zero), `P` (towards +Inf) or `M` (towards –Inf), `Z` (towards zero), or `X` (nearest, ties to even, signal inexactness)

```
VRINT<rx>.F32.F32      Qd, Qm
VRINT<rx>.F32.F32      Dd, Dm
```
> Round a single-precision value to nearest integral single-precision value (vector), with half-way cases rounding according to `<rx>`.

6.5 Crypto Extension

Equivalent to the A64 cryptographic instructions listed in section 5.7.24.

AESD.8   Qd, Qm

>   AES single round decryption.

AESE.8   Qd, Qm

>   AES single round encryption.

AESIMC.8   Qd, Qm

>   AES inverse mix columns.

AESMC.8   Qd, Qm

>   AES mix columns.

SHA1C.32   Qd, Qn, Qm

>   SHA1 hash update accelerator (choose).

SHA1M.32   Qd, Qn, Qm

>   SHA1 hash update accelerator (majority).

SHA1P.32   Qd, Qn, Qm

>   SHA1 hash update accelerator (parity).

SHA1H.32   Qd, Qm

>   SHA1 hash update accelerator (rotate left by 30).

SHA1SU0.32   Qd, Qn, Qm

>   SHA1 schedule update accelerator, first part

SHA1SU1.32   Qd, Qm

>   SHA1 schedule update accelerator, second part

SHA256H.32   Qd, Qn, Qm

>   SHA256 hash update accelerator.

SHA256H2.32   Qd, Qn, Qm

>   SHA256 hash update accelerator upper part.

SHA256SU0.32 Qd, Qm

>   SHA256 schedule update accelerator, first part

SHA256SU1.32 Qd, Qn, Qm

>   SHA256 schedule update accelerator, second part

VMULL.P64   Qd, Dn, Dm

>   Polynomial multiply long, AES-GCM acceleration 64x64 to 128-bit.

6.6 System Instructions

6.6.1 Halting Debug

New halting mode debug support instructions.

DCPS1

>   Debug switch to EL1 (valid in halting mode debug state only).

DCPS2

>   Debug switch to EL2 (valid in halting mode debug state only).

DCPS3

>   Debug switch to EL3 (valid in halting mode debug state only).

HLT #uimm6

>   Halting mode software breakpoint: enters halting mode debug state if enabled, else treated as
>   UNALLOCATED. With 6-bit payload in uimm6.

6.6.2 Barriers and Hints

New barrier options and hint instructions to match those in A64, as described in section 5.8.5.

DMB <option>

>   Data Memory Barrier is extended to support the new A64 Load-Load/Store options.

DSB <option>

>   Data Synchronization Barrier is extended to support the new A64 Load-Load/Store options.

SEVL

>   Send Event Locally without being required to affect other processors, for example to prime a wait-loop
>   which starts with a WFE instruction.